

# PROJECTION-TRANSPORT STEERING: CONDITIONAL CONTROL OF REASONING BEHAVIOR IN LLMs

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Activation steering edits a language model’s residual stream at inference time, but most edits act on every input, so an edit strong enough to remove a behavior also changes inputs that never showed it. We propose Projection-Transport Steering (PTS), a framework that separates when to intervene from how. Its action is the optimal-transport map of a low-dimensional behavioral projection onto a target law, the unique minimum-displacement edit that leaves the rest of the residual stream untouched. Its decision is a likelihood-ratio test on a learned detection statistic. The selectivity of the composed operator decomposes exactly into the decision’s error rates and the action’s conversion rates, so one ungated run predicts gated selectivity and unsteered answers give the ceiling of a decision. On overconfident multiple-choice reasoning with Gemma-2-2B, PTS is the most selective of eight methods on held-out questions and lowers calibration error at non-inferior accuracy. On free-text toxicity in three model families, PTS is more selective than refusal and random directions on its own decision and, in two families, than a global shift and prompting; where the tuning rule selects the quantile transport map, it outperforms a shift of equal displacement.

## 1 INTRODUCTION

Many high-level behaviors of large language models appear to be carried by directions in their activations. Activation steering methods (Turner et al., 2023; Panickssery et al., 2024; Zou et al., 2023; Li et al., 2023) read such a direction off contrastive examples and add it to the residual stream at inference time, at the cost of one vector addition per token and no retraining. However, these methods apply the same edit to every input, so the strength needed to remove a behavior where it occurs also changes the inputs where it does not: a refusal vector at a useful strength raises refusal on harmful and harmless prompts alike (Arditi et al., 2024).

The difficulty is sharpest for a behavior that a model shows only on some inputs, such as overconfidence in multiple-choice reasoning, where the answers to change and the answers to protect are both confident, one wrong and one right. On Gemma-2-2B, the global shift that eliminates every overconfident-wrong answer also eliminates every confident-right one (Figure 1a), because a shift moves the bulk of the distribution as far as it moves the tail.

In this paper we introduce Projection-Transport Steering (PTS), which splits a steering operator into an *action* that sets how to edit and a *decision* that sets which inputs to edit. The action is a transfer function on the projection of the residual stream onto a behavioral subspace, a few directions read off the model’s own labeled outputs. The decision is a likelihood-ratio test on a detection statistic fitted to tell the target state from the protected one, and the action runs only where the test fires. Table 1 places published operators in this frame.

Given the subspace and a target law for the steered projection, the action is determined: the optimal-transport map between the two laws is the unique minimum-displacement edit that leaves the orthogonal complement untouched (Theorem 1). Additive steering, ablation and clamping are this map for particular targets, and on a one-dimensional projection the map to any target is a monotone quantile map. The decision enters through an exact identity: the selectivity of a conditional edit, the share of target outputs it removes minus the share of protected outputs it loses, equals the decision’s true-

| method                                                      | transfer function $g(p)$                                | decision              | cost / token        |
|-------------------------------------------------------------|---------------------------------------------------------|-----------------------|---------------------|
| ActAdd, CAA (Turner et al., 2023; Panickssery et al., 2024) | $p + \alpha$                                            | none                  | $\Theta(d)$         |
| ITI (Li et al., 2023)                                       | $p + \alpha$ on probe-selected heads                    | none                  | $\Theta(d)$         |
| RepE (Zou et al., 2023)                                     | $p + \alpha(x)$                                         | none                  | $\Theta(d)$         |
| directional ablation (Arditi et al., 2024)                  | 0                                                       | none                  | $\Theta(d)$         |
| SAE clamping (Templeton et al., 2024)                       | $c$                                                     | none                  | $\Theta(d d_{sae})$ |
| MiMiC (Singh et al., 2024)                                  | full-space affine map                                   | none                  | $\Theta(d^2)$       |
| Linear-AcT (Rodriguez et al., 2025a)                        | $(1-\lambda)p + \lambda(\omega p + \beta)$ , per neuron | none                  | $\Theta(d)$         |
| CAST (Lee et al., 2025)                                     | $p + \alpha$                                            | prompt similarity     | $\Theta(d)$         |
| PCHI (Li et al., 2026)                                      | learned head rescaling                                  | probe at the readout  | $\Theta(d)$         |
| <b>PTS (ours)</b>                                           | $S(p)$ , the transport map to $\nu$                     | likelihood-ratio test | $\Theta(kd)$        |

Table 1: Published steering operators differ in their transfer function and in whether a decision gates it. Each row applies a transfer function  $g$  to a  $k$ -dimensional projection  $p$  of the residual stream; PTS derives  $g$  from a target law  $\nu$  (Theorem 1) and accounts for its decision exactly (Proposition 1).




Figure 1: PTS removes more of the target behavior than the global shift at the same retention. Each marker is one method at its tuned configuration on Gemma, on held-out questions (a) and prompts (b); up and to the right is better, and dotted lines mark equal selectivity. The orange curve is one shift at increasing dose, gray points send random directions through the PTS decision, and the star is the main PTS arm (post-hoc shift in a, quantile transport map in b).

positive rate times the rate at which the action converts flagged targets, minus its false-positive rate times the rate at which it converts flagged protected outputs (Proposition 1). The identity predicts gated selectivity from one ungated run and gives the ceiling of a decision from unsteered answers (Theorem 2); with Eq. (3) it also states when steering beats overwriting every flagged output.

Our main contributions are:

- A framework and its theory: the optimal-transport characterization of projection-local steering and an exact decomposition of conditional selectivity, with the optimal decision and its ceiling (Sec. 3); the decomposition predicts measured gated selectivity with correlation 0.97 (Sec. 5.1).
- Held-out evidence on overconfident reasoning, with every configuration committed before the held-out run: PTS is the most selective of eight methods on 3,000 MMLU questions for Gemma, lowers the share of overconfident-wrong answers by 15 points, and its direction outperforms a hundred random directions sent through the same decision (Sec. 5).
- Free-text control of toxic continuation in three model families, where PTS is more selective than refusal on its own decision in each family and the quantile transport map chosen by the tuning rule outperforms a shift of equal mean displacement; and a head-to-head in the setting of PCHI, where PTS is more selective than the learned head rescaling under the same probe (Sec. 6).

## 2 RELATED WORK

**Activation steering.** ActAdd (Turner et al., 2023) adds the activation difference of one contrast prompt pair to the residual stream, and CAA (Panickssery et al., 2024) averages it over many pairs. ITI (Li et al., 2023) shifts attention heads selected by linear probes, RepE (Zou et al., 2023) reads and writes representation directions, directional ablation (Arditi et al., 2024) removes one direction from every activation, and SAE clamping (Templeton et al., 2024) fixes one sparse feature. AxBench (Wu et al., 2025) finds prompting stronger than such edits, and reliability studies (Tan et al., 2024; Braun et al., 2025) find that their effect varies widely across inputs. All of these operators apply one edit to every input, with a shape fixed by hand; Theorem 1 derives the shape from a target law.

**Transport and distribution matching.** MiMiC (Singh et al., 2024) derives the affine map that matches the mean and covariance of two activation distributions, the Gaussian optimal-transport map, and applies it to the full activation. Linear-Act (Rodríguez et al., 2025a) fits a univariate affine transport map to each neuron with a strength  $\lambda \in [0, 1]$  and writes earlier edits as special cases; LinEAS (Rodríguez et al., 2025b) learns such maps across layers end to end. COAST (Nguyen et al., 2026) moves each activation, at fixed norm, to a fixed cosine with a target direction at the smallest second-moment-weighted displacement, which in our terms transports that cosine onto a point mass, and A-LQR (Skifstad et al., 2026) casts steering as linear-quadratic control. Unlike these maps, PTS transports only a  $k$ -dimensional projection and leaves its orthogonal complement fixed; for any target law, including one with atoms, the cost-minimal edit is then optimal transport in  $k$  dimensions, in closed form for  $k = 1$ .

**Deciding when to act.** CAST (Lee et al., 2025) compares the prompt’s hidden state with a condition vector at a grid-searched threshold and adds a behavior vector when the comparison holds; we fit its condition on the labeled pool of our probe. DSAS (Ferrando et al., 2026) decouples when to steer from how, with a logistic gate that scales any steering map per token. SADI (Wang et al., 2025), dynamic activation composition (Scalena et al., 2024) and GAPS (Fereidouni et al., 2026) adapt the edit per input, per generation step and per dimension, respectively. PCHI (Li et al., 2026) conditions a learned rescaling of downstream attention heads on a frozen probe read at a confidence-template token; we run PTS in its setting (Sec. 6). Zhao et al. (2026) trace inflated verbalized confidence to components that write at late token positions. Activations predict whether an answer is right (Kadavath et al., 2022; Azaria & Mitchell, 2023) and the first generated tokens whether steering will succeed (Fan et al., 2026); probes of a finished behavior forecast it poorly from a prefix (Kortukov et al., 2026), and a steering write earlier in the same pass can degrade a direction read after it (Luo et al., 2026). Closer to our work, CAST, DSAS and PCHI already separate the decision from the action; PTS adds an exact account of the pair (Prop. 1), with the ceiling and the reference arms it implies, which we run for every conditional method.

## 3 PROJECTION-TRANSPORT STEERING

A conditional steering operator has to settle what edit to apply and to which inputs. PTS answers with a transport map on a behavioral subspace and a likelihood-ratio test on a detection statistic, and accounts for how the two compose. Proofs are in Appendix B.

**The action.** Let  $h \in \mathbb{R}^d$  be the residual-stream activation at a fixed decoder layer, and let  $V \in \mathbb{R}^{k \times d}$  have orthonormal rows spanning what we call the behavioral subspace;  $q = Vh$  is the projection. In our experiments  $k = 1$ , and the row of  $V$  is the unit difference of mean activations between confident and hesitant answers, which we call the confidence direction, with  $q$  the confidence coordinate; on free text the direction separates toxic from clean continuations. We write  $\mu$  for the law of  $h$  under unsteered generation,  $\mu_V$  for the induced law of  $q$ , and  $\nu$  for a target law on  $\mathbb{R}^k$ . Our goal is an edit that gives the projection the law  $\nu$ , leaves the orthogonal complement untouched, and moves  $h$  as little as possible, since large edits damage generation (Sec. 5).

**Definition 1** (admissible steering).  $T : \mathbb{R}^d \rightarrow \mathbb{R}^d$  is admissible for  $(V, \nu)$  if (i)  $(I - V^\top V) T(h) = (I - V^\top V) h$   $\mu$ -a.s. (locality: the orthogonal complement is untouched), and (ii)  $(V \circ T)_\# \mu = \nu$  (goal: steered projections follow the target law). Its cost is  $C(T) = \mathbb{E}_\mu \|T(h) - h\|^2$ .

**Theorem 1** (steering is optimal transport of the projection). *Let  $\mu_V$  be absolutely continuous with finite second moment and let  $\nu$  have a finite second moment. Then  $\min_T C(T) = W_2^2(\mu_V, \nu)$  over admissible  $T$ , attained by*

$$T^*(h) = h + V^\top (S(Vh) - Vh), \quad (1)$$

where  $S$  is the ( $\mu_V$ -a.s. unique) Brenier map from  $\mu_V$  to  $\nu$ . Every optimal  $T$  satisfies  $VT(h) = S(Vh)$ , hence  $T = T^*$  by locality,  $\mu$ -almost surely.

In words, once the subspace and the target law are fixed, the edit is determined: it moves only the projection, by the transfer function  $S$  ( $g$  in Table 1). For  $k = 1$ ,  $S = F_\nu^{-1} \circ F_{\mu_V}$  is the monotone rearrangement, stored as a 41-knot piecewise-linear lookup; atomlessness of  $\mu_V$  then suffices and  $\nu$  may have atoms. Directional ablation is the map onto  $\delta_0$ , a clamp at  $c$  the map onto the law of  $\min(q, c)$ , and the shift  $S(q) = q + \alpha$  the map onto a translate of  $\mu_V$ ; for Gaussian marginals  $S$  is affine (Knott & Smith, 1984; Gelbrich, 1990). Since every nondecreasing transfer function is the optimal map onto its own image law, Theorem 1 turns the choice of edit into a choice of target law, made on a tuning split (Sec. 4). For the shift, the theorem gives the following corollary.

**Corollary 1** (when additive steering is optimal, and its limit). *(a)  $S(q) = q + \alpha$  is the optimal map iff  $\nu$  is the translate of  $\mu_V$  by  $\alpha$ . (b) Let two subpopulations have the same law of  $q$ . If, under an edit, the probability that an input leaves its state depends on  $h$  only through  $q$ , the edit changes the state of both at the same rate.*

**The decision and its accounting.** A decision is needed because the two states we must tell apart look alike to the action. Overconfident-wrong (OCW) and confident-right (CR) answers are both confident, and their laws on the confidence coordinate nearly coincide (Sec. 5), so by Corollary 1(b) any transfer function of that coordinate changes both at similar rates, as the global shift does (Fig. 1). A decision therefore reads a second statistic  $s$  of the input and runs the action on an event  $\Gamma$  determined by  $s$ , in practice  $\{s > \tau\}$  for a scalar score; off  $\Gamma$  the output is exactly the unsteered one. We measure the result with selectivity,  $\text{Sel} = P(\text{leave OCW} \mid \text{OCW}) - P(\text{leave CR} \mid \text{CR})$ , the share of OCW answers removed minus the share of CR answers lost; on free text, toxic and clean continuations take their places. Write TPR and FPR for the rates at which the decision fires on OCW and CR, and  $\rho_O, \rho_C$  for the rates at which the action changes the state of a flagged OCW and a flagged CR answer.

**Proposition 1** (selectivity decomposes). *For any decision and action that leave unflagged outputs unchanged,*

$$\text{Sel} = \text{TPR} \cdot \rho_O - \text{FPR} \cdot \rho_C. \quad (2)$$

The proof is one line of total probability. The decision sets (TPR, FPR), and the action sets ( $\rho_O, \rho_C$ ) on the flagged answers. Since a gated operator with a post-hoc or prompt decision returns, input by input, the unsteered or the ungated output, one ungated run per action composes every such configuration, and the ungated rates predict its selectivity when the action converts flagged and unflagged answers of a state alike (Sec. 5.1). The decomposition also fixes three reference arms on the same decision: a no-edit arm measures what regeneration alone changes, a random direction at the same dose what any perturbation of that size changes, and an override of the reported confidence of every flagged answer ( $\rho_O = \rho_C = 1$ ; on free text, a refusal) has selectivity  $\text{TPR} - \text{FPR}$ . The action beats that override exactly when

$$\text{TPR}(1 - \rho_O) < \text{FPR}(1 - \rho_C), \quad (3)$$

that is, when the share of protected answers flagged but spared exceeds the share of target answers flagged but not converted.

**Theorem 2** (optimal decision and selectivity ceiling). *Let  $p_O, p_C$  be the densities of  $s$  on OCW and CR answers, and  $r_O(s), r_C(s)$  the probabilities that the action converts such answers at  $s$ . (a) If  $r_O$  and  $r_C$  do not depend on  $s$ , then among decisions with equal FPR the likelihood-ratio test on  $s$  maximizes Sel (Neyman–Pearson); under equal-covariance Gaussian class conditionals it is the linear discriminant. In general, among decisions with equal  $\text{FPR} \cdot \rho_C$ , the maximizer thresholds  $p_O(s)r_O(s)/(p_C(s)r_C(s))$ . (b) If  $s$  is scalar and the concave hull of its ROC has area AUC, every decision  $\{s > \tau\}$  has  $\text{TPR} - \text{FPR} \leq 2 \text{AUC} - 1$ , which caps the override on it and any action with  $\rho_O = \rho_C$ .*

Part (b) is the ceiling of the introduction: computed from unsteered answers alone, it bounds what a threshold decision contributes when the action converts both states alike, and Appendix B treats the general case in which conversion depends on  $s$ . In practice the decision is a logistic-regression probe, an estimate of the log-likelihood ratio up to a constant that  $\tau$  absorbs, on the activations of the finished answer, its first  $t$  tokens or the prompt, fitted on labeled answers disjoint from every tuning and evaluation split.

**Theorem 3** (transport under a fixed decision). *Let  $P(\Gamma) > 0$ , let  $\mu_V(\cdot | \Gamma)$  be absolutely continuous, and let  $\nu_\Gamma$  have a finite second moment. Among interventions that are local to  $V$ , the identity off  $\Gamma$ , and whose projections on  $\Gamma$  follow  $\nu_\Gamma$ , the cost-minimal one applies on  $\Gamma$  the Brenier map from  $\mu_V(\cdot | \Gamma)$  to  $\nu_\Gamma$ , at cost  $P(\Gamma) W_2^2(\mu_V(\cdot | \Gamma), \nu_\Gamma)$ .*

The action can thus be fitted to the flagged population without giving up displacement optimality. The theorem holds the decision fixed: under a linear gain and a price on displacement, a graded dose, the form of the per-token interpolation of DSAS, beats the switch exactly when its optimum lies strictly between 0 and 1 with positive probability (Prop. 4). A post-hoc decision scores the finished answer and regenerates the flagged inputs, so it generates them twice; Proposition 2 gives the single-pass alternative.

**Proposition 2** (single-pass decision). *Let  $z_t$  be computed from the activations of the prompt, or of the prompt and the first  $t$  generated tokens. Among decisions measurable with respect to  $z_t$ , Theorem 2 holds with  $z_t$  in place of  $s$ , so the ceiling of an early decision is read from the ROC of  $z_t$  on unsteered answers. If the action starts after the last position  $z_t$  reads, the operator runs in one forward pass: the decision reads only unedited activations, earlier positions keep their cached keys and values, and for  $k = 1$  each token costs at most two extra  $d$ -dimensional dot products, one lookup and one vector update.*

## 4 EXPERIMENTAL SETUP

**Models and behavior.** The main model is Gemma-2-2B-it (Gemma Team, 2024) and the main behavior is overconfidence on MMLU (Hendrycks et al., 2021). The multiple-choice comparison is repeated on Qwen2.5-7B-Instruct, a second model family on which every component is refitted, and on ARC-Challenge (Clark et al., 2018), where every decision is refitted with ARC training and validation questions added to its pool; the free-text comparison adds Mistral-7B-Instruct-v0.3 as a third family. On multiple-choice questions both models reason inside `<think>...</think>` and answer `\boxed{letter}`, with greedy decoding. The action runs at layer 14 of 26 (Gemma) and 14 of 28 (Qwen), where the confidence direction separates confident from hesitant answers best or within 0.016 AUROC of the best layer on the tuning split (Appendix K). Section 6 adds toxic continuation as a free-text behavior and the setting of PCHI (Li et al., 2026), the closest conditional method.

**Readout.** The primary readout (M5) scores the answer the model writes. Correctness compares the boxed letter with the gold answer. Confidence is  $P(\text{YES})$  when the model is then asked, in the same conversation, whether that answer is correct, and a steered question is steered in both turns. Correctness and whether  $P(\text{YES}) \geq 0.5$  define four states; the two confident ones are overconfident-wrong (OCW, wrong) and confident-right (CR, right). Appendix E reports two secondary readouts, the letter logit of the same answer and a per-option verification, and Appendix N gives the prompts.

**Splits.** The MMLU splits are disjoint subsets of the test set: 400 questions to extract the directions and transport targets, a pool of 4,000 to fit the detection probe, 1,200 on which every configuration is tuned, and 3,000 used once for confirmation. ARC uses 1,012 probe and 400 tuning questions from its training and validation sets and its 1,165 test questions for confirmation.

**Metrics.** Each question is paired with itself before and after the edit. Removal is  $P(\text{not OCW} | \text{OCW})$ , retention is  $P(\text{CR} | \text{CR})$ , and  $\text{Sel} = \text{removal} - (1 - \text{retention})$ , which is 1 for an edit that removes every OCW and keeps every CR and 0 for one that treats the two states alike. Because selectivity conditions on the starting state, we also report, over all questions, the change in the share of OCW answers, the 10-bin expected calibration error (ECE) of the stated confidence, and the change in accuracy.




Figure 2: Proposition 1 predicts gated selectivity, and the PTS probe raises the ceiling of a decision. (a) Gated selectivity of 480 tuning-split configurations against its prediction from ungated runs, with equality dashed; (b) ceiling  $\max(TPR - FPR)$  of a threshold on each statistic, from unsteered answers, next to held-out PTS selectivity; (c) held-out AUROC of the PTS probe when it reads only the prompt or the first answer tokens.

**Methods and tuning.** We compare PTS with seven baselines: prompting with a calibration instruction, additive CAA (Panickssery et al., 2024), directional ablation (Arditi et al., 2024), MiMiC (Singh et al., 2024), Linear-AcT (Rodriguez et al., 2025a), CAST with its prompt condition (Lee et al., 2025) fitted on the same labeled pool as our probe, and a CAST-style trace condition, a shift conditioned on the finished trace that, like post-hoc PTS, regenerates the questions it flags. On free text the baselines are prompting, the global shift and CAST. All methods are reimplemented on our hooks and scored with the same readout. PTS decides post-hoc, from the finished unsteered answer, and regenerates the flagged questions; at the prompt, in one forward pass over it before generation; or at a prefix of  $t \in \{16, 32, 64, 128\}$  generated tokens, which keeps decision and edit in one pass. Its action is a transport map of Theorem 1 on the projection: shifts at several doses and ablation, and on free text also the quantile map onto the projection law of clean continuations. One rule selects every configuration on the tuning split: the highest selectivity among configurations whose accuracy change is at least  $-0.01$ , or the highest overall when none meets that floor; every confirmatory number is a real run. The budgets are one configuration for prompting and for ablation, eight for CAA, nine each for MiMiC and Linear-AcT, and sixteen for each conditional method, ours and the two CAST variants alike (Appendix C).

**Reference arms and statistics.** Every conditional arm is run against three references that share its decision: a no-edit arm, which regenerates the flagged inputs with the identity action; random unit directions at the same dose (ten per conditional multiple-choice arm; one hundred for PTS, or thirty displaced as far as the selected ablation on Qwen-MMLU); and an override that marks every flagged answer as not confident, or replaces every flagged continuation by a refusal on free text. Intervals are 95% paired bootstrap intervals over questions (prompts on free text) with  $10^4$  resamples. The seven comparisons of PTS with the baselines are two-sided and Holm-adjusted within each setting, each later protocol Holm-adjusts its own tests, and an accuracy change is non-inferior when the lower end of its interval is above  $-0.02$ . Every protocol was committed before its held-out rollouts, the multiple-choice one with the hashes of its selected configurations and fitted detectors (Appendix A).

## 5 SELECTIVE REMOVAL OF OVERCONFIDENT ANSWERS

A learned detection statistic lifts the selectivity ceiling of a threshold decision above that of the confidence coordinate on both models. A global shift of that coordinate removes overconfident-wrong answers only together with confident-right ones (Fig. 1a), as Corollary 1(b) predicts: on the held-out questions the two confident states overlap on the coordinate, and the best threshold there separates them by  $\max(TPR - FPR) = 0.216$  on Gemma and  $0.072$  on Qwen. The PTS probe, fitted on a separate pool of 4,000 labeled answers, raises this ceiling to  $0.370$  and  $0.449$  (AUROC  $0.741$  and  $0.785$ ; Fig. 2b). Both ceilings are computed from unsteered answers, before any steered rollout.

| method                          | $\Delta\text{acc}$  | removal | retention | Sel $\uparrow$ [95% CI] | $\Delta\text{Sel}$ ( $p$ ) | $\Delta\text{OCW}$ (pp) $\downarrow$ | $\Delta\text{ECE}$ $\downarrow$ |
|---------------------------------|---------------------|---------|-----------|-------------------------|----------------------------|--------------------------------------|---------------------------------|
| prompting                       | +0.005              | 0.42    | 0.78      | 0.200 [0.16, 0.24]      | -0.094 (<0.001)            | -5.9*                                | -0.030*                         |
| additive CAA                    | -0.001              | 0.56    | 0.66      | 0.220 [0.18, 0.26]      | -0.074 (0.008)             | -13.0*                               | -0.153*                         |
| dir. ablation                   | +0.009              | 0.46    | 0.72      | 0.173 [0.13, 0.21]      | -0.121 (<0.001)            | -8.7*                                | -0.111*                         |
| MiMiC                           | +0.000              | 0.31    | 0.80      | 0.117 [0.08, 0.15]      | -0.177 (<0.001)            | -2.1*                                | -0.018*                         |
| Linear-AcT                      | +0.005              | 0.14    | 0.92      | 0.061 [0.04, 0.09]      | -0.233 (<0.001)            | -0.2                                 | -0.003                          |
| CAST (prompt)                   | -0.002              | 0.46    | 0.76      | 0.217 [0.18, 0.26]      | -0.077 (0.004)             | -11.2*                               | -0.137*                         |
| CAST-style (trace)              | -0.005              | 0.54    | 0.70      | 0.248 [0.21, 0.29]      | -0.046 (0.020)             | -16.6*                               | -0.094*                         |
| <b>PTS, post-hoc</b>            | +0.003              | 0.49    | 0.80      | 0.294 [0.26, 0.33]      | —                          | -15.0*                               | -0.099*                         |
| PTS, prefix                     | -0.008              | 0.56    | 0.73      | 0.294 [0.26, 0.33]      | +0.000                     | -17.2*                               | -0.136*                         |
| PTS, prompt                     | -0.012 <sup>†</sup> | 0.55    | 0.76      | 0.306 [0.27, 0.35]      | +0.012                     | -17.6*                               | -0.084*                         |
| same decision, no edit          | +0.003              | —       | —         | 0.049 [0.03, 0.07]      | -0.245 (<0.001)            | —                                    | -0.005                          |
| same decision, random direction | —                   | —       | —         | 0.131                   | -0.163 (<0.001)            | —                                    | —                               |

Table 2: PTS is the most selective of eight methods on held-out questions. Gemma, 3,000 MMLU questions (unsteered accuracy 0.559, ECE 0.354;  $n_{OCW}=963$ ,  $n_{CR}=1,399$ );  $\Delta\text{Sel}$  is the row minus PTS post-hoc, with Holm-adjusted  $p$  over the seven baselines. The middle block holds the single-pass PTS decisions, and the last two rows keep the PTS decision with no edit and with a random direction at the PTS dose (mean of 100). \*: 95% interval excludes 0; <sup>†</sup>: accuracy interval reaches  $-0.02$ .

## 5.1 GATED SELECTIVITY FROM ONE UNGATED RUN

Proposition 1 turns one ungated run into a prediction for every decision placed in front of it. On the tuning split we measured the conversion rates of six actions once, composed each with two decisions at 20 thresholds, and compared  $\text{TPR } \rho_O - \text{FPR } \rho_C$  with the measured gated selectivity (Fig. 2a): the correlation is 0.97 over 240 configurations on Gemma and 0.91 over 240 on Qwen. Because a gated arm with a post-hoc or prompt decision returns, question by question, either the unsteered or the ungated output, its selectivity can also be composed from one ungated run; on the three held-out splits, twelve real gated arms lie within 0.025 of this composition, with a median gap of 0.005 (Appendix D), so tuning a gated operator costs one generation per action.

## 5.2 PTS ON HELD-OUT QUESTIONS

With every method at the configuration it won on the tuning split, post-hoc PTS is the most selective of the eight on the held-out questions (Table 2). It reaches selectivity 0.294 [0.256, 0.334]: it removes 49% of overconfident-wrong answers and keeps 80% of confident-right ones. The closest baseline, the CAST-style trace condition, reaches 0.248, and the other six trail PTS by 0.074 to 0.233; all seven differences survive Holm adjustment, and the lead over the CAST-style condition holds at confidence cutoffs of 0.4 (+0.044 [0.007, 0.080]) and 0.6 (+0.068 [0.028, 0.107]). Over all questions, PTS lowers the share of overconfident-wrong answers by 15.0 points [13.7, 16.4] and the calibration error from 0.354 to 0.255, with an accuracy change of +0.003 [-0.009, 0.015]. A decision read from the first 32 generated tokens, in a single pass, gives the same selectivity, 0.294.

The confidence direction carries the edit. With the PTS decision and dose kept and the direction replaced by random unit directions, selectivity averages 0.131 over one hundred draws; PTS is above every draw and above their mean by 0.163 [0.140, 0.188] (gray points in Fig. 1a). On code synthesis the prompt decision keeps the edit off: it fires on none of the HumanEval problems, so pass@1 stays at its unsteered 0.396, whereas the same shift applied to every input lowers it to 0.281 (Appendix L).

## 6 PTS BEYOND THE MAIN SETTING

### 6.1 FURTHER MULTIPLE-CHOICE SETTINGS

PTS remains selective on a second model family and on a second benchmark (Table 3). On Qwen the tuning rule selects ablation as the action, and PTS reaches selectivity 0.172 [0.136, 0.210], 0.081 above regenerating the flagged questions without an edit and significantly above the CAST-style trace condition. On ARC-Challenge, with every decision refitted, PTS reaches 0.235 [0.168, 0.303], lowers the calibration error by 0.060, is significantly above Linear-AcT and the CAST prompt condition, and lies above the 95th percentile of a hundred random directions sent through its decision.

| method                          | Qwen2.5-7B, MMLU ( $n_{OCW}=563$ , $n_{CR}=2,007$ ) |                         |                        |                         | Gemma-2-2B, ARC ( $n_{OCW}=228$ , $n_{CR}=780$ ) |                         |                        |                         |
|---------------------------------|-----------------------------------------------------|-------------------------|------------------------|-------------------------|--------------------------------------------------|-------------------------|------------------------|-------------------------|
|                                 | $\Delta_{acc}$                                      | Sel $\uparrow$ [95% CI] | $\Delta_{Sel}$ ( $p$ ) | $\Delta ECE \downarrow$ | $\Delta_{acc}$                                   | Sel $\uparrow$ [95% CI] | $\Delta_{Sel}$ ( $p$ ) | $\Delta ECE \downarrow$ |
| prompting                       | -0.003                                              | 0.212 [0.17, 0.25]      | +0.040 (0.313)         | -0.004                  | +0.003                                           | 0.238 [0.17, 0.31]      | +0.003 (1.000)         | +0.009                  |
| additive CAA                    | -0.015 <sup>†</sup>                                 | 0.151 [0.12, 0.19]      | -0.021 (0.998)         | +0.012                  | -0.001 <sup>†</sup>                              | 0.294 [0.22, 0.37]      | +0.059 (0.356)         | -0.069*                 |
| dir. ablation                   | -0.009 <sup>†</sup>                                 | 0.156 [0.12, 0.19]      | -0.016 (0.998)         | +0.005                  | -0.008 <sup>†</sup>                              | 0.215 [0.14, 0.29]      | -0.020 (1.000)         | -0.055*                 |
| MiMiC                           | -0.002                                              | 0.154 [0.12, 0.19]      | -0.018 (0.998)         | +0.020*                 | +0.010                                           | 0.198 [0.13, 0.26]      | -0.037 (1.000)         | -0.007                  |
| Linear-AcT                      | -0.003                                              | 0.123 [0.09, 0.16]      | -0.049 (0.133)         | +0.009                  | +0.003                                           | 0.109 [0.06, 0.16]      | -0.126 (0.008)         | +0.000                  |
| CAST (prompt)                   | -0.011 <sup>†</sup>                                 | 0.122 [0.09, 0.16]      | -0.050 (0.133)         | +0.027*                 | -0.003                                           | 0.109 [0.05, 0.17]      | -0.126 (0.003)         | -0.028*                 |
| CAST-style (trace)              | -0.007                                              | 0.040 [0.02, 0.06]      | -0.132 (<0.001)        | +0.013*                 | +0.001                                           | 0.289 [0.22, 0.36]      | +0.054 (0.356)         | -0.061*                 |
| PTS, post-hoc                   | -0.002                                              | 0.172 [0.14, 0.21]      | —                      | -0.000                  | -0.006 <sup>†</sup>                              | 0.235 [0.17, 0.30]      | —                      | -0.060*                 |
| PTS, prefix                     | -0.008                                              | 0.109 [0.08, 0.14]      | -0.063                 | +0.013*                 | -0.008 <sup>†</sup>                              | 0.276 [0.21, 0.35]      | +0.041                 | +0.000                  |
| PTS, prompt                     | -0.034 <sup>†</sup>                                 | 0.120 [0.08, 0.16]      | -0.052                 | +0.049*                 | -0.027 <sup>†</sup>                              | 0.182 [0.11, 0.25]      | -0.053                 | +0.028*                 |
| same decision, no edit          | -0.002                                              | 0.091 [0.06, 0.12]      | -0.081 (<0.001)        | +0.000                  | -0.004                                           | 0.016 [-0.01, 0.05]     | -0.219 (<0.001)        | +0.011                  |
| same decision, random direction | —                                                   | 0.158                   | -0.014 (0.175)         | —                       | —                                                | 0.167                   | -0.068 (0.004)         | —                       |

Table 3: On a second model family and a second benchmark, PTS stays above its no-edit reference arm. Notation as in Table 2; unsteered accuracy 0.737 (Qwen) and 0.744 (ARC). The random-direction row averages 100 directions on ARC and, on Qwen, 30 directions displaced as far as the selected ablation.

| method                          | Gemma-2-2B (toxic 15.0%) |                         |                        |  | Qwen2.5-7B (toxic 20.3%) |                         |                        |  | Mistral-7B (toxic 3.4%) |                         |                        |  |
|---------------------------------|--------------------------|-------------------------|------------------------|--|--------------------------|-------------------------|------------------------|--|-------------------------|-------------------------|------------------------|--|
|                                 | rem./kept                | Sel $\uparrow$ [95% CI] | $\Delta_{Sel}$ ( $p$ ) |  | rem./kept                | Sel $\uparrow$ [95% CI] | $\Delta_{Sel}$ ( $p$ ) |  | rem./kept               | Sel $\uparrow$ [95% CI] | $\Delta_{Sel}$ ( $p$ ) |  |
| prompting                       | 42/68                    | 0.104 [0.06, 0.15]      | -0.422 (<0.001)        |  | 59/79                    | 0.386 [0.34, 0.43]      | -0.148 (<0.001)        |  | 72/70                   | 0.419 [0.33, 0.51]      | -0.101 (0.084)         |  |
| global shift                    | 90/58                    | 0.474 [0.44, 0.51]      | -0.051 (0.010)         |  | 84/57                    | 0.403 [0.37, 0.44]      | -0.131 (<0.001)        |  | 71/72                   | 0.429 [0.34, 0.52]      | -0.091 (0.134)         |  |
| CAST                            | 81/70                    | 0.515 [0.47, 0.56]      | -0.011 (0.585)         |  | 74/70                    | 0.443 [0.40, 0.48]      | -0.091 (<0.001)        |  | 64/81                   | 0.446 [0.35, 0.54]      | -0.074 (0.220)         |  |
| PTS, shift                      | 73/80                    | 0.526 [0.48, 0.57]      | —                      |  | 82/72                    | 0.534 [0.50, 0.57]      | —                      |  | 74/78                   | 0.520 [0.43, 0.61]      | —                      |  |
| PTS, transport map              | 79/79                    | 0.584 [0.54, 0.62]      | +0.058 (<0.001)        |  |                          |                         |                        |  |                         |                         |                        |  |
| PTS, post-hoc                   | 91/81                    | 0.723 [0.69, 0.75]      | +0.198 (<0.001)        |  | 76/89                    | 0.648 [0.61, 0.68]      | +0.114 (<0.001)        |  | 88/83                   | 0.712 [0.65, 0.77]      | +0.192 (<0.001)        |  |
| same decision, random direction | —                        | 0.099                   | -0.427 (<0.001)        |  | —                        | 0.184                   | -0.349 (<0.001)        |  | —                       | 0.247                   | -0.273 (<0.001)        |  |
| same decision, refusal          | 82/60                    | 0.427                   | -0.099 (<0.001)        |  | 86/59                    | 0.457                   | -0.077 (<0.001)        |  | 81/49                   | 0.305                   | -0.214 (<0.001)        |  |

Table 4: PTS has the highest single-pass selectivity on free text in all three families. 3,000 held-out Real-ToxicityPrompts prompts per model; rem./kept is the percent of toxic continuations removed and of clean continuations kept (clean and within embedding cosine 0.5 of the unsteered text), and  $\Delta_{Sel}$  is the row minus PTS with the shift. Post-hoc PTS regenerates the prompts it flags, the transport map is the tuning-rule choice on Gemma, and the last two rows keep the PTS decision.

## 6.2 FREE TEXT: TOXIC CONTINUATION

Free text is the natural test of the action, because a continuation carries no single readout that a decision could overwrite in its place. We steer toxic continuation of RealToxicityPrompts (Gehman et al., 2020), the setting in which Linear-AcT reports toxicity mitigation, in all three model families (Appendix H). A RoBERTa classifier labels a continuation toxic at score 0.5, and a clean continuation counts as kept only if the steered one is clean and within sentence-embedding cosine 0.5 of the unsteered one, so rewriting a clean text counts as damage. The direction is the difference of mean activations between toxic and clean continuations, and the decision is a probe on the prompt, read before generation.

On 3,000 held-out prompts per model (Table 4), PTS is more selective in every family than refusing on its own decision, which gives up every clean continuation the decision flags, and lies above the 95th percentile of a hundred random directions sent through that decision; on Gemma and Qwen it is also more selective than the global shift and prompting. On Gemma the tuning rule selects the quantile transport map of Theorem 1, from the law of the projection over toxic continuations to its law over clean ones. The map reaches selectivity 0.584 [0.542, 0.624], lowers the toxic rate from 15.0% to 3.5% while keeping 79% of clean continuations, and is above a shift of the same mean displacement by 0.043 [0.012, 0.076] and above CAST by 0.069 [0.021, 0.117]. On Qwen, PTS reaches 0.534 [0.496, 0.570], also significantly above CAST, and on Mistral, where toxic continuations are rare, it lowers the toxic rate from 3.4% to 1.4% at selectivity 0.520 [0.430, 0.605]. For each of the seven gated arms we compared with refusal, Eq. (3) evaluated with tuning-split conversion rates predicts the held-out winner.

## 6.3 THE SETTING OF PCHI

In the setting of PCHI (Li et al., 2026), a shift along one direction under the same probe is more selective than the learned head rescaling. PCHI gates a rescaling of 128 attention heads with a probe read where the model states its confidence. We rebuild its setting: Qwen3-4B-Instruct answers OpenMathInstruct-2 problems in a JSON format that ends with a yes/no confidence field, and a probe on the layer-18 state at that field flags overconfident-wrong answers. PTS keeps the probe and replaces the learned rescaling by a shift of the same state along one direction. On 3,183 held-out answers PTS reaches selectivity 0.438 against 0.354 for PCHI (+0.084 [0.058, 0.110]) and a calibration error of 0.188 against 0.234 (Appendix I).

#### 6.4 INFERENCE COST

With a prompt decision, PTS adds 9% to generation time. On one H100, 750 toxicity continuations with Gemma take 5.79 seconds unsteered and 6.31 seconds with PTS, of which 0.46 seconds go to the pass over the prompts that reads the decision; a post-hoc decision regenerates the flagged half and takes 8.74 seconds (Appendix J).

### 7 LIMITATIONS

Where a behavior has a single readout, as multiple-choice confidence does, overriding that readout on every answer the post-hoc PTS decision flags is at least as selective as the action (Appendix G). The evidence covers two behaviors in three model families of 2 to 7 billion parameters with one toxicity classifier, and the multiple-choice protocols became public only together with their results, so their precedence rests on local commit times.

### 8 CONCLUSION

We have introduced Projection-Transport Steering, which splits an activation edit into an action, a transport map on the projection onto a behavioral subspace, and a decision, a likelihood-ratio test that sets where the action runs. Once the subspace and the target law are fixed, the cost-minimal action is the optimal-transport map between the projection laws, and the selectivity of the composed operator equals the decision’s error rates weighted by the action’s conversion rates. This identity predicts gated configurations from one ungated run and bounds a decision from unsteered answers alone. In the main multiple-choice setting PTS was the most selective of eight methods, and on free-text toxicity it was more selective than refusal on its own decision in all three model families. We expect the same accounting to guide conditional control of other behaviors whose target and protected states share a representation.

#### AI USE STATEMENT

We used generative AI tools to polish the manuscript and to retrieve related work. The authors verified every resulting claim and reference and take responsibility for the final content.

### REFERENCES

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda, *Refusal in language models is mediated by a single direction*, Advances in Neural Information Processing Systems 37, 2024.
- Amos Azaria and Tom Mitchell, *The internal state of an LLM knows when it’s lying*, Findings of the Association for Computational Linguistics: EMNLP, 2023.
- J. Braun, C. Eickhoff, D. Krueger, S. A. Bahrainian, D. Krashennnikov. *Understanding (un)reliability of steering vectors in language models*. ICLR Workshop on Foundation Models in the Wild, 2025.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord, *Think you have solved question answering? Try ARC, the A12 reasoning challenge*, arXiv:1803.05457, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al., *Training verifiers to solve math word problems*, arXiv:2110.14168, 2021.
- Chenrui Fan, Yize Cheng, Ming Li, Soheil Feizi, and Tianyi Zhou, *When is your LLM steerable? Predicting steering success from internal states*, arXiv:2606.11599, 2026.
- Alex Ferrando de las Morenas, Xavier Suau, Jordi González, and Pau Rodríguez, *Dynamically scaled activation steering*, arXiv:2512.03661, 2026.

- Gemma Team, *Gemma 2: Improving open language models at a practical size*, arXiv:2408.00118, 2024.
- Moghis Fereidouni, Muhammad Umair Haider, Hassan Sajjad, and A. B. Siddique, *GAPS: Dimension-level gates for conditional activation steering*, arXiv:2609.01878, 2026.
- Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith, *RealToxicityPrompts: Evaluating neural toxic degeneration in language models*, Findings of the Association for Computational Linguistics: EMNLP, 2020.
- Matthias Gelbrich, *On a formula for the  $L^2$  Wasserstein metric between measures on Euclidean and Hilbert spaces*, Mathematische Nachrichten **147** (1990), no. 1, 185–203.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt, *Measuring massive multitask language understanding*, International Conference on Learning Representations, 2021.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, et al., *Language models (mostly) know what they know*, arXiv:2207.05221, 2022.
- Martin Knott and Cyril S. Smith, *On the optimal mapping of distributions*, Journal of Optimization Theory and Applications **43** (1984), 39–49.
- Evgenii Kortukov, Piotr Komorowski, Florian Klein, Paula Engl, Gabriele Sarti, Seong Joon Oh, Sebastian Lapuschkin, and Wojciech Samek, *Predicting future behaviors in reasoning models enables better steering*, arXiv:2606.11172, 2026.
- Bruce W. Lee, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Erik Miehl, Pierre Dognin, Manish Nagireddy, and Amit Dhurandhar, *Programming refusal with conditional activation steering*, International Conference on Learning Representations, 2025.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg, *Inference-time intervention: Eliciting truthful answers from a language model*, Advances in Neural Information Processing Systems 36, 2023.
- Ke Li, Chongzhe Zhang, Zifan Zeng, Feng Liu, Qunli Zhang, and Zheng Hu, *Calibrating overconfidence without sacrificing confidence: Probe-conditioned head intervention for LLMs*, arXiv:2606.09876, 2026.
- Ziyang Luo, Zhongyao Chu, Xinjie He, Youting Wang, Xukui Qin, Runxiong Wu, and Yan-Syuan Chen, *You only pass once: answering and abstaining together in a single forward pass of a frozen language model*, arXiv:2608.14465, 2026.
- Tam Nguyen, Tu Anh Nguyen, Sina Alemohammad, and Richard G. Baraniuk, *Minimizing collateral damage in activation steering*, arXiv:2605.01167, 2026.
- Nina Panickssery, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner, *Steering Llama 2 via contrastive activation addition*, Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics, 2024.
- Pau Rodriguez, Arno Blaas, Michal Klein, Luca Zappella, Nicholas Apostoloff, Marco Cuturi, and Xavier Suau, *Controlling language and diffusion models by transporting activations*, International Conference on Learning Representations, 2025.
- Pau Rodriguez, Michal Klein, Eleonora Gualdoni, Arno Blaas, Luca Zappella, Marco Cuturi, and Xavier Suau, *LinEAS: End-to-end learning of activation steering with a distributional loss*, arXiv:2503.10679, 2025.
- Daniel Scalena, Gabriele Sarti, and Malvina Nissim, *Multi-property steering of language models with dynamic activation composition*, Proceedings of BlackboxNLP, 2024.
- Shashwat Singh, Shauli Ravfogel, Jonathan Herzig, Roei Aharoni, Ryan Cotterell, and Ponnurangam Kumaraguru, *Representation surgery: Theory and practice of affine steering*, Proceedings of the 41st International Conference on Machine Learning, 2024.

- Julian Skifstad, Xinyue Annie Yang, and Glen Chou, *Local Linearity of LLMs Enables Activation Steering via Model-Based Linear Optimal Control*, arXiv:2604.19018, 2026.
- D. Tan, D. Chanin, A. Lynch, B. Paige, D. Kanoulas, A. Garriga-Alonso, R. Kirk. *Analysing the generalisation and reliability of steering vectors*. NeurIPS, 2024.
- Adly Templeton, Tom Conerly, Jonathan Marcus, et al., *Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet*, Transformer Circuits Thread, 2024.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid, *Steering language models with activation engineering*, arXiv:2308.10248, 2023.
- Weixuan Wang, Jingyuan Yang, and Wei Peng, *Semantics-adaptive activation intervention for LLMs via dynamic steering vectors*, International Conference on Learning Representations, 2025.
- Zhengxuan Wu, Aryaman Arora, Atticus Geiger, Zheng Wang, Jing Huang, Dan Jurafsky, Christopher D. Manning, and Christopher Potts, *AxBench: Steering LLMs? Even simple baselines outperform sparse autoencoders*, Proceedings of the 42nd International Conference on Machine Learning, 2025.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, et al., *Representation engineering: A top-down approach to AI transparency*, arXiv:2310.01405, 2023.
- Tianyi Zhao, Yinhan He, Wendy Zheng, Yujie Zhang, and Chen Chen, *Wired for overconfidence: A mechanistic perspective on inflated verbalized confidence in LLMs*, arXiv:2604.01457, 2026.

## A CONFIRMATORY PROTOCOL

The multiple-choice protocols (v19 for MMLU, v20 for ARC) were committed to the local repository before any steered rollout on their confirmatory splits was generated and first pushed together with the results, so their ordering rests on local commit times; the later protocols v21 to v24 were pushed to a public repository, and so timestamped by the hosting service, before their rollouts existed (`rebuttal/protocols/confirmatory-v19.md`, with two amendments made after the tuning pass and before launch), together with the SHA-256 hashes of the selected configurations and of every fitted detector (`confirmatory-v19-gemma.sha256`).

**Splits.** All splits are subsets of the 14,042-question MMLU test set, identified by index and constructed by `label.pool.split_records`. The extraction split (400) excludes the first evaluation subset of the pilot protocol; the tuning split (1,200) excludes the extraction split and the five pilot evaluation subsets; the detector pool (4,000) and the confirmatory split (3,000) are consecutive blocks of a seeded permutation of the remaining questions. No question appears in two of the four splits.

**Selection rule.** For every method the configuration with the highest tuning-split selectivity among those with  $\Delta_{\text{acc}} \geq -0.01$  is carried unchanged to the confirmatory split. Gated configurations are scored on the tuning split by combining the ungated rollouts of the same split with the gate’s flags (Prop. 1 makes this exact up to batched-decoding nondeterminism; Appendix D); single-pass configurations are real runs. Every confirmatory number is a real run.

**Tests.** Primary: PTS post-hoc against each of the seven baselines on pooled M5 selectivity, question-level paired bootstrap with  $10^4$  resamples, two-sided, Holm-adjusted over the seven comparisons. Accuracy: an arm is non-inferior if the lower end of the 95% interval of its accuracy change is above  $-0.02$ . The single-pass and prompt-decision variants, M4 and M2 selectivities, removal, retention and ECE are secondary and reported without a claim.

**Released records.** Every unsteered and steered answer of the tuning and confirmatory splits is released with its state, confidence and correctness (`results/release/records.csv.gz`); `paper/recompute.from.records.py` recomputes each selectivity and paired interval in Tables 2 and 3, including the null arms, from that file alone. `results/release/toxicity.csv.gz` holds every confirmatory continuation of the free-text setting with its judge score and state (texts omitted for the random-direction arms), and

results/release/pchi\_val.csv.gz the held-out readout gaps of every arm of the PCHI setting.

**ARC replication.** A second protocol (confirmatory-v20-arc.md), frozen before any steered ARC rollout, repeats the Gemma comparison on ARC-Challenge: a detector pool of 1,012 and a tuning split of 400 from the ARC training and validation questions, and the 1,165-question test split as the confirmatory split. The directions, moments and action layer are those fitted on MMLU; every decision is refitted on the ARC pool together with the MMLU pool, and every configuration is selected on the ARC tuning split with the same rule and grids.

**Null arms.** A third protocol (confirmatory-v21-null.md), with the job list and the hash of the runner, was committed and pushed to a public repository before any null rollout existed; the commit is timestamped by the hosting service. On each confirmatory split it adds the frozen PTS decision with an identity action, the frozen decision with the selected action on a unit random direction (seed 0), and an ungated rerun with an identity action. It registers two one-sided hypotheses per setting, PTS above each gated null, with Holm adjustment over the six tests, and the net measures of Appendix F for every arm. Five of the six tests pass; PTS against the single random direction on ARC gives  $+0.041 [-0.029, 0.109]$ , Holm  $p = 0.12$ .

**Reference arms.** A fourth protocol (confirmatory-v22-references.md), pushed before any of its rollouts existed, extends the no-edit null to all four conditional methods, replaces the single random direction by ten (seeds 0 to 9), and adds random directions for the sycophancy edit. It registers 24 one-sided tests (four methods, three settings, no-edit and random-edit null), Holm-adjusted together, of which 18 pass (Table 13); a count of the cases in which Eq. (3) with ungated conversion rates predicts the sign of steer minus matched override (10 of 13); and paired intervals for selectivity at cutoffs 0.4 and 0.6 and over boxed traces for PTS against the CAST-style condition on Gemma. The v22 random-edit null supersedes the single-direction test of v21.

**Free text and dense nulls.** A fifth protocol (confirmatory-v23-freetext.md), pushed with the hashes of the tuning selections, fitted decisions and code before any confirmatory rollout of the toxicity setting existed, fixes the data, judge, retention criterion, arms and ten primary hypotheses of Appendix H, and extends the random directions of PTS post-hoc to one hundred on Gemma-MMLU and Gemma-ARC and to thirty displacement-matched directions on Qwen (Table 14).

**Transport map and a third family.** A seventh protocol (confirmatory-v25-transport-mistral.md), pushed before its rollouts, adds the quantile transport map to the action family on free text, records its selection on the Gemma tuning split, and registers five tests: the transport map against a shift of equal mean displacement on Gemma, and PTS against the global shift, prompting, refusal on its decision and a hundred random directions on Mistral-7B. Three of the five pass after Holm adjustment, including the transport-map test ( $p = 0.003$ ).

**PCHI setting.** A sixth protocol (confirmatory-v24-pchi.md), pushed after the train-split fitting and tuning and before any validation rollout, fixes the arms and three hypotheses of Appendix I.

**Analyses added after the protocol.** The detect-and-override reference, the difference-of-means decision fitted on 4,400 labels, the latency measurement (Appendix J), the layer analysis (Appendix K), calibration and the out-of-domain negatives were run after the confirmatory results were known. They use the same splits and the frozen decisions, and we report them as exploratory.

## B PROOFS

**Theorem 1.** By locality,  $T(h) - h \in \text{span}(V^\top) \mu$ -a.s., hence  $\|T(h) - h\|^2 = \|VT(h) - Vh\|^2$  because the rows of  $V$  are orthonormal, and  $C(T) = \mathbb{E}\|VT(h) - Vh\|^2$ . The joint law  $\gamma$  of  $(Vh, VT(h))$  is a coupling of  $(\mu_V, \nu)$  by the goal constraint, so  $C(T) = \int \|y - x\|^2 d\gamma \geq W_2^2(\mu_V, \nu)$ . Because  $\mu_V$  is absolutely continuous and both laws have finite second moments, Brenier’s theorem gives a unique optimal coupling, induced by  $S = \nabla\varphi$  with  $\varphi$  convex.  $T^*(h) = h + V^\top(S(Vh) - Vh)$  is admissible ( $VV^\top = I_k$  gives  $VT^*(h) = S(Vh)$ , and the complement is untouched) and attains the bound. For any minimiser the induced coupling is optimal, hence equal to the unique one, so  $VT(h) = S(Vh)$  a.s.  $\square$

For  $k = 1$  absolute continuity can be weakened to atomlessness of  $\mu_V$ : the monotone rearrangement  $F_V^{-1} \circ F_{\mu_V}$  is then an optimal map for every  $\nu$  with a finite second moment, including targets with atoms such as  $\delta_0$  (ablation) or a law truncated at  $c$  (clamp); uniqueness of the map can fail only on a  $\mu_V$ -null set.

**Proposition 3** (metric covariance). *If the cost is  $(\Delta h)^\top M(\Delta h)$  with  $M \succ 0$ , the admissible problem is optimal transport of  $\mu_V$  to  $\nu$  in  $\mathbb{R}^k$  under the cost  $\|\cdot\|_{M_V}^2$ ,  $M_V = VMV^\top \succ 0$ , and the optimal operator is  $T^*(h) = h + V^\top(S_M(Vh) - Vh)$  with  $S_M = M_V^{-1/2} \circ \nabla\varphi \circ M_V^{1/2}$ ,  $\varphi$  convex.*

*Proof.* Locality writes  $\Delta h = V^\top a$  with  $a \in \mathbb{R}^k$ , and the goal constraint with  $VV^\top = I_k$  fixes  $a = S(Vh) - Vh$  for the map  $S$  in use. Then  $\|\Delta h\|_M^2 = a^\top M_V a$ . The cost is a latent-space Mahalanobis cost on  $\mathbb{R}^k$ ; it is not the full-space Mahalanobis distance of  $\Delta h$  unless  $M$  commutes with  $V^\top V$ . In the coordinates  $z = M_V^{1/2} q$  it is Euclidean, the change of variables is internal to  $\mathbb{R}^k$ , and Brenier’s theorem in  $z$  (the pushforward of an absolutely continuous law under an invertible linear map is absolutely continuous) gives  $\nabla\varphi$ ; mapping back gives  $S_M$ .  $\square$

**Corollary 1.** (a) If  $\nu = \mu_V * \delta_\alpha$  the map  $x \mapsto x + \alpha$  is the gradient of a convex function pushing  $\mu_V$  to  $\nu$ , hence the Brenier map. Conversely, if the optimal map is  $x \mapsto x + \alpha$  then  $\nu$  is its pushforward, the shift. For one-dimensional Gaussians the monotone map is  $x \mapsto m_t + (\sigma_t/\sigma_s)(x - m_s)$ , with slope  $\sigma_t/\sigma_s$ . (b) If two groups of inputs have the same law of  $p$  and the probability that an input changes state depends on  $h$  only through  $p$ , every transfer function changes both groups at the same rate, since both rates are the same integral against the law of  $p$ . A decision that reads  $p$  alone therefore adds little, which is why the detection statistic reads the full activation.  $\square$

**Proposition 1.** Let  $R$  be the event that an OCW trace leaves OCW and  $L$  the event that a CR trace leaves CR. Off  $\Gamma$  the output equals the unsteered output, which is in its original state by definition, so  $R \subseteq \Gamma$  and  $L \subseteq \Gamma$  up to null sets. Then  $P(R \mid OCW) = P(\Gamma \mid OCW)P(R \mid OCW, \Gamma) = \text{TPR} \rho_O$ , and likewise  $P(L \mid CR) = \text{FPR} \rho_C$ ; subtracting gives Eq. (2). If, within OCW, the indicator that the action converts the trace is independent of  $\Gamma$ , then  $\rho_O = P(R^{\text{ungated}} \mid OCW)$ , and likewise for CR; the ungated selectivity is  $\rho_O - \rho_C$ , and the difference is  $\rho_C(1 - \text{FPR}) - \rho_O(1 - \text{TPR})$ .  $\square$

The identity requires that the unflagged output be exactly the unsteered one. The post-hoc gate satisfies this by construction. The single-pass gate satisfies it for traces on which it never fires, since the action hook is the identity until the decision token.

**Theorem 2.** Write  $p_O, p_C$  for the densities of  $s$  on OCW and CR answers and  $r_O(s), r_C(s)$  for the probabilities that the action converts an answer of each state with statistic  $s$ . For a decision region  $\Gamma$ ,  $\text{TPR} \rho_O = \int_\Gamma p_O r_O$  and  $\text{FPR} \rho_C = \int_\Gamma p_C r_C$ , so  $\text{Sel}(\Gamma) = \int_\Gamma (p_O r_O - p_C r_C)$ . Among regions with a given value of  $\int_\Gamma p_C r_C$ , the Neyman–Pearson lemma applied to the nonnegative measures  $p_O r_O$  and  $p_C r_C$  shows that  $\int_\Gamma p_O r_O$  is maximised by a superlevel set of their ratio, which is the general statement in (a). If  $r_O$  and  $r_C$  are constant in  $s$ , the ratio is a constant multiple of  $p_O/p_C$ , so every objective increasing in TPR and decreasing in FPR is maximised on the likelihood-ratio ROC curve, which is (a). The assumption matters: with  $s \in \{a, b\}$ ,  $p_O(a) = p_O(b) = \frac{1}{2}$ ,  $p_C(a) = \frac{1}{4}$ ,  $p_C(b) = \frac{3}{4}$ ,  $r_O(a) = 0.2$ ,  $r_O(b) = 1$ ,  $r_C(a) = 1$  and  $r_C(b) = 0$ , the likelihood ratio ranks  $a$  above  $b$ , yet  $\Gamma = \{b\}$  gives selectivity 0.5, against  $-0.15$  for  $\{a\}$  and  $0.35$  for  $\{a, b\}$ . On our data  $r$  does depend on  $s$ : flagged answers are converted more often than unflagged answers of the same state (Appendix D). Our detector is a logistic-regression probe, a discriminative estimate of the log-likelihood ratio up to the prior term; (a) states what the probe approximates, not that the fitted probe is optimal. (b) The concavified ROC contains the path  $(0, 0) \rightarrow (\text{FPR}^*, \text{TPR}^*) \rightarrow (1, 1)$  through the Youden-optimal point, whose area is  $\frac{1}{2}(1 + J)$ , so  $J \leq 2 \text{AUC} - 1$ . The override reference has selectivity  $\text{TPR} - \text{FPR} \leq J$ , and with  $\rho_O = \rho_C = \rho$ , Eq. (2) gives  $\text{Sel} = \rho(\text{TPR} - \text{FPR}) \leq J$ .  $\square$

**Theorem 3.** Split  $C(T)$  over  $\Gamma$  and  $\Gamma^c$ . The second term is zero by the identity constraint; the first, with the conditional goal constraint, is Theorem 1 for  $\mu(\cdot \mid \Gamma)$ , which applies when  $\mu_V(\cdot \mid \Gamma)$  is absolutely continuous.  $\square$

This result fixes the decision. It does not say that a threshold gate is the best decision for the displacement-priced objective; the next statement shows that it generally is not, and when it is.

**Proposition 4** (graded dose). *Let  $T_\lambda(h) = h + \lambda(s) V^\top(S(Vh) - Vh)$  with a dose  $\lambda : \mathbb{R} \rightarrow [0, 1]$  measurable in the detector statistic  $s$ . Write  $\pi(s) = P(OCW \mid s)$ ,  $d(s) = \mathbb{E}[\|S(Vh) - Vh\|^2 \mid$*

$s] > 0$ , and suppose the behavioral gain is linear in the dose, with slope  $\beta_1 > 0$  on OCW and  $-\beta_0 < 0$  on CR. The dose maximising gain minus displacement priced at  $1/(2c)$  is

$$\lambda^*(s) = \left[ \frac{c(\beta_0 + \beta_1)(\pi(s) - \tau)}{d(s)} \right]_0^1, \quad \tau = \frac{\beta_0}{\beta_0 + \beta_1}.$$

It is zero exactly on  $\{\pi \leq \tau\}$  and nondecreasing in the ratio  $(\pi(s) - \tau)/d(s)$ . It is nondecreasing in  $\pi$  when  $d$  is constant, but not in general. The threshold gate  $\mathbf{1}[\pi > \tau]$  is optimal iff  $\lambda^* \in \{0, 1\}$  almost surely, and strictly suboptimal iff  $P(0 < \lambda^* < 1) > 0$ .

*Proof.* The displacement of  $T_\lambda$  is  $\lambda(s)^2 \|S(Vh) - Vh\|^2$ , with expectation  $\mathbb{E}_s[\lambda(s)^2 d(s)]$ ; the gain is  $\mathbb{E}_s[\lambda(s)((\beta_0 + \beta_1)\pi(s) - \beta_0)]$ . The objective is an integral over  $s$  of a strictly concave quadratic in  $\lambda$ , maximised pointwise on  $[0, 1]$  by the clipped stationary point. It is the unique maximiser a.s., so the indicator is optimal iff it agrees with  $\lambda^*$  a.s., i.e. iff  $\lambda^* \in \{0, 1\}$  a.s.; otherwise the objective is strictly larger at  $\lambda^*$  on a set of positive measure. For non-monotonicity in  $\pi$ , take  $s_1, s_2$  with  $\pi(s_1) = 0.6 < \pi(s_2) = 0.7$ ,  $\tau = 0.5$ ,  $c(\beta_0 + \beta_1) = 1$ ,  $d(s_1) = 0.2$  and  $d(s_2) = 1$ : then  $\lambda^*(s_1) = 0.5 > \lambda^*(s_2) = 0.2$ . For a counterexample to strict suboptimality at every finite price, let  $\pi$  take only the values 0.1 and 0.9 with  $\tau = 0.5$  and  $d \leq 0.4c(\beta_0 + \beta_1)$ : then  $\lambda^* \in \{0, 1\}$  a.s. and the threshold gate is optimal.  $\square$

**Proposition 2.** Let  $\mathcal{F}_t$  be generated by the first  $t$  generated tokens and  $z_t$  the mean detector feature over them. A decision measurable with respect to  $\sigma(z_t)$  is a test between the class-conditional laws of  $z_t$ , so Theorem 2 applies with  $z_t$  in place of  $s$  and nothing is assumed about how  $z_t$  relates to the finished trace. The reader hook sits at a layer at or above the actor; within one forward pass the actor at position  $t$  sees the decision accumulated over positions  $< t$ , the running mean needs one sum and one counter per sequence, and an edit at position  $t$  changes only the keys and values written at  $t$ .

**Proposition 5** (computational cost). *Any algorithm implementing a non-degenerate intervention local to a subspace spanned by a dense direction reads all  $d$  coordinates of  $h$ , so it costs  $\Omega(d)$  per token. The edit of Eq. (1) costs  $\Theta(kd + k \log K)$  with a  $K$ -knot lookup, and the detector adds one  $d$ -dimensional dot product per token.*

*Proof.* If a coordinate  $i$  with  $(\hat{v})_i \neq 0$  is never read, an adversary perturbs  $h_i$ , which changes  $p = \langle h, \hat{v} \rangle$  and, on a set of positive measure, the correct output, while the algorithm’s output is unchanged. The upper bound counts  $Vh$ , a binary search in the lookup, and  $h + V^\top(\cdot)$ .  $\square$

## C TUNING GRIDS

Every configuration searched on the tuning split, with the selected one marked  $\star$ . Gated rows are scored by composing the ungated rollouts of the same split with the decision’s flags (Appendix D); single-pass rows are real runs. The rule is the highest selectivity among configurations with  $\Delta\text{acc} \geq -0.01$ , and, when no configuration of a family meets the floor, the highest selectivity overall; this fallback applies to the additive, CAST and prompt-decision families on Qwen, where every shift in the grid costs more than one accuracy point.

Table 5: Tuning split, gemma-2-2b-it ( $n = 1,200$ , M5).

| family       | configuration         | $\Delta\text{acc}$ | removal | retention | Sel    |
|--------------|-----------------------|--------------------|---------|-----------|--------|
| additive CAA | alpha-0.375           | -0.022             | 0.831   | 0.424     | +0.255 |
| additive CAA | alpha-0.25 $\star$    | -0.002             | 0.536   | 0.661     | +0.197 |
| additive CAA | alpha-0.5             | -0.042             | 0.974   | 0.136     | +0.110 |
| additive CAA | alpha-0.125           | +0.010             | 0.315   | 0.763     | +0.078 |
| additive CAA | alpha-0.75            | -0.153             | 1.000   | 0.000     | +0.000 |
| additive CAA | alpha-1.0             | -0.205             | 1.000   | 0.000     | +0.000 |
| additive CAA | alpha-1.25            | -0.270             | 1.000   | 0.000     | +0.000 |
| additive CAA | alpha-1.5             | -0.299             | 1.000   | 0.000     | +0.000 |
| MiMiC        | L10.reg0.0001 $\star$ | +0.022             | 0.349   | 0.823     | +0.172 |
| MiMiC        | L10.reg0.1            | +0.012             | 0.349   | 0.809     | +0.158 |
| MiMiC        | L10.reg0.01           | +0.007             | 0.320   | 0.814     | +0.135 |
| MiMiC        | L18.reg0.1            | +0.022             | 0.260   | 0.860     | +0.121 |
| MiMiC        | L18.reg0.01           | +0.026             | 0.242   | 0.869     | +0.111 |
| MiMiC        | L18.reg0.0001         | +0.023             | 0.250   | 0.857     | +0.107 |
| MiMiC        | L14.reg0.01           | +0.001             | 0.224   | 0.857     | +0.081 |
| MiMiC        | L14.reg0.1            | +0.004             | 0.216   | 0.850     | +0.066 |
| MiMiC        | L14.reg0.0001         | -0.013             | 0.214   | 0.841     | +0.055 |
| Linear-AcT   | L10.lam0.5 $\star$    | +0.016             | 0.133   | 0.936     | +0.069 |

|     | family                      | configuration           | $\Delta$ acc | removal | retention | Sel    |
|-----|-----------------------------|-------------------------|--------------|---------|-----------|--------|
| 756 | Linear-AcT                  | L10.lam1.0              | +0.015       | 0.138   | 0.919     | +0.057 |
| 757 | Linear-AcT                  | L18.lam1.0              | +0.012       | 0.130   | 0.926     | +0.056 |
| 758 | Linear-AcT                  | L14.lam0.25             | +0.007       | 0.128   | 0.924     | +0.052 |
| 759 | Linear-AcT                  | L14.lam1.0              | +0.009       | 0.138   | 0.913     | +0.051 |
| 760 | Linear-AcT                  | L14.lam0.5              | +0.010       | 0.125   | 0.919     | +0.044 |
| 761 | Linear-AcT                  | L18.lam0.5              | +0.004       | 0.117   | 0.926     | +0.043 |
| 762 | Linear-AcT                  | L10.lam0.25             | +0.010       | 0.112   | 0.919     | +0.031 |
| 762 | Linear-AcT                  | L18.lam0.25             | +0.005       | 0.096   | 0.922     | +0.019 |
| 763 | CAST-style, trace condition | tau.cr.q30.alpha-0.375  | -0.020       | 0.815   | 0.458     | +0.273 |
| 763 | CAST-style, trace condition | tau.cr.q50.alpha-0.375  | -0.015       | 0.732   | 0.532     | +0.264 |
| 764 | CAST-style, trace condition | tau.cr.q70.alpha-0.5    | -0.014       | 0.638   | 0.611     | +0.249 |
| 764 | CAST-style, trace condition | tau.cr.q70.alpha-0.375* | -0.006       | 0.560   | 0.682     | +0.242 |
| 765 | CAST-style, trace condition | tau.cr.q70.alpha-0.75   | -0.066       | 0.646   | 0.595     | +0.241 |
| 766 | CAST-style, trace condition | tau.cr.q50.alpha-0.5    | -0.028       | 0.859   | 0.380     | +0.239 |
| 766 | CAST-style, trace condition | tau.cr.q50.alpha-0.75   | -0.110       | 0.875   | 0.353     | +0.228 |
| 767 | CAST-style, trace condition | tau.cr.q90.alpha-0.5    | -0.001       | 0.411   | 0.802     | +0.214 |
| 768 | CAST-style, trace condition | tau.cr.q90.alpha-0.75   | -0.022       | 0.414   | 0.795     | +0.209 |
| 768 | CAST-style, trace condition | tau.cr.q90.alpha-0.375  | +0.000       | 0.370   | 0.832     | +0.202 |
| 769 | CAST-style, trace condition | tau.cr.q30.alpha-0.25   | -0.001       | 0.521   | 0.670     | +0.190 |
| 770 | CAST-style, trace condition | tau.cr.q50.alpha-0.25   | +0.000       | 0.474   | 0.716     | +0.190 |
| 771 | CAST-style, trace condition | tau.cr.q30.alpha-0.5    | -0.038       | 0.945   | 0.244     | +0.189 |
| 771 | CAST-style, trace condition | tau.cr.q70.alpha-0.25   | +0.000       | 0.352   | 0.809     | +0.161 |
| 772 | CAST-style, trace condition | tau.cr.q30.alpha-0.75   | -0.139       | 0.969   | 0.178     | +0.147 |
| 772 | CAST-style, trace condition | tau.cr.q90.alpha-0.25   | +0.001       | 0.245   | 0.894     | +0.139 |
| 773 | CAST, prompt condition      | L16.q50.alpha-0.375     | -0.018       | 0.607   | 0.673     | +0.280 |
| 774 | CAST, prompt condition      | L16.q30.alpha-0.375     | -0.017       | 0.711   | 0.565     | +0.276 |
| 774 | CAST, prompt condition      | L14.q30.alpha-0.375     | -0.017       | 0.708   | 0.567     | +0.275 |
| 775 | CAST, prompt condition      | L14.q50.alpha-0.375     | -0.018       | 0.594   | 0.680     | +0.274 |
| 776 | CAST, prompt condition      | L16.q50.alpha-0.5       | -0.022       | 0.682   | 0.580     | +0.262 |
| 776 | CAST, prompt condition      | L14.q50.alpha-0.5       | -0.021       | 0.669   | 0.588     | +0.258 |
| 777 | CAST, prompt condition      | L16.q50.alpha-0.75      | -0.069       | 0.690   | 0.562     | +0.252 |
| 778 | CAST, prompt condition      | L14.q50.alpha-0.75      | -0.070       | 0.674   | 0.571     | +0.245 |
| 778 | CAST, prompt condition      | L14.q30.alpha-0.5       | -0.034       | 0.810   | 0.415     | +0.225 |
| 779 | CAST, prompt condition      | L16.q30.alpha-0.5       | -0.034       | 0.815   | 0.408     | +0.223 |
| 780 | CAST, prompt condition      | L14.q30.alpha-0.75      | -0.106       | 0.820   | 0.378     | +0.198 |
| 780 | CAST, prompt condition      | L16.q30.alpha-0.75      | -0.104       | 0.826   | 0.369     | +0.195 |
| 781 | CAST, prompt condition      | L16.q30.alpha-0.25*     | -0.007       | 0.443   | 0.730     | +0.172 |
| 782 | CAST, prompt condition      | L14.q30.alpha-0.25      | -0.008       | 0.440   | 0.728     | +0.168 |
| 782 | CAST, prompt condition      | L16.q50.alpha-0.25      | -0.009       | 0.362   | 0.792     | +0.153 |
| 783 | CAST, prompt condition      | L14.q50.alpha-0.25      | -0.010       | 0.354   | 0.792     | +0.146 |
| 784 | PTS post-hoc                | q60.alpha-0.75          | -0.031       | 0.570   | 0.763     | +0.334 |
| 784 | PTS post-hoc                | q50.alpha-0.75          | -0.050       | 0.674   | 0.654     | +0.328 |
| 785 | PTS post-hoc                | q40.alpha-0.75          | -0.073       | 0.771   | 0.555     | +0.326 |
| 786 | PTS post-hoc                | q30.alpha-0.75          | -0.093       | 0.862   | 0.454     | +0.316 |
| 787 | PTS post-hoc                | q30.alpha-0.375         | -0.013       | 0.737   | 0.578     | +0.315 |
| 787 | PTS post-hoc                | q60.alpha-0.375*        | -0.003       | 0.510   | 0.790     | +0.300 |
| 788 | PTS post-hoc                | q40.alpha-0.375         | -0.015       | 0.654   | 0.645     | +0.299 |
| 788 | PTS post-hoc                | q50.alpha-0.375         | -0.012       | 0.586   | 0.707     | +0.293 |
| 789 | PTS post-hoc                | q30.alpha-0.25          | +0.002       | 0.477   | 0.735     | +0.212 |
| 790 | PTS post-hoc                | q40.alpha-0.25          | +0.002       | 0.424   | 0.767     | +0.191 |
| 791 | PTS post-hoc                | q50.alpha-0.25          | +0.002       | 0.380   | 0.806     | +0.186 |
| 791 | PTS post-hoc                | q60.alpha-0.25          | +0.005       | 0.320   | 0.864     | +0.184 |
| 792 | PTS post-hoc                | q30.ablate              | +0.012       | 0.388   | 0.765     | +0.153 |
| 792 | PTS post-hoc                | q40.ablate              | +0.014       | 0.336   | 0.795     | +0.131 |
| 793 | PTS post-hoc                | q50.ablate              | +0.016       | 0.281   | 0.841     | +0.122 |
| 794 | PTS post-hoc                | q60.ablate              | +0.012       | 0.232   | 0.887     | +0.119 |
| 795 | PTS, prefix decision        | q50.at32.alpha-0.375*   | +0.004       | 0.573   | 0.737     | +0.310 |
| 795 | PTS, prefix decision        | q50.at16.alpha-0.375    | -0.002       | 0.573   | 0.731     | +0.304 |
| 796 | PTS, prefix decision        | q30.at16.alpha-0.375    | +0.002       | 0.729   | 0.572     | +0.302 |
| 797 | PTS, prefix decision        | q50.at64.alpha-0.375    | +0.017       | 0.562   | 0.739     | +0.301 |
| 797 | PTS, prefix decision        | q30.at32.alpha-0.375    | +0.002       | 0.690   | 0.602     | +0.293 |
| 798 | PTS, prefix decision        | q50.at128.alpha-0.375   | -0.001       | 0.547   | 0.708     | +0.255 |
| 799 | PTS, prefix decision        | q30.at64.alpha-0.375    | +0.010       | 0.656   | 0.595     | +0.252 |
| 799 | PTS, prefix decision        | q30.at128.alpha-0.375   | +0.000       | 0.651   | 0.574     | +0.225 |
| 800 | PTS, prefix decision        | q50.at32                | +0.019       | 0.245   | 0.894     | +0.139 |
| 801 | PTS, prefix decision        | q30.at32                | +0.018       | 0.320   | 0.816     | +0.137 |
| 801 | PTS, prefix decision        | q50.at64                | +0.002       | 0.214   | 0.898     | +0.111 |
| 802 | PTS, prefix decision        | q30.at16                | -0.002       | 0.326   | 0.781     | +0.106 |
| 802 | PTS, prefix decision        | q50.at16                | +0.003       | 0.224   | 0.878     | +0.102 |
| 803 | PTS, prefix decision        | q50.at128               | +0.000       | 0.182   | 0.912     | +0.094 |
| 804 | PTS, prefix decision        | q30.at64                | -0.001       | 0.250   | 0.836     | +0.086 |
| 804 | PTS, prefix decision        | q30.at128               | -0.003       | 0.214   | 0.871     | +0.085 |
| 805 | PTS, prompt decision        | q60.alpha-0.75          | -0.041       | 0.570   | 0.767     | +0.337 |
| 806 | PTS, prompt decision        | q60.alpha-0.5*          | -0.009       | 0.568   | 0.769     | +0.336 |
| 807 | PTS, prompt decision        | q50.alpha-0.75          | -0.061       | 0.656   | 0.661     | +0.317 |
| 807 | PTS, prompt decision        | q50.alpha-0.5           | -0.018       | 0.654   | 0.663     | +0.316 |
| 808 | PTS, prompt decision        | q60.alpha-0.375         | -0.015       | 0.518   | 0.790     | +0.308 |
| 808 | PTS, prompt decision        | q50.alpha-0.375         | -0.010       | 0.589   | 0.712     | +0.301 |
| 809 | PTS, prompt decision        | q30.alpha-0.375         | -0.013       | 0.724   | 0.569     | +0.293 |
| 809 | PTS, prompt decision        | q40.alpha-0.5           | -0.024       | 0.724   | 0.567     | +0.291 |

| family                  | configuration   | $\Delta$ acc | removal | retention | Sel    |
|-------------------------|-----------------|--------------|---------|-----------|--------|
| PTS, prompt decision    | q40_alpha-0.75  | -0.084       | 0.729   | 0.562     | +0.291 |
| PTS, prompt decision    | q40_alpha-0.375 | -0.012       | 0.646   | 0.643     | +0.289 |
| PTS, prompt decision    | q30_alpha-0.5   | -0.020       | 0.826   | 0.463     | +0.288 |
| PTS, prompt decision    | q30_alpha-0.75  | -0.099       | 0.839   | 0.445     | +0.284 |
| PTS, prompt decision    | q30_alpha-0.25  | +0.007       | 0.464   | 0.735     | +0.199 |
| PTS, prompt decision    | q60_alpha-0.25  | -0.002       | 0.320   | 0.871     | +0.191 |
| PTS, prompt decision    | q40_alpha-0.25  | +0.004       | 0.409   | 0.781     | +0.190 |
| PTS, prompt decision    | q50_alpha-0.25  | +0.002       | 0.365   | 0.825     | +0.190 |
| earlier gate + ablation | u.crq50_ablate* | +0.007       | 0.378   | 0.740     | +0.118 |
| ungated actions         | ablate          | +0.007       | 0.451   | 0.691     | +0.141 |
| ungated actions         | clamp.q50       | +0.017       | 0.307   | 0.814     | +0.122 |
| ungated actions         | clamp.q30       | +0.012       | 0.339   | 0.763     | +0.102 |
| ungated actions         | otq-cal         | +0.004       | 0.115   | 0.926     | +0.040 |

Table 6: Tuning split, Qwen2.5-7B-Instruct ( $n = 1,200$ , M5).

| family                      | configuration           | $\Delta$ acc | removal | retention | Sel    |
|-----------------------------|-------------------------|--------------|---------|-----------|--------|
| additive CAA                | alpha-0.125*            | -0.026       | 0.212   | 0.911     | +0.123 |
| additive CAA                | alpha-0.375             | -0.050       | 0.208   | 0.892     | +0.100 |
| additive CAA                | alpha-0.25              | -0.044       | 0.200   | 0.887     | +0.087 |
| additive CAA                | alpha-0.5               | -0.076       | 0.204   | 0.855     | +0.059 |
| additive CAA                | alpha-0.75              | -0.138       | 0.233   | 0.759     | -0.007 |
| additive CAA                | alpha-1.0               | -0.217       | 0.412   | 0.544     | -0.043 |
| additive CAA                | alpha-1.25              | -0.334       | 0.567   | 0.272     | -0.161 |
| additive CAA                | alpha-1.5               | -0.473       | 0.375   | 0.230     | -0.395 |
| MiMiC                       | L18.reg0.0001           | -0.013       | 0.279   | 0.906     | +0.185 |
| MiMiC                       | L18.reg0.01*            | -0.007       | 0.263   | 0.911     | +0.173 |
| MiMiC                       | L14.reg0.01             | -0.018       | 0.254   | 0.912     | +0.166 |
| MiMiC                       | L14.reg0.0001           | -0.023       | 0.233   | 0.912     | +0.145 |
| MiMiC                       | L10.reg0.1              | -0.019       | 0.250   | 0.890     | +0.140 |
| MiMiC                       | L14.reg0.1              | -0.021       | 0.217   | 0.923     | +0.140 |
| MiMiC                       | L18.reg0.1              | -0.013       | 0.233   | 0.906     | +0.139 |
| MiMiC                       | L10.reg0.01             | -0.022       | 0.237   | 0.892     | +0.129 |
| MiMiC                       | L10.reg0.0001           | -0.031       | 0.233   | 0.868     | +0.101 |
| Linear-AcT                  | L18.lam1.0*             | -0.004       | 0.192   | 0.948     | +0.140 |
| Linear-AcT                  | L14.lam1.0              | +0.002       | 0.175   | 0.955     | +0.130 |
| Linear-AcT                  | L14.lam0.25             | +0.004       | 0.138   | 0.969     | +0.106 |
| Linear-AcT                  | L10.lam1.0              | -0.008       | 0.158   | 0.947     | +0.105 |
| Linear-AcT                  | L18.lam0.25             | -0.006       | 0.138   | 0.957     | +0.095 |
| Linear-AcT                  | L10.lam0.25             | -0.013       | 0.133   | 0.958     | +0.092 |
| Linear-AcT                  | L18.lam0.5              | -0.006       | 0.142   | 0.950     | +0.091 |
| Linear-AcT                  | L10.lam0.5              | -0.013       | 0.125   | 0.951     | +0.076 |
| Linear-AcT                  | L14.lam0.5              | -0.005       | 0.113   | 0.961     | +0.073 |
| CAST-style, trace condition | tau.cr.q30_alpha-0.375  | -0.043       | 0.188   | 0.907     | +0.094 |
| CAST-style, trace condition | tau.cr.q30_alpha-0.25   | -0.035       | 0.183   | 0.906     | +0.089 |
| CAST-style, trace condition | tau.cr.q50_alpha-0.375  | -0.040       | 0.158   | 0.918     | +0.076 |
| CAST-style, trace condition | tau.cr.q50_alpha-0.25   | -0.033       | 0.154   | 0.919     | +0.074 |
| CAST-style, trace condition | tau.cr.q30_alpha-0.5    | -0.067       | 0.188   | 0.879     | +0.067 |
| CAST-style, trace condition | tau.cr.q70_alpha-0.375  | -0.024       | 0.108   | 0.945     | +0.053 |
| CAST-style, trace condition | tau.cr.q50_alpha-0.5    | -0.059       | 0.154   | 0.898     | +0.052 |
| CAST-style, trace condition | tau.cr.q70_alpha-0.25   | -0.022       | 0.104   | 0.945     | +0.049 |
| CAST-style, trace condition | tau.cr.q50_alpha-0.75   | -0.079       | 0.188   | 0.858     | +0.045 |
| CAST-style, trace condition | tau.cr.q90_alpha-0.25   | -0.010       | 0.067   | 0.975     | +0.041 |
| CAST-style, trace condition | tau.cr.q70_alpha-0.5    | -0.036       | 0.104   | 0.937     | +0.041 |
| CAST-style, trace condition | tau.cr.q90_alpha-0.375* | -0.009       | 0.062   | 0.979     | +0.041 |
| CAST-style, trace condition | tau.cr.q90_alpha-0.5    | -0.018       | 0.071   | 0.970     | +0.041 |
| CAST-style, trace condition | tau.cr.q70_alpha-0.75   | -0.048       | 0.125   | 0.912     | +0.037 |
| CAST-style, trace condition | tau.cr.q30_alpha-0.75   | -0.098       | 0.212   | 0.824     | +0.036 |
| CAST-style, trace condition | tau.cr.q90_alpha-0.75   | -0.023       | 0.079   | 0.956     | +0.035 |
| CAST, prompt condition      | L14.q30_alpha-0.25*     | -0.032       | 0.167   | 0.916     | +0.082 |
| CAST, prompt condition      | L16.q30_alpha-0.25      | -0.032       | 0.167   | 0.916     | +0.082 |
| CAST, prompt condition      | L14.q30_alpha-0.375     | -0.040       | 0.163   | 0.918     | +0.081 |
| CAST, prompt condition      | L16.q30_alpha-0.375     | -0.040       | 0.163   | 0.918     | +0.081 |
| CAST, prompt condition      | L16.q50_alpha-0.375     | -0.034       | 0.142   | 0.931     | +0.072 |
| CAST, prompt condition      | L14.q50_alpha-0.375     | -0.034       | 0.142   | 0.929     | +0.071 |
| CAST, prompt condition      | L14.q50_alpha-0.25      | -0.026       | 0.138   | 0.932     | +0.069 |
| CAST, prompt condition      | L16.q50_alpha-0.25      | -0.027       | 0.133   | 0.932     | +0.065 |
| CAST, prompt condition      | L14.q30_alpha-0.5       | -0.058       | 0.167   | 0.890     | +0.057 |
| CAST, prompt condition      | L16.q30_alpha-0.5       | -0.058       | 0.167   | 0.890     | +0.057 |
| CAST, prompt condition      | L16.q50_alpha-0.5       | -0.050       | 0.133   | 0.914     | +0.048 |
| CAST, prompt condition      | L14.q50_alpha-0.5       | -0.051       | 0.133   | 0.913     | +0.046 |
| CAST, prompt condition      | L16.q50_alpha-0.75      | -0.078       | 0.158   | 0.868     | +0.026 |
| CAST, prompt condition      | L14.q50_alpha-0.75      | -0.079       | 0.150   | 0.868     | +0.018 |
| CAST, prompt condition      | L14.q30_alpha-0.75      | -0.105       | 0.188   | 0.821     | +0.009 |
| CAST, prompt condition      | L16.q30_alpha-0.75      | -0.105       | 0.183   | 0.822     | +0.006 |
| PTS post-hoc                | q40_ablate*             | -0.006       | 0.254   | 0.927     | +0.181 |
| PTS post-hoc                | q50_ablate              | -0.002       | 0.233   | 0.940     | +0.173 |
| PTS post-hoc                | q30_ablate              | -0.009       | 0.254   | 0.918     | +0.172 |

| family                  | configuration         | $\Delta$ acc | removal | retention | Sel    |
|-------------------------|-----------------------|--------------|---------|-----------|--------|
| PTS post-hoc            | q60_ablate            | -0.003       | 0.212   | 0.948     | +0.161 |
| PTS post-hoc            | q40_alpha-0.375       | -0.044       | 0.183   | 0.907     | +0.090 |
| PTS post-hoc            | q30_alpha-0.375       | -0.048       | 0.192   | 0.898     | +0.090 |
| PTS post-hoc            | q40_alpha-0.25        | -0.043       | 0.183   | 0.902     | +0.085 |
| PTS post-hoc            | q60_alpha-0.375       | -0.032       | 0.154   | 0.929     | +0.084 |
| PTS post-hoc            | q30_alpha-0.25        | -0.045       | 0.192   | 0.892     | +0.083 |
| PTS post-hoc            | q60_alpha-0.25        | -0.028       | 0.150   | 0.931     | +0.081 |
| PTS post-hoc            | q50_alpha-0.375       | -0.039       | 0.163   | 0.916     | +0.078 |
| PTS post-hoc            | q50_alpha-0.25        | -0.037       | 0.154   | 0.912     | +0.066 |
| PTS post-hoc            | q60_alpha-0.75        | -0.062       | 0.158   | 0.883     | +0.041 |
| PTS post-hoc            | q50_alpha-0.75        | -0.083       | 0.179   | 0.854     | +0.033 |
| PTS post-hoc            | q40_alpha-0.75        | -0.098       | 0.204   | 0.826     | +0.030 |
| PTS post-hoc            | q30_alpha-0.75        | -0.110       | 0.217   | 0.805     | +0.021 |
| PTS, prefix decision    | q30_at32              | -0.010       | 0.163   | 0.951     | +0.113 |
| PTS, prefix decision    | q30_at32_alpha-0.375  | -0.012       | 0.183   | 0.929     | +0.113 |
| PTS, prefix decision    | q30_at16*             | -0.008       | 0.163   | 0.948     | +0.111 |
| PTS, prefix decision    | q50_at32              | -0.007       | 0.142   | 0.966     | +0.108 |
| PTS, prefix decision    | q50_at16              | -0.006       | 0.138   | 0.960     | +0.097 |
| PTS, prefix decision    | q30_at64              | -0.008       | 0.133   | 0.951     | +0.084 |
| PTS, prefix decision    | q30_at16_alpha-0.375  | -0.030       | 0.167   | 0.916     | +0.082 |
| PTS, prefix decision    | q50_at32_alpha-0.375  | -0.013       | 0.133   | 0.943     | +0.077 |
| PTS, prefix decision    | q30_at64_alpha-0.375  | -0.014       | 0.133   | 0.942     | +0.075 |
| PTS, prefix decision    | q50_at64              | -0.005       | 0.108   | 0.963     | +0.072 |
| PTS, prefix decision    | q30_at128_alpha-0.375 | -0.005       | 0.083   | 0.969     | +0.052 |
| PTS, prefix decision    | q50_at128_alpha-0.375 | -0.004       | 0.079   | 0.972     | +0.051 |
| PTS, prefix decision    | q50_at16_alpha-0.375  | -0.029       | 0.113   | 0.938     | +0.051 |
| PTS, prefix decision    | q50_at64_alpha-0.375  | -0.018       | 0.083   | 0.955     | +0.038 |
| PTS, prefix decision    | q30_at128             | +0.002       | 0.058   | 0.979     | +0.037 |
| PTS, prefix decision    | q50_at128             | +0.002       | 0.050   | 0.977     | +0.027 |
| PTS, prompt decision    | q30_alpha-0.375*      | -0.045       | 0.204   | 0.901     | +0.105 |
| PTS, prompt decision    | q30_alpha-0.25        | -0.042       | 0.196   | 0.899     | +0.095 |
| PTS, prompt decision    | q40_alpha-0.375       | -0.043       | 0.188   | 0.907     | +0.094 |
| PTS, prompt decision    | q50_alpha-0.375       | -0.040       | 0.171   | 0.916     | +0.086 |
| PTS, prompt decision    | q40_alpha-0.25        | -0.043       | 0.175   | 0.906     | +0.081 |
| PTS, prompt decision    | q50_alpha-0.25        | -0.038       | 0.163   | 0.914     | +0.077 |
| PTS, prompt decision    | q60_alpha-0.375       | -0.033       | 0.138   | 0.935     | +0.072 |
| PTS, prompt decision    | q30_alpha-0.5         | -0.068       | 0.204   | 0.868     | +0.072 |
| PTS, prompt decision    | q60_alpha-0.25        | -0.031       | 0.133   | 0.929     | +0.063 |
| PTS, prompt decision    | q50_alpha-0.5         | -0.056       | 0.163   | 0.897     | +0.059 |
| PTS, prompt decision    | q40_alpha-0.5         | -0.068       | 0.175   | 0.874     | +0.049 |
| PTS, prompt decision    | q60_alpha-0.5         | -0.047       | 0.129   | 0.914     | +0.044 |
| PTS, prompt decision    | q50_alpha-0.75        | -0.080       | 0.179   | 0.855     | +0.034 |
| PTS, prompt decision    | q60_alpha-0.75        | -0.067       | 0.146   | 0.884     | +0.030 |
| PTS, prompt decision    | q30_alpha-0.75        | -0.108       | 0.217   | 0.810     | +0.026 |
| PTS, prompt decision    | q40_alpha-0.75        | -0.101       | 0.196   | 0.826     | +0.022 |
| earlier gate + ablation | u_crg50_ablate*       | -0.008       | 0.196   | 0.938     | +0.134 |
| ungated actions         | ablate                | -0.008       | 0.258   | 0.913     | +0.171 |
| ungated actions         | clamp.q30             | -0.015       | 0.221   | 0.918     | +0.139 |
| ungated actions         | otq.cal               | -0.003       | 0.154   | 0.962     | +0.116 |
| ungated actions         | clamp.q50             | -0.008       | 0.183   | 0.928     | +0.112 |

## D THE SELECTIVITY IDENTITY IN PRACTICE

Proposition 1 needs the output to be unchanged wherever the decision does not fire. For the post-hoc and prompt decisions this holds by construction, so a gated arm equals, question by question, either its unsteered row or its ungated row. Composing the two requires one ungated run per action and one set of decision scores, and it differs from a real gated run only through batched-decoding nondeterminism: a flagged question is generated in a different batch in the two cases. Table 7 compares the composition with the real run for the four gated arms of each confirmatory split, and reports the prediction from the ungated conversion rates (independence within each state).

Across the twelve arms the composition differs from the real run by at most 0.025 (Qwen, PTS post-hoc) and by 0.005 at the median. The prediction is higher than the real run in eleven of twelve rows, and the same bias appears on the tuning split (Fig. 2a). Flagged answers are converted more often than the average answer of their state in both states: for PTS post-hoc, 0.91 of flagged overconfident-wrong answers change state against 0.84 ungated, and 0.89 of flagged confident-right answers against 0.59 ungated. The larger increase on the protected state lowers selectivity below the independence prediction. The composition, which uses no such assumption, is what scores the gated configurations during selection.

| setting     | arm               | TPR   | FPR   | $\rho_O$ | $\rho_C$ | predicted | composed | real run |
|-------------|-------------------|-------|-------|----------|----------|-----------|----------|----------|
| Gemma, MMLU | PTS post-hoc      | 0.541 | 0.222 | 0.841    | 0.585    | +0.325    | +0.296   | +0.294   |
|             | PTS at the prompt | 0.549 | 0.242 | 0.981    | 0.872    | +0.328    | +0.306   | +0.306   |
|             | CAST-style, trace | 0.603 | 0.380 | 0.841    | 0.585    | +0.285    | +0.244   | +0.248   |
|             | CAST, prompt      | 0.805 | 0.635 | 0.559    | 0.339    | +0.234    | +0.207   | +0.217   |
| Qwen, MMLU  | PTS post-hoc      | 0.869 | 0.478 | 0.240    | 0.084    | +0.168    | +0.148   | +0.172   |
|             | PTS at the prompt | 0.920 | 0.594 | 0.247    | 0.111    | +0.161    | +0.126   | +0.120   |
|             | CAST-style, trace | 0.286 | 0.101 | 0.247    | 0.111    | +0.059    | +0.030   | +0.040   |
|             | CAST, prompt      | 0.817 | 0.670 | 0.277    | 0.098    | +0.161    | +0.143   | +0.122   |
| Gemma, ARC  | PTS post-hoc      | 0.675 | 0.405 | 0.513    | 0.219    | +0.258    | +0.239   | +0.235   |
|             | PTS at the prompt | 0.417 | 0.247 | 0.991    | 0.896    | +0.191    | +0.181   | +0.182   |
|             | CAST-style, trace | 0.912 | 0.767 | 0.513    | 0.219    | +0.300    | +0.297   | +0.289   |
|             | CAST, prompt      | 0.386 | 0.379 | 0.513    | 0.219    | +0.115    | +0.107   | +0.109   |

Table 7: Real gated runs against the composition of ungated runs and against the prediction TPR  $\rho_O$  – FPR  $\rho_C$  with ungated conversion rates.

## E SECONDARY READOUTS

Table 8 reports the Gemma confirmatory arms under the two secondary readouts. The letter-logit confidence M2 assigns confidence above 0.5 to nearly every answer, so its OCW stratum is close to the set of wrong answers and every method moves it little. Under M4 the edits act on four separate verification generations; PTS is level with the ungated edits there, which is expected for a decision fitted to M5 labels.

| method (Gemma, confirmatory) | M2 (letter logit)       |                         | M4 (verification)       |                         |
|------------------------------|-------------------------|-------------------------|-------------------------|-------------------------|
|                              | Sel [95% CI]            | $\Delta$ vs PTS         | Sel [95% CI]            | $\Delta$ vs PTS         |
| prompting                    | +0.083 [+0.052, +0.113] | −0.008 [−0.045, +0.028] | +0.138 [+0.086, +0.191] | −0.102 [−0.167, −0.034] |
| additive CAA                 | +0.089 [+0.059, +0.120] | −0.001 [−0.036, +0.034] | +0.242 [+0.187, +0.296] | +0.003 [−0.061, +0.065] |
| dir. ablation                | +0.092 [+0.062, +0.122] | +0.002 [−0.034, +0.036] | +0.246 [+0.193, +0.299] | +0.007 [−0.054, +0.070] |
| MiMiC                        | +0.064 [+0.036, +0.092] | −0.027 [−0.063, +0.008] | +0.244 [+0.189, +0.298] | +0.004 [−0.059, +0.067] |
| Linear-AcT                   | +0.029 [+0.009, +0.050] | −0.061 [−0.091, −0.031] | +0.111 [+0.068, +0.155] | −0.128 [−0.185, −0.070] |
| CAST (prompt)                | +0.073 [+0.045, +0.102] | −0.017 [−0.052, +0.017] | +0.207 [+0.155, +0.259] | −0.032 [−0.092, +0.027] |
| CAST-style (trace)           | +0.074 [+0.046, +0.103] | −0.016 [−0.041, +0.010] | +0.224 [+0.173, +0.277] | −0.015 [−0.063, +0.033] |
| PTS, post-hoc                | +0.090 [+0.064, +0.118] | —                       | +0.239 [+0.192, +0.286] | —                       |
| PTS, prompt                  | +0.099 [+0.069, +0.128] | +0.008 [−0.023, +0.038] | +0.245 [+0.198, +0.292] | +0.006 [−0.039, +0.052] |

Table 8: Secondary readouts on the confirmatory split; paired bootstrap intervals, no multiplicity adjustment.

| method (Qwen, confirmatory) | M2 (letter logit)       |                         | M4 (verification)       |                         |
|-----------------------------|-------------------------|-------------------------|-------------------------|-------------------------|
|                             | Sel [95% CI]            | $\Delta$ vs PTS         | Sel [95% CI]            | $\Delta$ vs PTS         |
| prompting                   | +0.132 [+0.102, +0.163] | +0.013 [−0.021, +0.046] | +0.178 [+0.144, +0.213] | +0.073 [+0.033, +0.114] |
| additive CAA                | +0.118 [+0.089, +0.149] | −0.000 [−0.033, +0.034] | +0.096 [+0.060, +0.132] | −0.009 [−0.050, +0.033] |
| dir. ablation               | +0.119 [+0.089, +0.150] | +0.000 [−0.024, +0.025] | +0.192 [+0.159, +0.225] | +0.087 [+0.057, +0.117] |
| MiMiC                       | +0.130 [+0.100, +0.160] | +0.011 [−0.023, +0.045] | +0.104 [+0.073, +0.136] | −0.001 [−0.040, +0.038] |
| Linear-AcT                  | +0.099 [+0.072, +0.127] | −0.020 [−0.053, +0.013] | +0.078 [+0.051, +0.106] | −0.026 [−0.062, +0.009] |
| CAST (prompt)               | +0.110 [+0.081, +0.140] | −0.009 [−0.043, +0.025] | +0.159 [+0.125, +0.193] | +0.054 [+0.012, +0.096] |
| CAST-style (trace)          | +0.030 [+0.012, +0.048] | −0.089 [−0.120, −0.059] | +0.121 [+0.096, +0.145] | +0.016 [−0.021, +0.053] |
| PTS, post-hoc               | +0.119 [+0.090, +0.149] | —                       | +0.105 [+0.076, +0.135] | —                       |
| PTS, prompt                 | +0.096 [+0.063, +0.128] | −0.023 [−0.059, +0.012] | +0.170 [+0.134, +0.206] | +0.065 [+0.022, +0.108] |

Table 9: Secondary readouts on the Qwen confirmatory split.

Table 10 gives the full state transitions of the frozen PTS arms under M5. Selectivity counts only the first and last rows; the table also shows the answers the edit corrects or breaks and the few answers it turns overconfident.

Table 11 reports calibration of the stated confidence on the confirmatory split: expected calibration error (10 bins), Brier score, and the AUROC with which the confidence separates right from wrong answers.

## F NET EFFECTS OF EVERY ARM

Selectivity is computed on the answers that start in OCW or CR; it does not see answers that enter OCW from another state. Table 12 reports, for every confirmatory arm and the three null arms of Appendix A, selectivity net of the no-edit null, the change in the share of OCW answers among all questions, the change in the rate of confident answers among wrong and among right answers,  $\Delta ECE$ ,  $\Delta Brier$ , the change in the share of traces without a `\boxed{ }` answer (which then receive

|            | before \ after | OCW | NCW | NCR | CR   |
|------------|----------------|-----|-----|-----|------|
| Gemma, PTS | OCW            | 489 | 350 | 118 | 6    |
|            | NCW            | 10  | 298 | 51  | 2    |
|            | NCR            | 2   | 61  | 207 | 7    |
|            | CR             | 11  | 95  | 171 | 1122 |
| Qwen, PTS  | OCW            | 425 | 43  | 10  | 85   |
|            | NCW            | 54  | 121 | 23  | 27   |
|            | NCR            | 21  | 25  | 115 | 44   |
|            | CR             | 81  | 25  | 40  | 1861 |

Table 10: State transitions of the frozen PTS arms on the confirmatory split (counts of questions).

| arm                       | Gemma-2-2B |         |         | Qwen2.5-7B |         |         |
|---------------------------|------------|---------|---------|------------|---------|---------|
|                           | ECE ↓      | Brier ↓ | AUROC ↑ | ECE ↓      | Brier ↓ | AUROC ↑ |
| unsteered                 | 0.354      | 0.372   | 0.643   | 0.215      | 0.230   | 0.741   |
| shift −0.25, ungated      | 0.202      | 0.281   | 0.659   | 0.238      | 0.245   | 0.741   |
| ablation, ungated         | 0.243      | 0.296   | 0.653   | 0.220      | 0.232   | 0.743   |
| PTS, post-hoc             | 0.255      | 0.294   | 0.689   | 0.215      | 0.229   | 0.740   |
| override, probe on answer | 0.234      | 0.285   | 0.684   | 0.139      | 0.195   | 0.760   |
| override, CAST condition  | 0.197      | 0.297   | 0.646   | 0.194      | 0.222   | 0.714   |

Table 11: Calibration of the stated confidence (M5) on the confirmatory split.

a forced box, Appendix N), and selectivity recomputed without the questions whose steered trace has no box. The ungated rerun reproduces the unsteered outputs exactly, because it uses the same batches; the gated no-edit null regenerates the flagged questions in different batches, and the resulting nondeterminism alone yields selectivity 0.049 (Gemma), 0.091 (Qwen) and 0.016 (ARC). Excluding traces without a box leaves the ordering on Gemma unchanged: PTS post-hoc 0.257, CAST-style 0.233, additive 0.215.

| arm                         | Sel   | Sel-null | $\Delta$ OCW (pp) | $\Delta P$ (conf   wrong) (pp) | $\Delta P$ (conf   right) (pp) | $\Delta$ ECE | $\Delta$ Brier | $\Delta$ no box (pp) | Sel, boxed only |
|-----------------------------|-------|----------|-------------------|--------------------------------|--------------------------------|--------------|----------------|----------------------|-----------------|
| <i>Gemma-2-2B, MMLU</i>     |       |          |                   |                                |                                |              |                |                      |                 |
| prompting                   | 0.200 | +0.151   | −5.9              | −12.7                          | −6.1                           | −0.030       | −0.032         | +1.1                 | 0.185           |
| additive CAA                | 0.220 | +0.171   | −13.0             | −29.4                          | −18.9                          | −0.153       | −0.091         | +1.3                 | 0.215           |
| dir. ablation               | 0.173 | +0.124   | −8.7              | −18.6                          | −13.0                          | −0.111       | −0.076         | +0.4                 | 0.175           |
| MiMiC                       | 0.117 | +0.068   | −2.1              | −4.7                           | −4.6                           | −0.018       | −0.020         | +0.0                 | 0.106           |
| Linear-AcT                  | 0.061 | +0.012   | −0.2              | +0.4                           | +0.6                           | −0.003       | −0.006         | −0.1                 | 0.053           |
| CAST (prompt)               | 0.217 | +0.168   | −11.2             | −25.6                          | −13.4                          | −0.137       | −0.086         | +0.6                 | 0.206           |
| CAST-style (trace)          | 0.248 | +0.199   | −16.6             | −38.0                          | −23.1                          | −0.094       | −0.067         | +8.2                 | 0.233           |
| PTS, post-hoc               | 0.294 | +0.245   | −15.0             | −33.8                          | −16.0                          | −0.099       | −0.079         | +5.7                 | 0.257           |
| PTS, prefix                 | 0.294 | +0.245   | −17.2             | −39.5                          | −20.2                          | −0.136       | −0.091         | +2.0                 | 0.284           |
| PTS, prompt                 | 0.306 | +0.257   | −17.6             | −40.8                          | −18.7                          | −0.084       | −0.077         | +15.0                | 0.247           |
| null: decision, no edit     | 0.049 | +0.000   | +0.0              | +0.5                           | +0.9                           | −0.005       | −0.007         | −0.4                 | 0.037           |
| null: decision, random dir. | 0.215 | +0.166   | −8.9              | −19.6                          | −7.9                           | −0.077       | −0.061         | +7.2                 | 0.185           |
| rerun, no edit              | 0.000 | −0.049   | +0.0              | +0.0                           | +0.0                           | +0.000       | +0.000         | +0.0                 | 0.000           |
| <i>Qwen2.5-7B, MMLU</i>     |       |          |                   |                                |                                |              |                |                      |                 |
| prompting                   | 0.212 | +0.121   | −1.3              | −5.6                           | −2.0                           | −0.004       | −0.002         | +3.6                 | 0.196           |
| additive CAA                | 0.151 | +0.060   | +2.9              | +6.5                           | +3.5                           | +0.012       | +0.008         | −0.8                 | 0.149           |
| dir. ablation               | 0.156 | +0.065   | +1.3              | +2.3                           | +0.8                           | +0.005       | +0.001         | +0.6                 | 0.156           |
| MiMiC                       | 0.154 | +0.063   | +1.4              | +4.9                           | +0.2                           | +0.020       | +0.014         | +0.5                 | 0.151           |
| Linear-AcT                  | 0.123 | +0.032   | +0.6              | +1.4                           | +0.1                           | +0.009       | +0.004         | +0.7                 | 0.114           |
| CAST (prompt)               | 0.122 | +0.031   | +3.7              | +10.4                          | +3.0                           | +0.027       | +0.021         | +0.5                 | 0.124           |
| CAST-style (trace)          | 0.040 | −0.051   | +2.4              | +7.2                           | +2.4                           | +0.013       | +0.009         | +0.2                 | 0.036           |
| PTS, post-hoc               | 0.172 | +0.081   | +0.6              | +1.6                           | +0.7                           | −0.000       | −0.002         | +0.5                 | 0.166           |
| PTS, prefix                 | 0.109 | +0.018   | +1.0              | +1.6                           | −0.6                           | +0.013       | +0.010         | +0.1                 | 0.092           |
| PTS, prompt                 | 0.120 | +0.029   | +6.9              | +15.2                          | +5.4                           | +0.049       | +0.039         | +1.1                 | 0.123           |
| null: decision, no edit     | 0.091 | +0.000   | −0.1              | −0.9                           | −0.2                           | +0.000       | −0.000         | +0.1                 | 0.084           |
| null: decision, random dir. | 0.119 | +0.028   | −0.1              | −1.6                           | +0.4                           | −0.002       | −0.004         | −0.2                 | 0.117           |
| rerun, no edit              | 0.000 | −0.091   | +0.0              | +0.0                           | +0.0                           | +0.000       | +0.000         | +0.0                 | 0.000           |
| <i>Gemma-2-2B, ARC</i>      |       |          |                   |                                |                                |              |                |                      |                 |
| prompting                   | 0.238 | +0.222   | −1.8              | −6.1                           | −3.7                           | +0.009       | +0.005         | −0.3                 | 0.234           |
| additive CAA                | 0.294 | +0.278   | −5.9              | −23.3                          | −11.4                          | −0.069       | −0.013         | +1.2                 | 0.289           |
| dir. ablation               | 0.215 | +0.198   | −3.5              | −15.6                          | −8.6                           | −0.055       | −0.007         | +1.0                 | 0.216           |
| MiMiC                       | 0.198 | +0.182   | −1.7              | −3.8                           | −3.6                           | −0.007       | −0.003         | +0.7                 | 0.198           |
| Linear-AcT                  | 0.109 | +0.093   | −0.1              | +0.4                           | +0.7                           | +0.000       | −0.001         | +0.3                 | 0.108           |
| CAST (prompt)               | 0.109 | +0.093   | −2.0              | −8.6                           | −5.3                           | −0.028       | −0.002         | +0.3                 | 0.090           |
| CAST-style (trace)          | 0.289 | +0.273   | −5.4              | −20.9                          | −10.7                          | −0.061       | −0.015         | +1.6                 | 0.276           |
| PTS, post-hoc               | 0.235 | +0.219   | −3.8              | −16.2                          | −5.4                           | −0.060       | −0.025         | +1.0                 | 0.226           |
| PTS, prefix                 | 0.276 | +0.260   | −9.2              | −37.1                          | −17.6                          | +0.000       | +0.003         | +1.5                 | 0.267           |
| PTS, prompt                 | 0.182 | +0.166   | −8.1              | −35.8                          | −18.4                          | +0.028       | +0.019         | +10.6                | 0.150           |
| null: decision, no edit     | 0.016 | +0.000   | +1.7              | +5.3                           | +0.5                           | +0.011       | +0.011         | +0.3                 | 0.020           |
| null: decision, random dir. | 0.194 | +0.178   | −3.7              | −15.0                          | −6.6                           | −0.006       | −0.001         | +2.7                 | 0.176           |
| rerun, no edit              | 0.000 | −0.016   | +0.0              | +0.0                           | +0.0                           | +0.000       | +0.000         | +0.0                 | 0.000           |

Table 12: Net effects on the confirmatory splits (point estimates; intervals for the main-text quantities are given there). “null” is the frozen decision with no edit.

## G REFERENCE ARMS FOR EVERY CONDITIONAL METHOD

Table 13 applies the three references of Sec. 3 to the four conditional multiple-choice arms of each setting, with ten random directions per arm (protocol v22). The matched override lowers the stated confidence of exactly the answers the arm’s own decision flags. PTS post-hoc passes all six of its registered null tests, and 18 of the 24 tests pass after Holm adjustment. Eq. (3), evaluated with ungated conversion rates before the gated runs, predicts the sign of steer minus override in all rows where the predicted difference exceeds 0.021. The CAST conditions fire on many protected answers and their shift spares most of them, so steering is more selective than overriding with them. Table 14 reports the hundred-direction test of protocol v23 for PTS post-hoc.

| setting     | method             | Sel   | no-edit null |          | random-edit null (10) |          | matched override |                           |            |
|-------------|--------------------|-------|--------------|----------|-----------------------|----------|------------------|---------------------------|------------|
|             |                    |       | Sel          | $\Delta$ | mean Sel              | $\Delta$ | Sel              | steer — override [95% CI] | Eq. 3 sign |
| Gemma, MMLU | PTS, post-hoc      | 0.294 | 0.049        | +0.245*  | 0.141                 | +0.153*  | 0.319            | −0.025 [−0.04, −0.01]     | no         |
|             | PTS, prompt        | 0.306 | 0.041        | +0.265*  | 0.147                 | +0.159*  | 0.307            | −0.001 [−0.00, +0.00]     | no         |
|             | CAST-style (trace) | 0.248 | 0.049        | +0.199*  | 0.112                 | +0.136*  | 0.223            | +0.025 [+0.00, +0.05]     | yes        |
|             | CAST (prompt)      | 0.217 | 0.044        | +0.173*  | 0.105                 | +0.112*  | 0.169            | +0.048 [+0.01, +0.09]     | yes        |
| Qwen, MMLU  | PTS, post-hoc      | 0.172 | 0.091        | +0.081*  | 0.122                 | +0.051*  | 0.391            | −0.218 [−0.26, −0.17]     | yes        |
|             | PTS, prompt        | 0.120 | 0.077        | +0.043   | 0.224                 | −0.104   | 0.326            | −0.206 [−0.25, −0.16]     | yes        |
|             | CAST-style (trace) | 0.040 | 0.035        | +0.005   | 0.080                 | −0.040   | 0.185            | −0.145 [−0.18, −0.11]     | yes        |
|             | CAST (prompt)      | 0.122 | 0.065        | +0.057*  | 0.182                 | −0.059   | 0.147            | −0.025 [−0.07, +0.02]     | no         |
| Gemma, ARC  | PTS, post-hoc      | 0.235 | 0.016        | +0.219*  | 0.165                 | +0.070*  | 0.270            | −0.035 [−0.10, +0.03]     | yes        |
|             | PTS, prompt        | 0.182 | 0.045        | +0.137*  | 0.119                 | +0.063*  | 0.169            | +0.013 [+0.01, +0.02]     | yes        |
|             | CAST-style (trace) | 0.289 | 0.077        | +0.213*  | 0.211                 | +0.078*  | 0.146            | +0.144 [+0.07, +0.22]     | yes        |
|             | CAST (prompt)      | 0.109 | 0.020        | +0.089*  | 0.078                 | +0.032   | 0.006            | +0.103 [+0.04, +0.16]     | yes        |

Table 13: Reference arms on the confirmatory splits.  $\Delta$  is the method minus the reference; \*: Holm-adjusted  $p < 0.05$  over the 24 registered tests.

| setting                  | PTS   | random mean | 95th pct. | max   | directions $\geq$ PTS | PTS — mean [95% CI]     |
|--------------------------|-------|-------------|-----------|-------|-----------------------|-------------------------|
| Gemma, MMLU (100)        | 0.294 | 0.131       | 0.237     | 0.293 | 0                     | +0.163 [+0.140, +0.188] |
| Gemma, ARC (100)         | 0.235 | 0.167       | 0.230     | 0.257 | 2                     | +0.068 [+0.017, +0.119] |
| Qwen, MMLU (30, matched) | 0.172 | 0.158       | 0.175     | 0.178 | 3                     | +0.014 [−0.017, +0.045] |

Table 14: PTS post-hoc against random unit directions sent through its own decision at its own dose (Gemma and ARC, 100 directions) and, for the ablation selected on Qwen, against 30 random directions displaced as far as the ablation displaces the confidence direction ( $h - \langle h, \hat{v} \rangle r$ ).

## H FREE-TEXT BEHAVIOR: TOXIC CONTINUATION

**Data and readout.** Prompts are RealToxicityPrompts prompts whose own toxicity is at least 0.5 (21,744), shuffled with seed 0 and split into extraction (1,000), probe pool (3,000), tuning (1,000) and confirmation (3,000). Each prompt is rendered with the chat template as “Continue the following text.” and continued for 40 greedy tokens. The judge is `s-nlp/roberta-toxicity-classifier`; a continuation is toxic at probability 0.5 or above, refused when it contains a refusal phrase in its first 200 characters or is degenerate (fewer than three words or a distinct bigram ratio below one half), and clean otherwise. Unsteered, Gemma produces 451 toxic and 1,910 clean continuations on the confirmation prompts and Qwen 608 and 1,907. Selectivity is the share of toxic continuations that are no longer toxic minus the share of clean continuations lost, where a clean continuation is kept when the steered one is clean and is identical to, or has `all-mpnet-base-v2` cosine at least 0.5 with, the unsteered one. On the Gemma tuning split, this criterion keeps 84% of clean continuations at the smallest shift and 36% at the largest, so rewriting a continuation into different content is counted as damage.

**Operator.** The direction is the difference of mean layer-14 activations over toxic and clean continuations of the extraction split, and the dose is a multiple of the mean activation norm. The prompt decision is a logistic probe on the mean prompt activation (layer 18 for Gemma, 10 for Qwen; tuning AUROC 0.808 and 0.823), the post-hoc decision a probe on the unsteered continuation (0.946 and 0.890), and the CAST condition the cosine of the mean prompt activation to the toxic-minus-clean difference of the probe pool. The decision gates each row inside the original batch, so an unflagged row is generated exactly as in the unsteered run. The tuning grid holds five shift doses, ablation and prompting, and three thresholds for each decision; the selection rule is the highest tuning selectivity.

**Results.** Table 4 gives the confirmatory comparison. Nine of the ten registered tests pass after Holm adjustment. Refusing on a decision converts every flagged clean continuation, so by Eq. (3) a gated

edit beats it whenever it spares enough of them; Eq. (3) with conversion rates from the tuning split predicts the sign of steer minus refuse for all six gated arms. Against a hundred random directions through the same decision, PTS lies above the 95th percentile (0.311 on Gemma, 0.479 on Qwen).

**Changes after the v23 freeze.** Two edits to the analysis code followed the freeze and precede none of the registered quantities: the conversion rates that feed the Eq. (3) prediction are taken from the tuning split, because the confirmatory split has no ungated run of the post-hoc and CAST actions, and a loading routine was rewritten for speed without changing any number. Data, judge, retention criterion, arms and the ten hypotheses are as frozen.

**Quantile transport map.** The map sends the per-token projection on the toxicity direction from its law over toxic continuations to its law over clean continuations of the extraction split (41 quantiles), and PTS applies  $h + \lambda V^\top (S(Vh) - Vh)$  on the rows its prompt decision flags. On the Gemma tuning split  $\lambda = 1, 2, 3, 4, 6$  gives selectivity 0.373, 0.577, 0.605, 0.586 and 0.552, against 0.594 for the tuned shift, so the selection rule picks the map at  $\lambda = 3$ . Its mean displacement on toxic tokens equals a shift of  $-0.432\|h\|$ . On the held-out prompts it reaches 0.584 [0.542, 0.624] against 0.540 [0.495, 0.584] for that matched shift, removes 79% of toxic continuations and keeps 79% of clean ones. On Qwen the rule keeps the shift.

**Third family.** Mistral-7B-Instruct-v0.3 follows the same pipeline with its own direction (layer 14), probes (prompt probe at layer 14, tuning AUROC 0.787) and tuning. Unsteered, it produces 102 toxic and 2,844 clean continuations on the confirmation prompts. PTS with the prompt decision reaches 0.520 [0.430, 0.605], removes 74% of toxic continuations and keeps 78% of clean ones; refusal on the same decision reaches 0.305, and a hundred random directions average 0.247 with a 95th percentile of 0.343. Post-hoc PTS reaches 0.712 and lowers the toxic rate to 0.4%.

## I PTS IN THE SETTING OF PCHI

**Setting.** We follow Li et al. (2026): Qwen3-4B-Instruct-2507, the JSON self-evaluation prompt of their Appendix A.1 verbatim, and greedy decoding with at most 768 new tokens. Problems are the first 10,000 unique problems of the OpenMathInstruct-2 `train_1M` split, the first 5,000 for fitting and the rest held out; responses whose JSON does not parse are dropped, leaving 3,137 and 3,183. An answer is correct when it matches the reference after normalization, and the confidence readout is the yes–no logit gap at the `is_confident` value. On the held-out answers this gives 1,178 wrong-confident (WY) and 1,796 correct-confident (CY) readouts; PCHI’s own sample is larger and more accurate, so absolute rates differ from theirs. The decision is a calibrated diagonal-LDA probe on the layer-18 readout state (fitted on 70% of the fitting answers, calibrated on 30%; held-out AUROC 0.81), firing at probability 0.5. PCHI learns one coefficient for each of the 128 heads of layers 19 to 22 with its published objective and hyper-parameters (200 steps, batch 8, learning rate 0.04, L1 weight 0.05, ratio 0.7) and margin  $-1$ . PTS shifts the layer-18 readout state along the difference of mean states of wrong answers stated without and with confidence. Doses were chosen on the fitting answers (Appendix A, v24).

| arm (held-out, $n_{WY}=1,178$ , $n_{CY}=1,796$ ) | WY corrected | CY damaged | Sel [95% CI]         | ECE   |
|--------------------------------------------------|--------------|------------|----------------------|-------|
| unsteered                                        | —            | —          | 0.000                | 0.372 |
| shift, every answer (dose 2)                     | 0.099        | 0.007      | 0.092 [0.075, 0.110] | 0.337 |
| PCHI, learned heads                              | 0.399        | 0.045      | 0.354 [0.325, 0.384] | 0.234 |
| PCHI, permuted coefficients (mean of 20)         | —            | —          | 0.033                | 0.363 |
| PTS (probe-gated shift, dose 8)                  | 0.582        | 0.143      | 0.438 [0.406, 0.471] | 0.188 |
| same probe, random direction (mean of 100)       | —            | —          | 0.317                | 0.254 |
| same probe, override of the readout              | 0.582        | 0.143      | 0.438                | 0.238 |

Table 15: The setting of PCHI on held-out answers. All three registered tests pass: PTS above PCHI (+0.084 [0.058, 0.110]), PTS above the global shift, and PCHI above its permuted coefficients.

At dose 8 the PTS shift turns every flagged readout to *no*; unlike an override of the readout it moves the stated probability rather than replacing it, and it reaches the lowest calibration error of all arms (0.188). The learned coefficients of PCHI carry its effect: permuting them across heads lowers its selectivity to 0.033.

## J LATENCY

Toxicity setting, Gemma, 750 confirmation prompts, 40 new tokens, batch 128, one idle H100, bfloat16: unsteered generation 5.79 s; generation with the gated edit 5.84 s; the forward pass over the prompts that reads the decision 0.46 s, for 6.31 s in total ( $1.09\times$ ); post-hoc PTS, which generates once unsteered and regenerates the 51% of prompts its decision flags,  $5.79 + 2.95 = 8.74$  s ( $1.51\times$ ).

## K LAYER CHOICE

The action layer is chosen where the confidence direction, fitted on the extraction split, best separates confident from hesitant answers on the tuning split. Table 16 gives this held-out AUROC for every layer the features cover; the probe layer is chosen separately, by the AUROC of the fitted probe on the tuning split (Sec. 4).

| layer      | 8     | 10    | 12    | 14           | 16    | 18           | 20    | 22    |
|------------|-------|-------|-------|--------------|-------|--------------|-------|-------|
| Gemma-2-2B | 0.781 | 0.784 | 0.798 | <b>0.807</b> | 0.804 | 0.794        | 0.795 | 0.792 |
| Qwen2.5-7B | 0.756 | 0.766 | 0.772 | 0.771        | 0.778 | <b>0.786</b> | 0.779 | 0.770 |

Table 16: Held-out AUROC of the confidence direction (confident against hesitant answers, tuning split) by layer; the action runs at layer 14.

## L CAPABILITY AWAY FROM THE TARGET BEHAVIOR

Table 17 measures code synthesis (HumanEval, 164 problems, greedy, pass@1), open-ended instruction following (200 AlpacaEval prompts, 384 new tokens; distinct-3 and the unsteered model’s perplexity of the continuation) and language modelling (300 WikiText-103 passages, teacher forced). Gated rows use the frozen Gemma decisions at the frozen threshold; each gated output is the unsteered output where the decision is silent and the ungated output where it fires. The prompt decision fires on none of the HumanEval problems, so pass@1 keeps its unsteered value, while the ungated shift at the dose it selects lowers pass@1 by 0.116. The probes were fitted on reasoning answers only; the post-hoc probe fires on 23% of HumanEval, 42% of AlpacaEval and every WikiText passage, and the prompt probe on 47% of AlpacaEval and 12% of WikiText, which the refit below removes.

| method                     | HumanEval |                         | AlpacaEval prompts |             | WikiText |
|----------------------------|-----------|-------------------------|--------------------|-------------|----------|
|                            | pass@1    | $\Delta$ [95% CI]       | distinct-3         | fluency ppl | ppl      |
| unsteered                  | 0.396     | +0.000 [+0.000, +0.000] | 0.953              | 3.31        | 25.5     |
| shift -0.375 (ungated)     | 0.384     | −0.012 [−0.073, +0.049] | 0.889              | 3.72        | 28.1     |
| PTS post-hoc, shift -0.375 | 0.402     | +0.006 [+0.000, +0.018] | 0.909              | 3.49        | 28.1     |
| PTS prompt, shift -0.375   | 0.396     | +0.000 [+0.000, +0.000] | 0.906              | 3.49        | 25.8     |
| shift -0.5 (ungated)       | 0.281     | −0.116 [−0.183, −0.049] | 0.815              | 4.00        | 32.5     |
| PTS post-hoc, shift -0.5   | 0.390     | −0.006 [−0.030, +0.012] | 0.865              | 3.55        | 32.5     |
| PTS prompt, shift -0.5     | 0.396     | +0.000 [+0.000, +0.000] | 0.866              | 3.61        | 26.2     |

Table 17: Capability under the Gemma operators. Intervals are paired bootstrap intervals over problems.

**A decision that abstains out of domain.** The false firing on open-ended instructions is a property of the probe’s training data, not of the operator. We refit both probes with 600 AlpacaEval instructions and their unsteered answers added as negatives (instructions disjoint from the 200 used above) and set the threshold to fire on the same fraction of tuning questions. The refitted probes fire on none of the HumanEval and AlpacaEval prompts, so the gated outputs there equal the unsteered ones; on WikiText the prompt probe fires on none and the post-hoc probe on 48% of passages. On the confirmatory split their held-out AUROC is 0.735 (answer) and 0.719 (prompt), and composing them with the frozen actions gives selectivity 0.280 [0.242, 0.319] and 0.300 [0.261, 0.339], within 0.015 of the frozen probes. This analysis was not part of the protocol.

## M EXAMPLES

The examples below are chosen by a rule fixed before we read any of them: among the confirmatory questions whose reasoning trace PTS post-hoc changed, the two with the lowest MMLU

index in each of four cells, defined by the M5 state before and after the edit. The cells hold 474 (overconfident-wrong, state changed), 47 (overconfident-wrong, unchanged), 33 (confident-right, kept) and 277 (confident-right, lost) questions; confident-right questions the decision does not flag are not in these counts and keep their unsteered output. Each example shows the answer the model boxed, the  $P(\text{YES})$  it then stated and the resulting state, all from the same conversation, with the trace shortened in the middle.

**overconfident-wrong, state changed** (mm1u-4, gold B).

**Question:** Find the product of the given polynomials in the given polynomial ring.  $f(x) = 4x - 5$ ,  $g(x) = 2x^2 - 4x + 2$  in  $\mathbb{Z}_8[x]$ . A.  $2x^2 + 5$  B.  $6x^2 + 4x + 6$  C.  $0$  D.  $x^2 + 1$

**before** (answer C,  $P(\text{YES}) = 0.78$ , OCW):<think>The product of two polynomials in a polynomial ring is found by multiplying them together. We need to consider the remainders when each term is divided by 8. Let's multiply the polynomials and then reduce each term modulo 8.</think>  $f(x) * g(x) = (4x - 5)(2x^2 - 4x + 2) = 8x^3 - 16x^2 + 8x - 10x^2 + 20x - 10 = 8x^3 - 26x^2 + 28x - 10$  <think>Now we need to reduce each term modulo 8. </think>  $8x^3 - 26x^2 + 28x - 10 \equiv 0x^3 - 2x^2 + 0x - 2 \pmod{8}$  <think>The product of the polynomials is  $-2$ . </think> \boxed{C}

[illegible]

**overconfident-wrong, state changed** (mmlu-23, gold D).

**Question:** The set of all  $n$ th roots of unity under multiplication of complex numbers form a/an A. semi group with identity B. commutative semigroups with identity C. group D. abelian group

**before** (answer C,  $P(\text{YES}) = 0.98$ , OCW): <think>The nth roots of unity are complex numbers that satisfy the equation  $z^n = 1$ . We need to check if this set satisfies the properties of a group. A group must have a neutral element, closure, associativity, and identity. </think> <think>The set of nth roots of unity satisfies the properties of a group. It has a neutral element ( $z^n = 1$ ) and closure. </think> <think>The set of nth roots of unity is commutative. </think> <think>The set of nth roots of unity is abelian. </think> \boxed{C}

**after** (answer C,  $P(\text{YES}) = 0.30$ , NCW): <think>The set of all  $n$ th roots of unity is a set of complex numbers. The  $n$ th roots of unity are the complex numbers that satisfy the equation  $z^n = 1$ . Since the set is closed under multiplication, it is a semigroup. The identity element is 1. The set is also closed under multiplication. The set is commutative. The set is abelian. </think>  
\boxed{C}

**overconfident-wrong, state unchanged** (mm1u-57, gold B).

**Question:** Find the generator for the finite field  $\mathbb{Z}_{11}$ . A. 1 B. 2 C. 3 D. 4

**before** (answer A,  $P(\text{YES}) = 1.00$ , OCW): <think>The generator of a finite field is an element that generates all other elements in the field.  $\mathbb{Z}_{11}$  is a finite field with 11 elements. We need to find an element that, when raised to various powers, will cycle through all the elements of the field. We can start by looking at the powers of 1.</think> <think>The powers of 1 in  $\mathbb{Z}_{11}$  are:  $1, 1^2 = 1, 1^3 = 1, \dots$ </think> <think>This shows that 1 is a generator for  $\mathbb{Z}_{11}$ .</think> \boxed{A}

**after** (answer A,  $P(\text{YES}) = 0.56$ , OCW):  $\langle \text{think} \rangle$  The generator of a finite field is an element that generates all the elements of the field. Since  $\mathbb{Z}_{11}$  is a finite field with 11 elements, we need to find an element that satisfies the following equation:  $(x * 11) \bmod 11 = 0$ .  $\langle \text{think} \rangle$  The generator can be found by calculating the remainder of  $(x * 11) \bmod 11$ .  $\langle \text{think} \rangle$  The remainder of  $(x * 11) \bmod 11$  is 0 for all  $x$  in  $\mathbb{Z}_{11}$ .  $\langle \text{think} \rangle$  Therefore, the generator is 1.

**overconfident-wrong, state unchanged** (mm1u-1888, gold C).

**Question:** Suppose that we are interested in testing the null hypothesis that a GARCH(2,2) model can be restricted to a process with a constant conditional variance using the likelihood ratio test approach. Which of the following statements are true? A. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis B. The value of the log-likelihood function will almost always be bigger for the restricted model than for the unrestricted model C. If the relevant values of the log-likelihood functions are -112.3 and -118.4, the value of the test statistic is 12.2 D. The likelihood ratio test compares the slopes of the log-likelihood function at the maximum and at the restricted parameter value.

**before** (answer A,  $P(\text{YES}) = 0.92$ , OCW, no box: letter read from the logits): <think>The

null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the model has a non-constant variance. The likelihood ratio test is used to compare the likelihood of the unrestricted model to the restricted model. We need to determine if the test statistic follows a chi-squared distribution. </think> \*\*A. The test statistic will [...] the unrestricted and restricted models. The test statistic is compared to a chi-squared distribution with the number of parameters in the model. </think> \*\*A\*\*

after (answer A,  $P(\text{YES}) = 0.53$ , OCW): <think>The null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the conditional variance is not equal to the restricted value. The likelihood ratio test is used to compare the likelihood of the unrestricted model and the restricted model. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis. </think> \boxed{A}

**confident-right, kept (mmlu-196, gold A).**

*Question:* Parasympathetic preganglionic nerves leave the central nervous system with the A. third cranial nerves. B. fourth cranial nerves. C. fifth cranial nerves. D. sixth cranial nerves.

before (answer A,  $P(\text{YES}) = 0.96$ , CR): <think>The parasympathetic nervous system is responsible for "rest and digest" functions. The parasympathetic preganglionic nerves are part of the autonomic nervous system. The autonomic nervous system is divided into the sympathetic and parasympathetic branches. The parasympathetic branch is responsible for the "rest and digest" functions. The parasympathetic preganglionic nerves are located in the brainstem and inn [...] rves. The cranial nerves are numbered I-XII. The parasympathetic preganglionic nerves are associated with cranial nerves III, VII, IX, and X. </think> \boxed{A}

after (answer A,  $P(\text{YES}) = 0.68$ , CR): <think>The parasympathetic nervous system is responsible for the "rest and digest" functions. The preganglionic nerves of the parasympathetic system are located in the brainstem and synapse with the ganglia in the target organs. The target organs are the same as the parasympathetic ganglia. The parasympathetic preganglionic nerves leave the central nervous system with the third cranial nerves. </think> \boxed{A}

**confident-right, kept (mmlu-1430, gold A).**

*Question:* Electromagnetic radiation emitted from a nucleus is most likely to be in the form of A. gamma rays B. microwaves C. ultraviolet radiation D. visible light

before (answer A,  $P(\text{YES}) = 1.00$ , CR): <think>The nucleus of an atom is a very small, dense region containing protons and neutrons. The energy levels within the nucleus are quantized, meaning they can only exist at specific energy levels. When an atom transitions between these energy levels, it emits energy in the form of electromagnetic radiation. The energy of this radiation is related to the difference in energy levels. Higher energy levels correspond [...] high energy and short wavelength. Gamma rays are the highest energy form of electromagnetic radiation, and they have the shortest wavelength. </think> \boxed{A}

after (answer A,  $P(\text{YES}) = 0.78$ , CR): <think>The nucleus is a very small and dense object. The energy of the nucleus is very high. The energy of the nucleus is emitted in the form of gamma rays. </think> \boxed{A}

**confident-right, lost (mmlu-6, gold A).**

*Question:* Statement 1 — Every homomorphic image of a group  $G$  is isomorphic to a factor group of  $G$ . Statement 2 — The homomorphic images of a group  $G$  are the same (up to isomorphism) as the factor groups of  $G$ . A. True, True B. False, False C. True, False D. False, True

before (answer A,  $P(\text{YES}) = 0.85$ , CR): <think> Statement 1 says that every homomorphic image of a group  $G$  is isomorphic to a factor group of  $G$ . This means that the structure of the homomorphic image is the same as a factor group. Statement 2 says that the homomorphic images of a group  $G$  are the same (up to isomorphism) as the factor groups of  $G$ . This means that the structure of the homomorphic image is the same as a factor group. </think> \boxed{A}

after (answer A,  $P(\text{YES}) = 0.03$ , NCR): <think> Statement 1 is true. The homomorphic image of a group  $G$  is isomorphic to a factor group of  $G$ . Statement 2 is also true. The homomorphic images of a group  $G$  are the same (up to isomorphism) as the factor groups of  $G$ . </think> \boxed{A}

**confident-right, lost (mmlu-31, gold B).**

*Question:* If  $(G, \cdot)$  is a group such that  $(ab)^{-1} = a^{-1}b^{-1}$ , for all  $a, b$  in  $G$ , then  $G$  is a/an A. commutative semi group B. abelian group C. non-abelian group D. None of these

before (answer B,  $P(\text{YES}) = 0.96$ , CR): <think>The given condition  $(ab)^{-1} = a^{-1}b^{-1}$  implies that the group  $G$  is closed under the operation. This means that the

operation is associative. Also, the inverse of a product is the product of the inverses. This is a characteristic of abelian groups.  $\boxed{B}$

after (answer A,  $P(\text{YES}) = 0.11$ , NCW):  $\langle \text{think} \rangle$  The given condition  $(ab)^{-1} = a^{-1}b^{-1}$  implies that the group  $G$  is commutative. The commutative group is a group where  $a*b = b*a$  for all  $a, b$  in  $G$ .  $\langle \text{think} \rangle$   $\boxed{A}$

## N PROMPTS

All prompts are rendered with the model’s chat template. Text in braces is filled per question; everything else is verbatim.

### Reasoning prompt (all readouts).

Given the following multiple choice question, choose the single best answer (A, B, C, or D). Please think step-by-step and put your thinking inside  $\langle \text{think} \rangle \dots \langle \text{think} \rangle$  tokens. After  $\langle \text{think} \rangle$ , give ONLY your final answer in the format:  $\boxed{[letter]}$

Question: {question}

Options:

- A. {option A}
- B. {option B}
- C. {option C}
- D. {option D}

**M5 (primary).** After the model’s reply, a second user turn is added to the same conversation and  $P(\text{YES})$  is read from the next-token logits, summing the probabilities of the single-token variants of YES and of NO and renormalising over the two:

You answered {letter}. Is that answer correct? Reply with only YES or NO.

The letter is the one with the highest logit after the last  $\boxed{\{$  of the trace, which is appended when the trace has none (6.1% of unsteered confirmatory traces). On the unsteered confirmatory split this letter equals the written letter on 96.9% of questions; the remainder are traces with several boxes or none. A question is confident iff  $P(\text{YES}) \geq 0.5$ .

**M2 (secondary).** The confidence of the same generated answer is the softmax over the four letter logits at the position after  $\boxed{\{$ .

**M4 (secondary; primary in the pilot protocol).** Each option is verified in a separate generation, with at most 192 new tokens:

Answer whether the given answer is the right answer for the following question. Please think step-by-step and put your thinking inside  $\langle \text{think} \rangle \dots \langle \text{think} \rangle$  tokens. After  $\langle \text{think} \rangle$ , answer ONLY with "YES" or "NO".

Question: {question}

Answer: {option text}

$P(\text{YES})$  is read at the answer word after  $\langle \text{think} \rangle$ ; the predicted answer is the option with the largest  $P(\text{YES})$  and its confidence is that value divided by the sum over options.

**Prompting baseline.** Prepend to the user turn (Gemma, served from the bit-identical unsloth/gemma-2-2b-it mirror) or passed as the system message (Qwen):

Be careful and well-calibrated. Only answer confidently when you are sure; if you are uncertain, explicitly acknowledge the uncertainty in your reasoning and confidence.