

PROJECTION-TRANSPORT STEERING: CONDITIONAL CONTROL OF REASONING BEHAVIOR IN LLMs

Anonymous authors

Paper under double-blind review

ABSTRACT

Activation steering changes a language model’s behavior at inference time by editing its residual stream, but most edits in use act on every input. An edit strong enough to remove a behavior where it occurs therefore also changes the inputs that never showed it. We propose Projection-Transport Steering (PTS), a framework that separates when to intervene from how. The action is the optimal-transport map from the projection on a low-dimensional behavioral subspace to a target law, which we show is the unique minimum-displacement edit that leaves the rest of the residual stream untouched. The decision is a likelihood-ratio test on a learned detection statistic. We show that the selectivity of any conditional edit decomposes exactly into the decision’s error rates and the action’s conversion rates, so one ungated run predicts the selectivity of every gated configuration. On overconfident multiple-choice reasoning with Gemma-2-2B, PTS is the most selective of eight methods on held-out questions and lowers calibration error. On free-text toxicity in two model families, PTS decides from the prompt before generating and is more selective than a global shift, prompting, and refusal on the same decision.

1 INTRODUCTION

Many high-level behaviors of large language models appear to be carried by directions in their activations. Activation steering methods (Turner et al., 2023; Panickssery et al., 2024; Zou et al., 2023; Li et al., 2023) read such a direction off contrastive examples and add it to the residual stream at inference time, at the cost of one vector addition per token and no retraining. However, these methods apply the same edit to every input, so an edit strong enough to remove a behavior also acts on the inputs that never showed it: a refusal vector at a useful strength raises refusal on harmful and harmless prompts alike (Arditi et al., 2024).

The problem is sharper for a behavior that the model shows only on some inputs, such as overconfidence in multiple-choice reasoning. The answers to change and the answers to protect are both confident, one wrong and one right, and on `gemma-2-2b-it` the global shift that eliminates every overconfident-wrong answer also eliminates every confident-right one (Figure 1), at a cost of fifteen accuracy points. A smaller coefficient trades one failure for the other, because a shift moves the bulk of the distribution as far as it moves the tail.

In this paper we introduce Projection-Transport Steering (PTS), a framework that separates when to intervene from how. The *action* is a transfer function on the projection of the residual stream onto a behavioral subspace, i.e., a few directions read off the model’s own labeled answers. The *decision* is a likelihood-ratio test on a detection statistic fitted to tell the two confident states apart, and the action runs only where the test fires. Table 1 places published methods in this frame: they differ in the transfer function and in the subspace it acts on, and of those listed only CAST and PCHI decide when to act.

Given the subspace and a target law, i.e., the distribution the steered projection should follow, the action is determined: the optimal-transport map from the unsteered law of the projection to the target law is the unique minimum-displacement edit that leaves the orthogonal complement untouched (Theorem 1). Additive steering is this map exactly when the target is a translate of the unsteered law, and ablation and clamping are the same map for other targets.

method	transfer function $g(p)$	decides when	cost / token
ActAdd, CAA (Turner et al., 2023; Panickssery et al., 2024)	$p + \alpha$	—	$\Theta(d)$
ITI (Li et al., 2023)	$p + \alpha$ on probe-selected heads	—	$\Theta(d)$
RepE (Zou et al., 2023)	$p + \alpha(x)$	—	$\Theta(d)$
directional ablation (Arditi et al., 2024)	0	—	$\Theta(d)$
SAE clamping (Templeton et al., 2024)	c	—	$\Theta(d d_{\text{sae}})$
MiMiC (Singh et al., 2024)	full-space affine map	—	$\Theta(d^2)$
Linear-AcT (Rodriguez et al., 2025a)	$(1-\lambda)p + \lambda(\omega p + \beta)$, per neuron	—	$\Theta(d)$
CAST (Lee et al., 2025)	$p + \alpha$	prompt similarity	$\Theta(d)$
PCHI (Li et al., 2026)	learned head rescaling	probe at the readout	$\Theta(d)$
PTS (ours)	$S(p)$, the transport map to ν	likelihood-ratio test	$\Theta(kd)$

Table 1: **Most published steering operators apply one transfer function to every input.** Each row applies a transfer function g to a k -dimensional projection p of the residual stream, with $k=1$ on a behavioral direction for most rows, $k=d_{\text{sae}}$ for SAE clamping and $k=d$ for the full-space maps. The rows differ in the shape of g and in the k it acts on, and only CAST and PCHI decide when to apply it. Theorem 1 identifies the cost-minimal g for a target law ν , and Proposition 1 writes the selectivity of every row with a decision in terms of that decision and its action.

We score an edit by its selectivity, the share of overconfident-wrong answers it removes minus the share of confident-right answers it destroys. The selectivity of any conditional edit equals the decision’s true-positive rate times the rate at which the action converts flagged target answers, minus its false-positive rate times the rate at which it converts flagged protected ones (Proposition 1). With conversion rates from one ungated run, this identity predicts the selectivity of every gated configuration. When the action converts both states at the same rate, the separability of the detection statistic on unsteered answers caps the selectivity (Theorem 2), and the same quantities determine when a steering action beats overriding every flagged output (Eq. (3)).

A decision that scores a finished answer has to regenerate every input it flags. The same test can read the prompt or a short prefix, so PTS decides once per input and edits in the same pass, at two dot products and one table lookup per token (Proposition 2). Inputs it does not flag are generated without any edit; on HumanEval the prompt decision flags no problem, and pass@1 is exactly its unsteered value (Appendix O).

Our main contributions are:

- A framework and its theory: the optimal-transport characterization of projection-local steering and an exact decomposition of conditional selectivity (Sec. 3), which predicts measured gated selectivity on both model families (Sec. 5.1).
- Held-out evidence on overconfident reasoning (Sec. 5): with every configuration committed before the held-out run, PTS is the most selective of eight methods on 3,000 MMLU questions of `gemma-2-2b-it` and lowers the share of overconfident-wrong answers and the calibration error at non-inferior accuracy. On `Qwen2.5-7B-Instruct` and on ARC-Challenge it is level with or above every baseline after Holm adjustment (Sec. 6).
- Free-text control decided from the prompt (Sec. 6): on toxic continuation in two model families, PTS is more selective than the global shift, prompting, and refusal on the same decision, and than CAST on Qwen, at $1.09\times$ the unsteered wall clock on Gemma.

2 RELATED WORK

Activation steering. ActAdd (Turner et al., 2023) adds the activation difference of one contrast prompt pair to the residual stream, and CAA (Panickssery et al., 2024) averages it over many pairs. ITI (Li et al., 2023) shifts attention heads selected by linear probes, RepE (Zou et al., 2023) reads and writes representation directions, directional ablation (Arditi et al., 2024) removes one direction from every activation, and SAE clamping (Templeton et al., 2024) fixes one sparse feature. AxBench (Wu et al., 2025) finds prompting stronger than such edits, and reliability studies (Tan et al., 2024; Braun et al., 2025) find that their effect varies widely across inputs. All of them apply one edit to every input, with a shape fixed by hand; Theorem 1 derives the shape from a target law, and the additive edits and ablation are its members for a translated target and for a point mass at zero.

Transport and distribution matching. MiMiC (Singh et al., 2024) derives the affine map that matches the mean and covariance of two activation distributions, the Gaussian optimal-transport map, and applies it to the full activation. Linear-AcT (Rodriguez et al., 2025a) fits a univariate affine




Figure 1: **Gating the action by a decision places PTS above the dose curve of the global shift.** Every method is shown on the 3,000 held-out MMLU questions of `gemma-2-2b-it`, at the configuration it won on a separate tuning split. The orange curve is the global shift at increasing dose: it gains removal only by losing confident-right answers, and its last point removes every overconfident-wrong answer and keeps none of the confident-right ones. Grey points keep the PTS decision and dose and replace its direction by one of a hundred random unit directions; post-hoc PTS is more selective than all of them. Dotted lines are levels of selectivity, removal minus the share of confident-right answers lost.

transport map to each neuron with a strength $\lambda \in [0, 1]$ and writes earlier edits as special cases; LinEAS (Rodriguez et al., 2025b) learns such maps across layers end to end. COAST (Nguyen et al., 2026) moves each activation, at fixed norm, to a fixed cosine with a target direction at the smallest second-moment-weighted displacement, which in our terms transports that cosine onto a point mass, and A-LQR (Skifstad et al., 2026) casts steering as linear-quadratic control. Unlike these maps, PTS transports only a k -dimensional projection and leaves its orthogonal complement fixed, which reduces the cost-minimal edit, for any target law and atoms included, to optimal transport in k dimensions, in closed form for $k = 1$.

Deciding when to act. CAST (Lee et al., 2025) compares the prompt’s hidden state with a condition vector at a grid-searched threshold and adds a behavior vector when the comparison holds; we fit its condition on the labeled pool of our probe. DSAS (Ferrando et al., 2026) decouples when to steer from how: a logistic gate scales any steering map per token. SADI (Wang et al., 2025), dynamic activation composition (Scalena et al., 2024) and GAPS (Haider et al., 2026) adapt the edit per input, per generation step and per dimension, respectively. PCHI (Li et al., 2026) conditions a learned rescaling of downstream attention heads on a frozen probe read at a confidence-template token; we run PTS in its setting (Appendix M). Activations predict whether an answer is right (Kadavath et al., 2022; Azaria & Mitchell, 2023) and the first generated tokens whether steering will succeed (Fan et al., 2026); probes of a finished behavior forecast it poorly from a prefix (Kortukov et al., 2026), and a steering write earlier in the same pass can degrade a direction read after it (Luo et al., 2026). Closer to our work, CAST, DSAS and PCHI already separate the decision from the action; PTS adds an exact account of the pair (Prop. 1), with the ceiling and the reference arms it implies, which we run for every conditional method, and it decides before its action starts.

3 METHOD: PROJECTION-TRANSPORT STEERING

A conditional steering operator has to settle what edit to apply and to which inputs. PTS answers with a transport map on a behavioral subspace and a likelihood-ratio test on a detection statistic, and accounts for how the two compose. Proofs are in Appendix C.

The action. Let $h \in \mathbb{R}^d$ be the residual-stream activation at a fixed decoder layer, and let $V \in \mathbb{R}^{k \times d}$ have orthonormal rows spanning what we call the behavioral subspace; $q = Vh$ is the projection. In our experiments $k = 1$, and the row of V is the unit difference of mean activations between confident and hesitant answers (toxic and clean continuations on free text), which we call the confidence coordinate. We write μ for the law of h under unsteered generation, μ_V for the induced law of q , and ν for a target law on $\mathbb{R}^k$. Our goal is an edit that gives the projection the law ν , leaves

the orthogonal complement untouched, and moves h as little as possible, since large edits are what damage generation (Sec. 5).

Definition 1 (admissible steering). $T : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is admissible for (V, ν) if (i) $(I - V^\top V) T(h) = (I - V^\top V) h$ μ -a.s. (locality: the orthogonal complement is untouched), and (ii) $(V \circ T)_\# \mu = \nu$ (goal: steered projections follow the target law). Its cost is $C(T) = \mathbb{E}_\mu \|T(h) - h\|^2$.

Theorem 1 (steering is optimal transport of the projection). Let μ_V be absolutely continuous with finite second moment and let ν have a finite second moment. Then $\min_T C(T) = W_2^2(\mu_V, \nu)$ over admissible T , attained by

$$T^*(h) = h + V^\top (S(Vh) - Vh), \quad (1)$$

where S is the $(\mu_V$ -a.s. unique) Brenier map from μ_V to ν . Every optimal T satisfies $VT(h) = S(Vh)$, hence $T = T^*$ by locality, μ -almost surely.

In words, once the subspace and the target law are fixed, the edit is determined: it moves only the projection, by the transfer function S (g in Table 1). For $k = 1$, $S = F_\nu^{-1} \circ F_{\mu_V}$ is the monotone rearrangement, stored as a 41-knot piecewise-linear lookup; atomlessness of μ_V then suffices and ν may have atoms. Directional ablation is the map onto δ_0 , a clamp at c the map onto the law of $\min(q, c)$, and the shift $S(q) = q + \alpha$ the map onto a translate of μ_V ; for Gaussian marginals S is affine (Knott & Smith, 1984; Gelbrich, 1990). Since every nondecreasing transfer function is the optimal map onto its own image law, Theorem 1 does not rank these edits: it turns the choice of edit into a choice of target law, made on a tuning split (Sec. 4), and it concerns displacement alone. Read backwards, it says when the shift is optimal.

Corollary 1 (when additive steering is optimal, and its limit). (a) $S(q) = q + \alpha$ is the optimal map iff ν is the translate of μ_V by α . (b) Let two subpopulations have the same law of q . If, under an edit, the probability that an input leaves its state depends on h only through q , the edit changes the state of both at the same rate.

The decision and its accounting. The states we must tell apart, overconfident-wrong (OCW) and confident-right (CR), are both confident, and their laws on the confidence coordinate nearly coincide (Sec. 5), so Corollary 1(b) predicts that any transfer function of that coordinate changes both at similar rates, as the global shift does (Fig. 1). A decision therefore reads a second statistic s of the input and runs the action on an event Γ determined by s , in practice $\{s > \tau\}$ for a scalar score; off Γ the output is exactly the unsteered one. We measure the result with selectivity, $\text{Sel} = P(\text{leave OCW} \mid \text{OCW}) - P(\text{leave CR} \mid \text{CR})$, the share of OCW answers removed minus the share of CR answers lost; on free text, toxic and clean continuations take their places. Write TPR and FPR for the rates at which the decision fires on OCW and CR, and ρ_O, ρ_C for the rates at which the action changes the state of a flagged OCW and a flagged CR answer.

Proposition 1 (selectivity decomposes). For any decision and action that leave unflagged outputs unchanged,

$$\text{Sel} = \text{TPR} \cdot \rho_O - \text{FPR} \cdot \rho_C. \quad (2)$$

The proof is one line of total probability. The decision sets (TPR, FPR), and the action sets (ρ_O, ρ_C) on the flagged answers. Since a post-hoc gated operator returns, input by input, the unsteered or the ungated output, one ungated run per action composes every gated configuration, and the ungated rates predict it when the action converts flagged and unflagged answers of a state alike (Sec. 5.1). The identity also fixes three reference arms on the same decision: an identity action measures what regeneration alone changes, a random direction at the same dose what any perturbation of that size changes, and an override of the reported confidence of every flagged answer ($\rho_O = \rho_C = 1$; on free text, a refusal) has selectivity $\text{TPR} - \text{FPR}$. The action beats that override exactly when

$$\text{TPR}(1 - \rho_O) < \text{FPR}(1 - \rho_C), \quad (3)$$

that is, when the share of protected answers flagged but spared exceeds the share of target answers flagged but not converted.

Theorem 2 (optimal decision and selectivity ceiling). Let p_O, p_C be the densities of s on OCW and CR answers, and $r_O(s), r_C(s)$ the probabilities that the action converts such answers at s . (a) If r_O and r_C do not depend on s , then among decisions with equal FPR the likelihood-ratio test on s maximizes Sel (Neyman–Pearson); under equal-covariance Gaussian class conditionals it is the linear discriminant. In general, among decisions with equal $\text{FPR} \cdot \rho_C$, the maximizer thresholds

$p_O(s) r_O(s) / (p_C(s) r_C(s))$. (b) If s is scalar and the concave hull of its ROC has area AUC, every decision $\{s > \tau\}$ has $\text{TPR} - \text{FPR} \leq 2 \text{AUC} - 1$, which caps the override on it and any action with $\rho_O = \rho_C$.

Part (b) is the ceiling of the introduction: computed from unsteered answers alone, it bounds what a threshold decision contributes when the action converts both states alike. The premise of (a) holds only approximately on our data, where flagged answers convert more often than unflagged ones of the same state (Appendix F); Appendix C shows that it matters, so (a) motivates the estimator without certifying it. The decision is a logistic-regression probe, an estimate of the log-likelihood ratio up to a constant that τ absorbs, on the activations of the finished answer, its first t tokens or the prompt, fitted on labeled answers disjoint from every tuning and evaluation split.

Theorem 3 (transport under a fixed decision). *Let $P(\Gamma) > 0$, let $\mu_V(\cdot | \Gamma)$ be absolutely continuous, and let ν_Γ have a finite second moment. Among interventions that are local to V , the identity off Γ , and whose projections on Γ follow ν_Γ , the cost-minimal one applies on Γ the Brenier map from $\mu_V(\cdot | \Gamma)$ to ν_Γ , at cost $P(\Gamma) W_2^2(\mu_V(\cdot | \Gamma), \nu_\Gamma)$.*

The action can thus be fitted to the flagged population without giving up displacement optimality. The theorem holds the decision fixed: under a linear gain and a price on displacement, a graded dose, the form of the per-token interpolation of DSAS, beats the switch exactly when its optimum lies strictly between 0 and 1 with positive probability (Prop. 4). Scoring a finished answer and regenerating the flagged ones doubles the cost of serving; deciding while the answer is written does not.

Proposition 2 (single-pass decision). *Let z_t be computed from the activations of the prompt, or of the prompt and the first t generated tokens. Among decisions measurable with respect to z_t , Theorem 2 holds with z_t in place of s , so the ceiling of an early decision is read from the ROC of z_t on unsteered answers. If the action starts after the last position z_t reads, the operator runs in one forward pass: the decision reads only unedited activations, earlier positions keep their cached keys and values, and for $k = 1$ each token costs at most two extra d -dimensional dot products, one lookup and one vector update.*

4 EXPERIMENTAL SETUP

Models and behavior. The main model is gemma-2-2b-it (Gemma Team, 2024). The confirmatory comparison is repeated on Qwen2.5-7B-Instruct, a second model family with its own directions, probes and tuning, and on ARC-Challenge (Clark et al., 2018), where every decision is refitted with ARC training and validation questions added to its pool. On MMLU (Hendrycks et al., 2021) both models are prompted to reason inside `<think>...</think>` and answer `\boxed{letter}`, with greedy decoding. The action runs at layer 14 of 26 (Gemma) and 14 of 28 (Qwen). Section 6 adds a free-text behavior, toxic continuation, and the setting of the closest conditional method, PCHI (Li et al., 2026).

Readout. The primary readout (M5) scores the answer the model writes. Correctness compares the boxed letter with the gold answer. Confidence is $P(\text{YES})$ when the model is then asked, in the same conversation, whether that answer is correct, and a steered question is steered in both turns. Each question falls into one of four states: overconfident-wrong (OCW, wrong with $P(\text{YES}) \geq 0.5$), nonconfident-wrong, nonconfident-right, and confident-right (CR, right with $P(\text{YES}) \geq 0.5$). Appendix G reports two secondary readouts, the letter logit of the same answer and a per-option verification, and Appendix S gives the prompts.

Splits. The MMLU splits are disjoint subsets of the test set: 400 questions to extract the directions and transport targets, a pool of 4,000 to fit the detection probe, 1,200 on which every configuration is tuned, and 3,000 used once for confirmation. ARC uses 1,012 probe and 400 tuning questions from its training and validation sets and its 1,165 test questions for confirmation.

Metrics. Each question is paired with itself before and after the edit. Removal is $P(\text{not OCW} | \text{OCW})$, retention is $P(\text{CR} | \text{CR})$, and Sel = removal $- (1 - \text{retention})$, which is 1 for an edit that removes every OCW and keeps every CR and 0 for one that treats the two states alike. Because selectivity conditions on the starting state, we also report, over all questions, the change in the share of OCW answers, the 10-bin expected calibration error (ECE) of the stated confidence, and the change in accuracy.




Figure 2: **Gated selectivity is predicted from one ungated run, and the decision can be read from the prompt.** (a) Each point is a decision–action–threshold combination on the tuning split (240 for Gemma in color, 240 for Qwen in gray). The horizontal axis is $\text{TPR}_{\rho_O} - \text{FPR}_{\rho_C}$ with ungated conversion rates, the vertical axis is the gated selectivity composed from the same ungated run, and the dashed line is the identity. Stars are the four gated arms of the Gemma confirmatory split, with their real gated runs on the vertical axis. (b) Held-out AUROC between OCW and CR on the 3,000 confirmatory questions of a probe that reads the prompt, the first t generated tokens, or the whole answer. Bars are 95% bootstrap intervals, and the dotted line is the AUROC of the trace direction used by the CAST-style condition.

Methods and tuning. We compare prompting with a calibration instruction, additive CAA (Panickssery et al., 2024), directional ablation (Arditi et al., 2024), MiMiC (Singh et al., 2024), Linear-AcT (Rodriguez et al., 2025a), CAST with its prompt condition (Lee et al., 2025) fitted on the same labeled pool as our probe, and a CAST-style shift conditioned on the finished trace, which, like post-hoc PTS, regenerates the questions it flags. All methods are reimplemented on our hooks and scored with the same readout. PTS decides from the finished unsteered answer and regenerates the flagged questions (post-hoc), or decides from the prompt, in one forward pass over it before generation, or from the first $t \in \{16, 32, 64, 128\}$ generated tokens, which keeps decision and edit in one pass. Its action is a transport map of Theorem 1 on the confidence coordinate, and its grids cover shifts at several doses and ablation. One rule selects every configuration on the tuning split: the highest selectivity among configurations whose accuracy change is at least -0.01 , or the highest overall when none meets that floor; every confirmatory number is a real run. The budgets are one configuration for prompting and for ablation, eight for CAA, nine each for MiMiC and Linear-AcT, and sixteen for each conditional method, ours and the two CAST variants alike (Appendix E).

Reference arms and statistics. Every conditional arm is run against three references that share its decision: an identity action, which regenerates the flagged inputs with no edit; random unit directions at the same dose (ten per conditional multiple-choice arm; one hundred for PTS, or thirty displaced as far as the selected ablation on Qwen-MMLU); and an override that marks every flagged answer as not confident, so that its selectivity is $\text{TPR} - \text{FPR}$, or replaces every flagged continuation by a refusal on free text. Intervals are 95% paired bootstrap intervals over questions (prompts on free text) with 10^4 resamples. The seven comparisons of PTS with the baselines are two-sided and Holm-adjusted within each setting, each later protocol Holm-adjusts its own tests, and an accuracy change is non-inferior when the lower end of its interval is above -0.02 . The protocol of the multiple-choice comparison, with the hashes of its selected configurations and fitted detectors, was committed to a local repository before its confirmatory rollouts and first pushed together with its results; the protocols of the reference arms, the free-text behavior and the PCHI setting were pushed to a public repository before any of their held-out rollouts (Appendix A).

5 SELECTIVE REMOVAL OF OVERCONFIDENT ANSWERS

A global edit strong enough to remove the behavior also removes the answers it should keep. On the Gemma tuning split the additive edit at $-0.75\|h\|$ removes every overconfident-wrong answer and every confident-right one and lowers accuracy by 15.3 points, and on the held-out questions the same dose keeps none of the 1,399 confident-right answers (Fig. 1). Corollary 1b suggests why: a transfer function of one coordinate converts two states that share its law at the same rate.

method	Δacc	removal	retention	Sel $\uparrow$ [95% CI]	ΔSel (p)	ΔOCW (pp) $\downarrow$	ΔECE $\downarrow$
prompting	+0.005	0.42	0.78	0.200 [0.16, 0.24]	-0.094 (<0.001)	-5.9*	-0.030*
additive CAA	-0.001	0.56	0.66	0.220 [0.18, 0.26]	-0.074 (0.008)	-13.0*	-0.153*
dir. ablation	+0.009	0.46	0.72	0.173 [0.13, 0.21]	-0.121 (<0.001)	-8.7*	-0.111*
MiMiC	+0.000	0.31	0.80	0.117 [0.08, 0.15]	-0.177 (<0.001)	-2.1*	-0.018*
Linear-AcT	+0.005	0.14	0.92	0.061 [0.04, 0.09]	-0.233 (<0.001)	-0.2	-0.003
CAST (prompt)	-0.002	0.46	0.76	0.217 [0.18, 0.26]	-0.077 (0.004)	-11.2*	-0.137*
CAST-style (trace)	-0.005	0.54	0.70	0.248 [0.21, 0.29]	-0.046 (0.020)	-16.6*	-0.094*
PTS, post-hoc	+0.003	0.49	0.80	0.294 [0.26, 0.33]	—	-15.0*	-0.099*
PTS, prefix	-0.008	0.56	0.73	0.294 [0.26, 0.33]	+0.000	-17.2*	-0.136*
PTS, prompt	-0.012 [†]	0.55	0.76	0.306 [0.27, 0.35]	+0.012	-17.6*	-0.084*
same decision, no edit	+0.003	—	—	0.049 [0.03, 0.07]	-0.245 (<0.001)	—	-0.005
same decision, random direction	—	—	—	0.131	-0.163 (<0.001)	—	—

Table 2: **PTS is the most selective of eight methods on held-out questions.** gemma-2-2b-it, 3,000 MMLU questions never used for fitting or tuning ($n_{OCW}=963$, $n_{CR}=1,399$), each method at the configuration selected for it on the tuning split under one rule. ΔSel is the row minus PTS post-hoc (paired bootstrap, 10^4 resamples), with Holm-adjusted p over the seven baselines; ΔOCW is the change in the share of overconfident-wrong answers among all questions. The middle block decides before the answer is written. The last two rows keep the PTS decision and replace its action by an identity or by a random unit direction at the same dose (mean over 100 directions). *: 95% interval excludes 0; [†]: the 95% interval of Δacc reaches below -0.02 .

By Eq. (2) and Theorem 2b, an override, or any action that converts both states at the same rate, reaches at most $\max(\text{TPR} - \text{FPR})$ of its decision statistic, a ceiling computed from unsteered answers before any steered rollout. On the confidence coordinate the edit moves, the two confident states overlap heavily: the held-out ceiling is 0.216 on Gemma and 0.072 on Qwen (AUROC 0.635 and 0.53). We read this as the coordinate tracking how confident the model is, with little information on whether the answer is right. The detection statistic of PTS, a probe fitted on the separate pool of 4,000 labeled answers, raises the ceiling to 0.370 and 0.449 (AUROC 0.741 and 0.785).

5.1 GATED SELECTIVITY FROM ONE UNGATED RUN

Proposition 1 turns one ungated run into a prediction for every decision placed in front of it. On the tuning split we measured the conversion rates of six actions once, composed each with two decisions at 20 thresholds, and compared $\text{TPR } \rho_O - \text{FPR } \rho_C$ with the selectivity composed from the same run (Fig. 2a). The two agree with correlation 0.97 over 240 combinations on Gemma and 0.91 over 240 on Qwen. A post-hoc gated arm is, question by question, either its unsteered or its ungated output, so its selectivity can be composed from one ungated run up to batched-decoding nondeterminism. Across twelve real gated arms on the three held-out splits the composition lies within 0.025 of the real run, with a median gap of 0.005 (Appendix F). We therefore score every gated configuration during tuning from one generation per action, and configurations that decide before the answer by real runs. With $\rho_O = \rho_C = \rho$, Eq. (2) is $\rho(\text{TPR} - \text{FPR})$, which the ceiling above bounds.

5.2 PTS ON HELD-OUT QUESTIONS

In Table 2, post-hoc PTS reaches selectivity 0.294 [0.256, 0.334]: it removes 49% of overconfident-wrong answers and keeps 80% of confident-right ones. The closest baseline, the CAST-style shift conditioned on the finished trace, reaches 0.248, and the other six trail PTS by 0.074 to 0.233; all seven differences remain significant after Holm adjustment. The lead over the CAST-style shift holds when the confidence cutoff moves from 0.5 to 0.4 (+0.044 [0.007, 0.080]) or to 0.6 (+0.068 [0.028, 0.107]).

Over all questions, PTS lowers the share of overconfident-wrong answers by 15.0 points [13.7, 16.4] and the expected calibration error from 0.354 to 0.255, with an accuracy change of +0.003 [-0.009, 0.015], within the non-inferiority margin of -0.02 . Deciding before the answer gives 0.294 from the first 32 generated tokens, in a single pass, and 0.306 from the prompt, the latter with an accuracy interval that reaches the margin ([†]); the prompt probe separates the two confident states nearly as well as the probe on the finished answer (Fig. 2b).

Reference arms. Random unit directions sent through the PTS decision at the PTS dose are less selective than the confidence direction. Over one hundred such directions the mean is 0.131, the 95th percentile 0.237 and the maximum 0.293; PTS is above all of them and above their mean by 0.163 [0.140, 0.188] (gray points in Fig. 1). Regenerating the flagged questions with no edit scores

method	Qwen2.5-7B, MMLU ($n_{OCW}=563$, $n_{CR}=2,007$)				Gemma-2-2B, ARC ($n_{OCW}=228$, $n_{CR}=780$)			
	Δ_{acc}	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	Δ_{ECE} $\downarrow$	Δ_{acc}	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	Δ_{ECE} $\downarrow$
prompting	-0.003	0.212 [0.17, 0.25]	+0.040 (0.313)	-0.004	+0.003	0.238 [0.17, 0.31]	+0.003 (1.000)	+0.009
additive CAA	-0.015 [†]	0.151 [0.12, 0.19]	-0.021 (0.998)	+0.012	-0.001 [†]	0.294 [0.22, 0.37]	+0.059 (0.356)	-0.069*
dir. ablation	-0.009 [†]	0.156 [0.12, 0.19]	-0.016 (0.998)	+0.005	-0.008 [†]	0.215 [0.14, 0.29]	-0.020 (1.000)	-0.055*
MiMiC	-0.002	0.154 [0.12, 0.19]	-0.018 (0.998)	+0.020*	+0.010	0.198 [0.13, 0.26]	-0.037 (1.000)	-0.007
Linear-AcT	-0.003	0.123 [0.09, 0.16]	-0.049 (0.133)	+0.009	+0.003	0.109 [0.06, 0.16]	-0.126 (0.008)	+0.000
CAST (prompt)	-0.011 [†]	0.122 [0.09, 0.16]	-0.050 (0.133)	+0.027*	-0.003	0.109 [0.05, 0.17]	-0.126 (0.003)	-0.028*
CAST-style (trace)	-0.007	0.040 [0.02, 0.06]	-0.132 (<0.001)	+0.013*	+0.001	0.289 [0.22, 0.36]	+0.054 (0.356)	-0.061*
PTS, post-hoc	-0.002	0.172 [0.14, 0.21]	—	-0.000	-0.006 [†]	0.235 [0.17, 0.30]	—	-0.060*
PTS, prefix	-0.008	0.109 [0.08, 0.14]	-0.063	+0.013*	-0.008 [†]	0.276 [0.21, 0.35]	+0.041	+0.000
PTS, prompt	-0.034 [†]	0.120 [0.08, 0.16]	-0.052	+0.049*	-0.027 [†]	0.182 [0.11, 0.25]	-0.053	+0.028*
same decision, no edit	-0.002	0.091 [0.06, 0.12]	-0.081 (<0.001)	+0.000	-0.004	0.016 [-0.01, 0.05]	-0.219 (<0.001)	+0.011
same decision, random direction	—	0.158	-0.014 (0.175)	—	—	0.167	-0.068 (0.004)	—

Table 3: **A second model family and a second benchmark.** Same protocol and notation as Table 2. Qwen refits everything; ARC reuses the MMLU directions and refits the decisions and configurations. The random-direction row averages 100 directions on ARC and, because Qwen selects ablation, 30 directions displaced as far as the ablation displaces the confidence direction.

method	Gemma-2-2B ($n_{tox}=451$, $n_{clean}=1,910$)				Qwen2.5-7B ($n_{tox}=608$, $n_{clean}=1,907$)			
	toxic removed (%)	clean kept (%)	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	toxic removed (%)	clean kept (%)	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)
prompting	42	68	0.104 [0.06, 0.15]	-0.422 (<0.001)	59	79	0.386 [0.34, 0.43]	-0.148 (<0.001)
global shift	90	58	0.474 [0.44, 0.51]	-0.051 (0.010)	84	57	0.403 [0.37, 0.44]	-0.131 (<0.001)
CAST (prompt condition)	81	70	0.515 [0.47, 0.56]	-0.011 (0.585)	74	70	0.443 [0.40, 0.48]	-0.091 (<0.001)
PTS, one pass	73	80	0.526 [0.48, 0.57]	—	82	72	0.534 [0.50, 0.57]	—
PTS, post-hoc	91	81	0.723 [0.69, 0.75]	+0.198 (<0.001)	76	89	0.648 [0.61, 0.68]	+0.114 (<0.001)
same decision, random direction (mean of 100)	—	—	0.099	-0.427 (<0.001)	—	—	0.184	-0.349 (<0.001)
same decision, refuse	82	60	0.427	-0.099 (<0.001)	86	59	0.457	-0.077 (<0.001)

Table 4: **PTS decided from the prompt removes most toxic continuations and keeps most clean ones.** 3,000 held-out RealToxicityPrompts prompts per model, every method at the configuration it won on a separate tuning split. A clean continuation is kept when the steered one is clean and stays within sentence-embedding cosine 0.5 of the unsteered one. Δ_{Sel} is the row minus PTS with the prompt decision. The last two rows keep that decision and replace the action by random unit directions at the same dose (mean of 100) or by a refusal.

0.049, which measures what regeneration in a new batch changes by itself. Appendix K reports the override arm for every conditional method.

Code generation. Gated by the prompt decision, the shift leaves HumanEval unchanged: the decision fires on none of the problems, so pass@1 stays at its unsteered 0.396. The same shift applied to every input at $-0.5\|h\|$, the dose selected for the prompt decision, lowers pass@1 to 0.281 (Appendix O).

6 PTS BEYOND THE MAIN SETTING

On a second family and a second benchmark, no baseline is significantly above PTS. On Qwen2.5-7B-Instruct the tuning split selects ablation as the action, and PTS reaches selectivity 0.172 [0.136, 0.210], 0.081 above its no-edit null. After Holm adjustment it is above the CAST-style trace condition and level with the other six methods, prompting at 0.212 included (Table 3). On ARC-Challenge PTS reaches 0.235 [0.168, 0.303] and lowers ECE by 0.060. It is above Linear-AcT and the CAST prompt condition, whose decision is at chance on ARC (AUROC 0.51), and level with the additive edit and the CAST-style trace condition, which reach 0.294 and 0.289 (Holm $p = 0.36$), and two of its hundred random directions score above it.

On free text, the transport action is above refusing on the same decision. Free text has no single readout that a decision could overwrite, so we test PTS on toxic continuation of RealToxicityPrompts (Gehman et al., 2020), a setting of Linear-AcT, in two model families. A RoBERTa classifier labels a continuation toxic at score 0.5, and a clean continuation is kept only if the steered one is clean and close to the unsteered one, so rewriting counts as damage. The action shifts along the difference of mean activations between toxic and clean continuations, and the decision is a probe on the prompt, read in one forward pass before generation.

On 3,000 held-out prompts per model, PTS reaches selectivity 0.526 and 0.534, and the toxic rate falls from 15.0% to 5.2% on Gemma and from 20.3% to 4.8% on Qwen (Table 4). PTS is above prompting and the global shift in both families and above CAST on Qwen; on Gemma it is level with CAST (+0.011, $p = 0.58$). Refusing on the same decision converts every clean continuation the decision flags, whereas the shift converts only part of them, and PTS is above refusing by 0.099 [0.066, 0.130] and 0.077 [0.054, 0.099], with the sign that Eq. (3) predicts from tuning-split conversion rates for all six gated arms. The direction of PTS lies above all hundred random directions

through the same decision on Gemma and above their 95th percentile on Qwen. Since the decision gates each row inside the original batch, its no-edit null leaves all 3,000 continuations bitwise unchanged.

In the setting of PCHI, a one-direction edit under the same probe is more selective than a learned head rescaling. PCHI (Li et al., 2026) gates a learned rescaling of attention heads with a probe read where the model states its confidence. We rebuild its setting: Qwen3-4B-Instruct answers OpenMathInstruct-2 problems in JSON ending with a yes/no confidence field, a probe on the layer-18 state at that field flags wrong-but-confident answers, and PCHI rescales 128 heads with coefficients learned under its published hyper-parameters. PTS keeps the probe and replaces the learned rescaling by a shift of the same state along one direction. On 3,183 held-out answers it reaches selectivity 0.438 against 0.354 for PCHI (+0.084 [0.058, 0.110]), and an expected calibration error of 0.188 against 0.234 (Appendix M).

With out-of-domain negatives, the decision stays silent on code and open instructions. Refitted with 600 AlpacaEval instructions as negatives, the probes fire on none of the HumanEval and AlpacaEval prompts, so the operator is the exact identity there, and they keep held-out selectivity at 0.280 and 0.300 (Appendix O).

Deciding from the prompt costs a tenth of the unsteered wall clock. On one idle H100, 750 toxicity continuations with Gemma take 5.79 seconds unsteered and 6.31 seconds with PTS, of which 0.46 seconds read the prompt for the decision; a post-hoc decision reaches 0.723 and 0.648 (Table 4) but regenerates the flagged half, at 8.74 seconds.

7 LIMITATIONS

Our study has several limitations. First, when a behavior is read from a single output, as multiple-choice confidence is, overriding that output on the answers the PTS decision flags is at least as selective as the action (0.319 against 0.294 on Gemma and 0.391 against 0.172 on Qwen, Appendix K), and in the setting of PCHI a large random edit at the readout matches the PTS selectivity in 46 of 100 draws (Appendix M); the action improves on an override only on free text (Sec. 6). Second, on Qwen PTS is level with prompting, and its direction is not distinguishable from random directions displaced as far. The ablation selected there creates about as many overconfident-wrong answers as it removes, so neither their share nor the calibration error moves. Third, on Gemma the edit raises the share of reasoning traces that end without a boxed answer by 5.7 points; counting only traces boxed in both arms, the lead over the CAST-style shift is +0.032 [−0.008, 0.071] (Appendix H). Fourth, a single classifier scores toxicity, and the evidence covers two behaviors in two instruction-tuned families of up to seven billion parameters. Fifth, the configuration hashes of the multiple-choice comparison were committed before its held-out run but made public only with its results, whereas every later protocol was public before its runs (Appendix A). Sixth, the free-text operator reads the prompt in a pass of its own; editing generated positions only lowers its selectivity to 0.269 on Gemma and 0.460 on Qwen (Appendix L).

8 CONCLUSION

We have introduced Projection-Transport Steering, which splits an activation edit into an action (a transfer function on the projection onto a behavioral subspace) and a decision (a test on a detection statistic that sets where the action runs). Once the subspace and the target law are fixed, the cost-minimal action is the optimal-transport map between the projection laws, and the additive edit is that map only when the target is a translate of the unsteered law. The selectivity of the composed operator equals the decision’s error rates weighted by the action’s conversion rates, an identity that predicts gated configurations from one ungated run; with Theorem 2 it caps, from unsteered answers alone, what a decision can reach with an action that converts both states alike.

In a comparison frozen before its held-out run, PTS is the most selective of eight methods on MMLU with Gemma, at 0.294 against 0.248 for the strongest baseline, and lowers both the share of overconfident-wrong answers and the calibration error. On free text, deciding from the prompt before generating, it is above global steering and above refusing on its own decision in two model families.

AI USE STATEMENT

We used generative AI tools to polish the manuscript and to retrieve related work. The authors verified every resulting claim and reference and take responsibility for the final content.

REFERENCES

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda, *Refusal in language models is mediated by a single direction*, Advances in Neural Information Processing Systems 37, 2024.
- Amos Azaria and Tom Mitchell, *The internal state of an LLM knows when it’s lying*, Findings of the Association for Computational Linguistics: EMNLP, 2023.
- J. Braun, C. Eickhoff, D. Krueger, S. A. Bahrainian, D. Krasheninnikov. *Understanding (un)reliability of steering vectors in language models*. ICLR Workshop on Foundation Models in the Wild, 2025.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord, *Think you have solved question answering? Try ARC, the AI2 reasoning challenge*, arXiv:1803.05457, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al., *Training verifiers to solve math word problems*, arXiv:2110.14168, 2021.
- Chenrui Fan, Yize Cheng, Ming Li, Soheil Feizi, and Tianyi Zhou, *When is your LLM steerable? Predicting steering success from internal states*, arXiv:2606.11599, 2026.
- Javier Ferrando, Xavier Suau, Pau Rodriguez, and Adrià Gonzalez, *Dynamically scaled activation steering*, arXiv:2512.03661, 2026.
- Gemma Team, *Gemma 2: Improving open language models at a practical size*, arXiv:2408.00118, 2024.
- Umair Haider, Hassan Sajjad, and A. B. Siddique, *GAPS: gated activation steering at dimension granularity*, arXiv:2609.01878, 2026.
- Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith, *RealToxicityPrompts: Evaluating neural toxic degeneration in language models*, Findings of the Association for Computational Linguistics: EMNLP, 2020.
- Matthias Gelbrich, *On a formula for the L^2 Wasserstein metric between measures on Euclidean and Hilbert spaces*, Mathematische Nachrichten **147** (1990), no. 1, 185–203.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt, *Measuring massive multitask language understanding*, International Conference on Learning Representations, 2021.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, et al., *Language models (mostly) know what they know*, arXiv:2207.05221, 2022.
- Martin Knott and Cyril S. Smith, *On the optimal mapping of distributions*, Journal of Optimization Theory and Applications **43** (1984), 39–49.
- Evgenii Kortukov, Piotr Komorowski, Florian Klein, Paula Engl, Gabriele Sarti, Seong Joon Oh, Sebastian Lapuschkin, and Wojciech Samek, *Predicting future behaviors in reasoning models enables better steering*, arXiv:2606.11172, 2026.
- Bruce W. Lee, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Erik Miehl, Pierre Dognin, Manish Nagireddy, and Amit Dhurandhar, *Programming refusal with conditional activation steering*, International Conference on Learning Representations, 2025.

- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg, *Inference-time intervention: Eliciting truthful answers from a language model*, Advances in Neural Information Processing Systems 36, 2023.
- Ke Li, Chongzhe Zhang, Zifan Zeng, Feng Liu, Qunli Zhang, and Zheng Hu, *Calibrating overconfidence without sacrificing confidence: Probe-conditioned head intervention for LLMs*, arXiv:2606.09876, 2026.
- Ziyang Luo, Zhongyao Chu, Xinjie He, Youting Wang, Xukui Qin, Runxiong Wu, and Yan-Syuan Chen, *You only pass once: answering and abstaining together in a single forward pass of a frozen language model*, arXiv:2608.14465, 2026.
- Tam Nguyen, Tu Anh Nguyen, Sina Alemohammad, and Richard G. Baraniuk, *Minimizing collateral damage in activation steering*, arXiv:2605.01167, 2026.
- Nina Panickssery, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner, *Steering Llama 2 via contrastive activation addition*, Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics, 2024.
- Pau Rodriguez, Arno Blaas, Michal Klein, Luca Zappella, Nicholas Apostoloff, Marco Cuturi, and Xavier Suau, *Controlling language and diffusion models by transporting activations*, International Conference on Learning Representations, 2025.
- Pau Rodriguez, Michal Klein, Eleonora Gualdoni, Arno Blaas, Luca Zappella, Marco Cuturi, and Xavier Suau, *LinEAS: End-to-end learning of activation steering with a distributional loss*, arXiv:2503.10679, 2025.
- Daniel Scalena, Gabriele Sarti, and Malvina Nissim, *Multi-property steering of language models with dynamic activation composition*, Proceedings of BlackboxNLP, 2024.
- Shashwat Singh, Shauli Ravfogel, Jonathan Herzig, Roei Aharoni, Ryan Cotterell, and Ponnurangam Kumaraguru, *Representation surgery: Theory and practice of affine steering*, Proceedings of the 41st International Conference on Machine Learning, 2024.
- Julian Skifstad, Xinyue Annie Yang, and Glen Chou, *Local Linearity of LLMs Enables Activation Steering via Model-Based Linear Optimal Control*, arXiv:2604.19018, 2026.
- D. Tan, D. Chanin, A. Lynch, B. Paige, D. Kanoulas, A. Garriga-Alonso, R. Kirk. *Analysing the generalisation and reliability of steering vectors*. NeurIPS, 2024.
- Adly Templeton, Tom Conerly, Jonathan Marcus, et al., *Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet*, Transformer Circuits Thread, 2024.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid, *Steering language models with activation engineering*, arXiv:2308.10248, 2023.
- Weixuan Wang, Jingyuan Yang, and Wei Peng, *Semantics-adaptive activation intervention for LLMs via dynamic steering vectors*, International Conference on Learning Representations, 2025.
- Zhengxuan Wu, Aryaman Arora, Atticus Geiger, Zheng Wang, Jing Huang, Dan Jurafsky, Christopher D. Manning, and Christopher Potts, *AxBench: Steering LLMs? Even simple baselines outperform sparse autoencoders*, Proceedings of the 42nd International Conference on Machine Learning, 2025.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, et al., *Representation engineering: A top-down approach to AI transparency*, arXiv:2310.01405, 2023.

A CONFIRMATORY PROTOCOL

The multiple-choice protocol was committed to the local repository before any steered rollout on the confirmatory split was generated and first pushed together with the results, so its ordering rests on local commit times; every later protocol below was pushed to a public repository, and so timestamped by the hosting service, before its rollouts existed (`rebuttal/protocols/confirmatory-v19.md`, with two amendments made after the tuning pass and before launch), together with the SHA-256 hashes of the selected configurations and of every fitted detector (`confirmatory-v19-gemma.sha256`).

Splits. All splits are subsets of the 14,042-question MMLU test set, identified by index and constructed by `label_pool.split_records`. The extraction split (400) excludes the first evaluation subset of the pilot protocol (Appendix R); the tuning split (1,200) excludes the extraction split and the five pilot evaluation subsets; the detector pool (4,000) and the confirmatory split (3,000) are consecutive blocks of a seeded permutation of the remaining questions. No question appears in two of the four splits.

Selection rule. For every method the configuration with the highest tuning-split selectivity among those with $\Delta_{\text{acc}} \geq -0.01$ is carried unchanged to the confirmatory split. Gated configurations are scored on the tuning split by combining the ungated rollouts of the same split with the gate’s flags (Prop. 1 makes this exact up to batched-decoding nondeterminism; Appendix F); single-pass configurations are real runs. Every confirmatory number is a real run.

Tests. Primary: PTS post-hoc against each of the seven baselines on pooled M5 selectivity, question-level paired bootstrap with 10^4 resamples, two-sided, Holm-adjusted over the seven comparisons. Accuracy: an arm is non-inferior if the lower end of the 95% interval of its accuracy change is above -0.02 . The single-pass and prompt-decision variants, M4 and M2 selectivities, removal, retention and ECE are secondary and reported without a claim.

Released records. Every unsteered and steered answer of the tuning and confirmatory splits is released with its state, confidence and correctness (`results/release/records.csv.gz`); `paper/recompute_from_records.py` recomputes each selectivity and paired interval in Tables 2 and 3, including the null arms, from that file alone. `results/release/toxicity.csv.gz` holds every confirmatory continuation of the free-text setting with its judge score and state (texts omitted for the random-direction arms), and `results/release/pchi_val.csv.gz` the held-out readout gaps of every arm of the PCHI setting.

ARC replication. A second protocol (`confirmatory-v20-arc.md`), frozen before any steered ARC rollout, repeats the Gemma comparison on ARC-Challenge: a detector pool of 1,012 and a tuning split of 400 from the ARC training and validation questions, and the 1,165-question test split as the confirmatory split. The directions, moments and action layer are those fitted on MMLU; every decision is refitted on the ARC pool together with the MMLU pool, and every configuration is selected on the ARC tuning split with the same rule and grids.

Null arms. A third protocol (`confirmatory-v21-null.md`), with the job list and the hash of the runner, was committed and pushed to a public repository before any null rollout existed; the commit is timestamped by the hosting service. On each confirmatory split it adds the frozen PTS decision with an identity action, the frozen decision with the selected action on a unit random direction (seed 0), and an ungated rerun with an identity action. It registers two one-sided hypotheses per setting, PTS above each gated null, with Holm adjustment over the six tests, and the net measures of Appendix H for every arm. Five of the six tests pass; PTS against the single random direction on ARC gives $+0.041 [-0.029, 0.109]$, Holm $p = 0.12$.

Reference arms. A fourth protocol (`confirmatory-v22-references.md`), pushed before any of its rollouts existed, extends the no-edit null to all four conditional methods, replaces the single random direction by ten (seeds 0 to 9), and adds random directions for the sycophancy edit. It registers 24 one-sided tests (four methods, three settings, no-edit and random-edit null), Holm-adjusted together, of which 18 pass (Table 14); a count of the cases in which Eq. (3) with ungated conversion rates predicts the sign of steer minus matched override (10 of 13); and paired intervals for selectivity at cutoffs 0.4 and 0.6 and over boxed traces for PTS against the CAST-style condition on Gemma. The v22 random-edit null supersedes the single-direction test of v21.

Free text and dense nulls. A fifth protocol (`confirmatory-v23-freetext.md`), pushed with the hashes of the tuning selections, fitted decisions and code before any confirmatory rollout of the toxicity setting existed, fixes the data, judge, retention criterion, arms and ten primary hypotheses of Appendix L, and extends the random directions of PTS post-hoc to one hundred on Gemma-MMLU and Gemma-ARC and to thirty displacement-matched directions on Qwen (Table 15).

PCHI setting. A sixth protocol (`confirmatory-v24-pchi.md`), pushed after the train-split fitting and tuning and before any validation rollout, fixes the arms and three hypotheses of Appendix M.

Analyses added after the protocol. The detect-and-override reference, the difference-of-means decision fitted on 4,400 labels, the latency measurement (Appendix N), the decision-by-action matrix and the trace-only and confidence-turn-only edits (Appendix J), the sycophancy experiment (Appendix I), calibration and the out-of-domain negatives were run after the confirmatory results were known. They use the same splits and the frozen decisions, and we report them as exploratory.

B A PILOT OPERATOR ON FURTHER MODELS

In a pilot study we froze a simplified operator, an action along a detection direction, after a development analysis on Qwen3, and tested it on held-out TruthfulQA panels of 237 questions (50 in-domain, 187 shifted) for Mistral-7B, OLMo-2-7B and Granite-8B. Each comparison was against a norm-matched spherical action, with 10,000 paired BCa resamples and Bonferroni-adjusted two-sided intervals, and a model passed only if calibrated MC2 and MC1 both improved with intervals excluding zero.

model / panel	comparison	Δ calibrated MC2 [adj. 95% CI]	Δ MC1 [adj. 95% CI]	decision
OLMo-2 / TruthfulQA	matched spherical	+0.137 [+0.105, +0.168]	+0.105 [+0.025, +0.177]	pass
OLMo-2 / TruthfulQA	no intervention	+0.299 [+0.251, +0.345]	+0.287 [+0.177, +0.384]	pass
OLMo-2 / TruthfulQA	strongest isotropic	+0.220 [+0.180, +0.258]	+0.194 [+0.097, +0.287]	pass
OLMo-2 / TruthfulQA	unprompted fit	+0.295 [+0.247, +0.343]	+0.287 [+0.177, +0.384]	pass

Table 5: Held-out OLMo-2 validation on 237 TruthfulQA questions, 50 in-domain and 187 shifted. Intervals use 10,000 paired BCa resamples and the predeclared Bonferroni family; the table shows the matched action and representative construction controls of the seven-control family, all of which pass.

On OLMo-2 the operator passed against the matched action and against all seven pre-registered construction controls (Table 5). Mistral improved calibrated MC2 by +0.153 [+0.108, +0.197], but its MC1 interval $[-0.059, +0.122]$ included zero, and Granite did not clear the screening threshold. A separate screen of the same three models on BBQ and ETHICS, six model-behavior cells with disjoint fitting, screening and test units, advanced one cell, Mistral on ETHICS, to a 1,000-scenario test. There the operator improved pooled calibrated MC2 by +0.107 [+0.090, +0.124] and lowered MC1 on label-1 scenarios by 0.093 [0.071, 0.124].

C PROOFS

Theorem 1. By locality, $T(h) - h \in \text{span}(V^\top) \mu$ -a.s., hence $\|T(h) - h\|^2 = \|VT(h) - Vh\|^2$ because the rows of V are orthonormal, and $C(T) = \mathbb{E}\|VT(h) - Vh\|^2$. The joint law γ of $(Vh, VT(h))$ is a coupling of (μ_V, ν) by the goal constraint, so $C(T) = \int \|y - x\|^2 d\gamma \geq W_2^2(\mu_V, \nu)$. Because μ_V is absolutely continuous and both laws have finite second moments, Brenier’s theorem gives a unique optimal coupling, induced by $S = \nabla\varphi$ with φ convex. $T^*(h) = h + V^\top(S(Vh) - Vh)$ is admissible ($VV^\top = I_k$ gives $VT^*(h) = S(Vh)$, and the complement is untouched) and attains the bound. For any minimiser the induced coupling is optimal, hence equal to the unique one, so $VT(h) = S(Vh)$ a.s. $\square$

For $k = 1$ absolute continuity can be weakened to atomlessness of μ_V : the monotone rearrangement $F_\nu^{-1} \circ F_{\mu_V}$ is then an optimal map for every ν with a finite second moment, including targets with atoms such as δ_0 (ablation) or a law truncated at c (clamp); uniqueness of the map can fail only on a μ_V -null set.

Proposition 3 (metric covariance). *If the cost is $(\Delta h)^\top M(\Delta h)$ with $M \succ 0$, the admissible problem is optimal transport of μ_V to ν in $\mathbb{R}^k$ under the cost $\|\cdot\|_{M_V}^2$, $M_V = VMV^\top \succ 0$, and the optimal operator is $T^*(h) = h + V^\top(S_M(Vh) - Vh)$ with $S_M = M_V^{-1/2} \circ \nabla\varphi \circ M_V^{1/2}$, φ convex.*

Proof. Locality writes $\Delta h = V^\top a$ with $a \in \mathbb{R}^k$, and the goal constraint with $VV^\top = I_k$ fixes $a = S(Vh) - Vh$ for the map S in use. Then $\|\Delta h\|_M^2 = a^\top M_V a$. The cost is a latent-space Mahalanobis cost on $\mathbb{R}^k$; it is not the full-space Mahalanobis distance of Δh unless M commutes with $V^\top V$. In the coordinates $z = M_V^{1/2} q$ it is Euclidean, the change of variables is internal to $\mathbb{R}^k$, and Brenier's theorem in z (the pushforward of an absolutely continuous law under an invertible linear map is absolutely continuous) gives $\nabla \varphi$; mapping back gives S_M . $\square$

Corollary 1. (a) If $\nu = \mu_V * \delta_\alpha$ the map $x \mapsto x + \alpha$ is the gradient of a convex function pushing μ_V to ν , hence the Brenier map. Conversely, if the optimal map is $x \mapsto x + \alpha$ then ν is its pushforward, the shift. For one-dimensional Gaussians the monotone map is $x \mapsto m_t + (\sigma_t/\sigma_s)(x - m_s)$, with slope σ_t/σ_s . (b) If two groups of inputs have the same law of p and the probability that an input changes state depends on h only through p , every transfer function changes both groups at the same rate, since both rates are the same integral against the law of p . A decision that reads p alone therefore adds little, which is why the detection statistic reads the full activation. $\square$

Proposition 1. Let R be the event that an OCW trace leaves OCW and L the event that a CR trace leaves CR. Off Γ the output equals the unsteered output, which is in its original state by definition, so $R \subseteq \Gamma$ and $L \subseteq \Gamma$ up to null sets. Then $P(R \mid OCW) = P(\Gamma \mid OCW)P(R \mid OCW, \Gamma) = \text{TPR } \rho_O$, and likewise $P(L \mid CR) = \text{FPR } \rho_C$; subtracting gives Eq. (2). If, within OCW, the indicator that the action converts the trace is independent of Γ , then $\rho_O = P(R^{\text{ungated}} \mid OCW)$, and likewise for CR; the ungated selectivity is $\rho_O - \rho_C$, and the difference is $\rho_C(1 - \text{FPR}) - \rho_O(1 - \text{TPR})$. $\square$

The identity requires that the unflagged output be exactly the unsteered one. The post-hoc gate satisfies this by construction. The single-pass gate satisfies it for traces on which it never fires, since the action hook is the identity until the decision token.

Theorem 2. Write p_O, p_C for the densities of s on OCW and CR answers and $r_O(s), r_C(s)$ for the probabilities that the action converts an answer of each state with statistic s . For a decision region Γ , $\text{TPR } \rho_O = \int_\Gamma p_O r_O$ and $\text{FPR } \rho_C = \int_\Gamma p_C r_C$, so $\text{Sel}(\Gamma) = \int_\Gamma (p_O r_O - p_C r_C)$. Among regions with a given value of $\int_\Gamma p_C r_C$, the Neyman–Pearson lemma applied to the nonnegative measures $p_O r_O$ and $p_C r_C$ shows that $\int_\Gamma p_O r_O$ is maximised by a superlevel set of their ratio, which is the general statement in (a). If r_O and r_C are constant in s , the ratio is a constant multiple of p_O/p_C , so every objective increasing in TPR and decreasing in FPR is maximised on the likelihood-ratio ROC curve, which is (a). The assumption matters: with $s \in \{a, b\}$, $p_O(a) = p_O(b) = \frac{1}{2}$, $p_C(a) = \frac{1}{4}$, $p_C(b) = \frac{3}{4}$, $r_O(a) = 0.2$, $r_O(b) = 1$, $r_C(a) = 1$ and $r_C(b) = 0$, the likelihood ratio ranks a above b , yet $\Gamma = \{b\}$ gives selectivity 0.5, against -0.15 for $\{a\}$ and 0.35 for $\{a, b\}$. On our data r does depend on s : flagged answers are converted more often than unflagged answers of the same state (Appendix F). Our detector is a logistic-regression probe, a discriminative estimate of the log-likelihood ratio up to the prior term; (a) states what the probe approximates, not that the fitted probe is optimal. (b) The concavified ROC contains the path $(0, 0) \rightarrow (\text{FPR}^*, \text{TPR}^*) \rightarrow (1, 1)$ through the Youden-optimal point, whose area is $\frac{1}{2}(1 + J)$, so $J \leq 2 \text{AUC} - 1$. The override reference has selectivity $\text{TPR} - \text{FPR} \leq J$, and with $\rho_O = \rho_C = \rho$, Eq. (2) gives $\text{Sel} = \rho(\text{TPR} - \text{FPR}) \leq J$. $\square$

Theorem 3. Split $C(T)$ over Γ and Γ^c . The second term is zero by the identity constraint; the first, with the conditional goal constraint, is Theorem 1 for $\mu(\cdot \mid \Gamma)$, which applies when $\mu_V(\cdot \mid \Gamma)$ is absolutely continuous. $\square$

This result fixes the decision. It does not say that a threshold gate is the best decision for the displacement-priced objective; the next statement shows that it generally is not, and when it is.

Proposition 4 (graded dose). Let $T_\lambda(h) = h + \lambda(s) V^\top (S(Vh) - Vh)$ with a dose $\lambda : \mathbb{R} \rightarrow [0, 1]$ measurable in the detector statistic s . Write $\pi(s) = P(OCW \mid s)$, $d(s) = \mathbb{E}[\|S(Vh) - Vh\|^2 \mid s] > 0$, and suppose the behavioral gain is linear in the dose, with slope $\beta_1 > 0$ on OCW and $-\beta_0 < 0$ on CR. The dose maximising gain minus displacement priced at $1/(2c)$ is

$$\lambda^*(s) = \left[\frac{c(\beta_0 + \beta_1)(\pi(s) - \tau)}{d(s)} \right]_0^1, \quad \tau = \frac{\beta_0}{\beta_0 + \beta_1}.$$

It is zero exactly on $\{\pi \leq \tau\}$ and nondecreasing in the ratio $(\pi(s) - \tau)/d(s)$. It is nondecreasing in π when d is constant, but not in general. The threshold gate $\mathbf{1}[\pi > \tau]$ is optimal iff $\lambda^* \in \{0, 1\}$ almost surely, and strictly suboptimal iff $P(0 < \lambda^* < 1) > 0$.

Proof. The displacement of T_λ is $\lambda(s)^2 \|S(Vh) - Vh\|^2$, with expectation $\mathbb{E}_s[\lambda(s)^2 d(s)]$; the gain is $\mathbb{E}_s[\lambda(s)((\beta_0 + \beta_1)\pi(s) - \beta_0)]$. The objective is an integral over s of a strictly concave quadratic in λ , maximised pointwise on $[0, 1]$ by the clipped stationary point. It is the unique maximiser a.s., so the indicator is optimal iff it agrees with λ^* a.s., i.e. iff $\lambda^* \in \{0, 1\}$ a.s.; otherwise the objective is strictly larger at λ^* on a set of positive measure. For non-monotonicity in π , take s_1, s_2 with $\pi(s_1) = 0.6 < \pi(s_2) = 0.7$, $\tau = 0.5$, $c(\beta_0 + \beta_1) = 1$, $d(s_1) = 0.2$ and $d(s_2) = 1$: then $\lambda^*(s_1) = 0.5 > \lambda^*(s_2) = 0.2$. For a counterexample to strict suboptimality at every finite price, let π take only the values 0.1 and 0.9 with $\tau = 0.5$ and $d \leq 0.4 c(\beta_0 + \beta_1)$: then $\lambda^* \in \{0, 1\}$ a.s. and the threshold gate is optimal. $\square$

Proposition 2. Let $\mathcal{F}_t$ be generated by the first t generated tokens and z_t the mean detector feature over them. A decision measurable with respect to $\sigma(z_t)$ is a test between the class-conditional laws of z_t , so Theorem 2 applies with z_t in place of s and nothing is assumed about how z_t relates to the finished trace. The reader hook sits at a layer at or above the actor; within one forward pass the actor at position t sees the decision accumulated over positions $< t$, the running mean needs one sum and one counter per sequence, and an edit at position t changes only the keys and values written at t .

Proposition 5 (computational cost). *Any algorithm implementing a non-degenerate intervention local to a subspace spanned by a dense direction reads all d coordinates of h , so it costs $\Omega(d)$ per token. The edit of Eq. (1) costs $\Theta(kd + k \log K)$ with a K -knot lookup, and the detector adds one d -dimensional dot product per token.*

Proof. If a coordinate i with $(\hat{v})_i \neq 0$ is never read, an adversary perturbs h_i , which changes $p = \langle h, \hat{v} \rangle$ and, on a set of positive measure, the correct output, while the algorithm’s output is unchanged. The upper bound counts Vh , a binary search in the lookup, and $h + V^\top(\cdot)$. $\square$

D GSM8K MCQ CONSTRUCTION

Each GSM8K test item with integer gold g receives three deterministic distractors drawn (seeded by item index) from $\{g \pm 1, g \pm 2, g \pm 10, 2g, \lfloor g/2 \rfloor, \text{round}(1.5g), \text{round}(0.8g), g \pm \text{U}\{3..9\}\}$, deduplicated, shuffled with the gold into four options. This converts free-form arithmetic into selection while keeping the reasoning burden; we report transfer numbers on it and do not claim GSM8K free-form results.

E TUNING GRIDS

Every configuration searched on the tuning split, with the selected one marked $\star$. Gated rows are scored by composing the ungated rollouts of the same split with the decision’s flags (Appendix F); single-pass rows are real runs. The rule is the highest selectivity among configurations with $\Delta_{\text{acc}} \geq -0.01$, and, when no configuration of a family meets the floor, the highest selectivity overall; this fallback applies to the additive, CAST and prompt-decision families on Qwen, where every shift in the grid costs more than one accuracy point.

Table 6: Tuning split, gemma-2-2b-it ($n = 1,200$, M5).

family	configuration	Δ_{acc}	removal	retention	Sel
additive CAA	alpha-0.375	-0.022	0.831	0.424	+0.255
additive CAA	alpha-0.25 $\star$	-0.002	0.536	0.661	+0.197
additive CAA	alpha-0.5	-0.042	0.974	0.136	+0.110
additive CAA	alpha-0.125	+0.010	0.315	0.763	+0.078
additive CAA	alpha-0.75	-0.153	1.000	0.000	+0.000
additive CAA	alpha-1.0	-0.205	1.000	0.000	+0.000
additive CAA	alpha-1.25	-0.270	1.000	0.000	+0.000
additive CAA	alpha-1.5	-0.299	1.000	0.000	+0.000
MiMiC	L10.reg0.0001 $\star$	+0.022	0.349	0.823	+0.172
MiMiC	L10.reg0.1	+0.012	0.349	0.809	+0.158
MiMiC	L10.reg0.01	+0.007	0.320	0.814	+0.135
MiMiC	L18.reg0.1	+0.022	0.260	0.860	+0.121
MiMiC	L18.reg0.01	+0.026	0.242	0.869	+0.111
MiMiC	L18.reg0.0001	+0.023	0.250	0.857	+0.107
MiMiC	L14.reg0.01	+0.001	0.224	0.857	+0.081
MiMiC	L14.reg0.1	+0.004	0.216	0.850	+0.066
MiMiC	L14.reg0.0001	-0.013	0.214	0.841	+0.055
Linear-Act	L10.lam0.5 $\star$	+0.016	0.133	0.936	+0.069

	family	configuration	Δ acc	removal	retention	Sel
810						
811	Linear-AcT	L10.lam1.0	+0.015	0.138	0.919	+0.057
812	Linear-AcT	L18.lam1.0	+0.012	0.130	0.926	+0.056
	Linear-AcT	L14.lam0.25	+0.007	0.128	0.924	+0.052
813	Linear-AcT	L14.lam1.0	+0.009	0.138	0.913	+0.051
814	Linear-AcT	L14.lam0.5	+0.010	0.125	0.919	+0.044
	Linear-AcT	L18.lam0.5	+0.004	0.117	0.926	+0.043
815	Linear-AcT	L10.lam0.25	+0.010	0.112	0.919	+0.031
	Linear-AcT	L18.lam0.25	+0.005	0.096	0.922	+0.019
816	CAST-style, trace condition	tau.cr.q30.alpha-0.375	-0.020	0.815	0.458	+0.273
817	CAST-style, trace condition	tau.cr.q50.alpha-0.375	-0.015	0.732	0.532	+0.264
818	CAST-style, trace condition	tau.cr.q70.alpha-0.5	-0.014	0.638	0.611	+0.249
	CAST-style, trace condition	tau.cr.q70.alpha-0.375*	-0.006	0.560	0.682	+0.242
819	CAST-style, trace condition	tau.cr.q70.alpha-0.75	-0.066	0.646	0.595	+0.241
820	CAST-style, trace condition	tau.cr.q50.alpha-0.5	-0.028	0.859	0.380	+0.239
	CAST-style, trace condition	tau.cr.q50.alpha-0.75	-0.110	0.875	0.353	+0.228
821	CAST-style, trace condition	tau.cr.q90.alpha-0.5	-0.001	0.411	0.802	+0.214
822	CAST-style, trace condition	tau.cr.q90.alpha-0.75	-0.022	0.414	0.795	+0.209
	CAST-style, trace condition	tau.cr.q90.alpha-0.375	+0.000	0.370	0.832	+0.202
823	CAST-style, trace condition	tau.cr.q30.alpha-0.25	-0.001	0.521	0.670	+0.190
824	CAST-style, trace condition	tau.cr.q50.alpha-0.25	+0.000	0.474	0.716	+0.190
	CAST-style, trace condition	tau.cr.q30.alpha-0.5	-0.038	0.945	0.244	+0.189
825	CAST-style, trace condition	tau.cr.q70.alpha-0.25	+0.000	0.352	0.809	+0.161
826	CAST-style, trace condition	tau.cr.q30.alpha-0.75	-0.139	0.969	0.178	+0.147
	CAST-style, trace condition	tau.cr.q90.alpha-0.25	+0.001	0.245	0.894	+0.139
827	CAST, prompt condition	L16.q50.alpha-0.375	-0.018	0.607	0.673	+0.280
828	CAST, prompt condition	L16.q30.alpha-0.375	-0.017	0.711	0.565	+0.276
	CAST, prompt condition	L14.q30.alpha-0.375	-0.017	0.708	0.567	+0.275
829	CAST, prompt condition	L14.q50.alpha-0.375	-0.018	0.594	0.680	+0.274
830	CAST, prompt condition	L16.q50.alpha-0.5	-0.022	0.682	0.580	+0.262
	CAST, prompt condition	L14.q50.alpha-0.5	-0.021	0.669	0.588	+0.258
831	CAST, prompt condition	L16.q50.alpha-0.75	-0.069	0.690	0.562	+0.252
	CAST, prompt condition	L14.q50.alpha-0.75	-0.070	0.674	0.571	+0.245
832	CAST, prompt condition	L14.q30.alpha-0.5	-0.034	0.810	0.415	+0.225
833	CAST, prompt condition	L16.q30.alpha-0.5	-0.034	0.815	0.408	+0.223
834	CAST, prompt condition	L14.q30.alpha-0.75	-0.106	0.820	0.378	+0.198
	CAST, prompt condition	L16.q30.alpha-0.75	-0.104	0.826	0.369	+0.195
835	CAST, prompt condition	L16.q30.alpha-0.25*	-0.007	0.443	0.730	+0.172
836	CAST, prompt condition	L14.q30.alpha-0.25	-0.008	0.440	0.728	+0.168
	CAST, prompt condition	L16.q50.alpha-0.25	-0.009	0.362	0.792	+0.153
837	CAST, prompt condition	L14.q50.alpha-0.25	-0.010	0.354	0.792	+0.146
838	PTS post-hoc	q60.alpha-0.75	-0.031	0.570	0.763	+0.334
	PTS post-hoc	q50.alpha-0.75	-0.050	0.674	0.654	+0.328
839	PTS post-hoc	q40.alpha-0.75	-0.073	0.771	0.555	+0.326
	PTS post-hoc	q30.alpha-0.75	-0.093	0.862	0.454	+0.316
840	PTS post-hoc	q30.alpha-0.375	-0.013	0.737	0.578	+0.315
841	PTS post-hoc	q60.alpha-0.375*	-0.003	0.510	0.790	+0.300
842	PTS post-hoc	q40.alpha-0.375	-0.015	0.654	0.645	+0.299
	PTS post-hoc	q50.alpha-0.375	-0.012	0.586	0.707	+0.293
843	PTS post-hoc	q30.alpha-0.25	+0.002	0.477	0.735	+0.212
844	PTS post-hoc	q40.alpha-0.25	+0.002	0.424	0.767	+0.191
	PTS post-hoc	q50.alpha-0.25	+0.002	0.380	0.806	+0.186
845	PTS post-hoc	q60.alpha-0.25	+0.005	0.320	0.864	+0.184
846	PTS post-hoc	q30.ablate	+0.012	0.388	0.765	+0.153
	PTS post-hoc	q40.ablate	+0.014	0.336	0.795	+0.131
847	PTS post-hoc	q50.ablate	+0.016	0.281	0.841	+0.122
	PTS post-hoc	q60.ablate	+0.012	0.232	0.887	+0.119
848	PTS, prefix decision	q50.at32.alpha-0.375*	+0.004	0.573	0.737	+0.310
849	PTS, prefix decision	q50.at16.alpha-0.375	-0.002	0.573	0.731	+0.304
	PTS, prefix decision	q30.at16.alpha-0.375	+0.002	0.729	0.572	+0.302
850	PTS, prefix decision	q50.at64.alpha-0.375	+0.017	0.562	0.739	+0.301
851	PTS, prefix decision	q30.at32.alpha-0.375	+0.002	0.690	0.602	+0.293
852	PTS, prefix decision	q50.at128.alpha-0.375	-0.001	0.547	0.708	+0.255
	PTS, prefix decision	q30.at64.alpha-0.375	+0.010	0.656	0.595	+0.252
853	PTS, prefix decision	q30.at128.alpha-0.375	+0.000	0.651	0.574	+0.225
854	PTS, prefix decision	q50.at32	+0.019	0.245	0.894	+0.139
	PTS, prefix decision	q30.at32	+0.018	0.320	0.816	+0.137
855	PTS, prefix decision	q50.at64	+0.002	0.214	0.898	+0.111
856	PTS, prefix decision	q30.at16	-0.002	0.326	0.781	+0.106
	PTS, prefix decision	q50.at16	+0.003	0.224	0.878	+0.102
857	PTS, prefix decision	q50.at128	+0.000	0.182	0.912	+0.094
858	PTS, prefix decision	q30.at64	-0.001	0.250	0.836	+0.086
	PTS, prefix decision	q30.at128	-0.003	0.214	0.871	+0.085
859	PTS, prompt decision	q60.alpha-0.75	-0.041	0.570	0.767	+0.337
860	PTS, prompt decision	q60.alpha-0.5*	-0.009	0.568	0.769	+0.336
	PTS, prompt decision	q50.alpha-0.75	-0.061	0.656	0.661	+0.317
861	PTS, prompt decision	q50.alpha-0.5	-0.018	0.654	0.663	+0.316
862	PTS, prompt decision	q60.alpha-0.375	-0.015	0.518	0.790	+0.308
	PTS, prompt decision	q50.alpha-0.375	-0.010	0.589	0.712	+0.301
863	PTS, prompt decision	q30.alpha-0.375	-0.013	0.724	0.569	+0.293
	PTS, prompt decision	q40.alpha-0.5	-0.024	0.724	0.567	+0.291

family	configuration	Δ acc	removal	retention	Sel
PTS, prompt decision	q40_alpha-0.75	-0.084	0.729	0.562	+0.291
PTS, prompt decision	q40_alpha-0.375	-0.012	0.646	0.643	+0.289
PTS, prompt decision	q30_alpha-0.5	-0.020	0.826	0.463	+0.288
PTS, prompt decision	q30_alpha-0.75	-0.099	0.839	0.445	+0.284
PTS, prompt decision	q30_alpha-0.25	+0.007	0.464	0.735	+0.199
PTS, prompt decision	q60_alpha-0.25	-0.002	0.320	0.871	+0.191
PTS, prompt decision	q40_alpha-0.25	+0.004	0.409	0.781	+0.190
PTS, prompt decision	q50_alpha-0.25	+0.002	0.365	0.825	+0.190
earlier gate + ablation	u_crq50_ablate*	+0.007	0.378	0.740	+0.118
ungated actions	ablate	+0.007	0.451	0.691	+0.141
ungated actions	clamp.q50	+0.017	0.307	0.814	+0.122
ungated actions	clamp.q30	+0.012	0.339	0.763	+0.102
ungated actions	otq-cal	+0.004	0.115	0.926	+0.040

Table 7: Tuning split, Qwen2.5-7B-Instruct ($n = 1,200$, M5).

family	configuration	Δ acc	removal	retention	Sel
additive CAA	alpha-0.125*	-0.026	0.212	0.911	+0.123
additive CAA	alpha-0.375	-0.050	0.208	0.892	+0.100
additive CAA	alpha-0.25	-0.044	0.200	0.887	+0.087
additive CAA	alpha-0.5	-0.076	0.204	0.855	+0.059
additive CAA	alpha-0.75	-0.138	0.233	0.759	-0.007
additive CAA	alpha-1.0	-0.217	0.412	0.544	-0.043
additive CAA	alpha-1.25	-0.334	0.567	0.272	-0.161
additive CAA	alpha-1.5	-0.473	0.375	0.230	-0.395
MiMiC	L18.reg0.0001	-0.013	0.279	0.906	+0.185
MiMiC	L18.reg0.01*	-0.007	0.263	0.911	+0.173
MiMiC	L14.reg0.01	-0.018	0.254	0.912	+0.166
MiMiC	L14.reg0.0001	-0.023	0.233	0.912	+0.145
MiMiC	L10.reg0.1	-0.019	0.250	0.890	+0.140
MiMiC	L14.reg0.1	-0.021	0.217	0.923	+0.140
MiMiC	L18.reg0.1	-0.013	0.233	0.906	+0.139
MiMiC	L10.reg0.01	-0.022	0.237	0.892	+0.129
MiMiC	L10.reg0.0001	-0.031	0.233	0.868	+0.101
Linear-AcT	L18.lam1.0*	-0.004	0.192	0.948	+0.140
Linear-AcT	L14.lam1.0	+0.002	0.175	0.955	+0.130
Linear-AcT	L14.lam0.25	+0.004	0.138	0.969	+0.106
Linear-AcT	L10.lam1.0	-0.008	0.158	0.947	+0.105
Linear-AcT	L18.lam0.25	-0.006	0.138	0.957	+0.095
Linear-AcT	L10.lam0.25	-0.013	0.133	0.958	+0.092
Linear-AcT	L18.lam0.5	-0.006	0.142	0.950	+0.091
Linear-AcT	L10.lam0.5	-0.013	0.125	0.951	+0.076
Linear-AcT	L14.lam0.5	-0.005	0.113	0.961	+0.073
CAST-style, trace condition	tau.cr.q30_alpha-0.375	-0.043	0.188	0.907	+0.094
CAST-style, trace condition	tau.cr.q30_alpha-0.25	-0.035	0.183	0.906	+0.089
CAST-style, trace condition	tau.cr.q50_alpha-0.375	-0.040	0.158	0.918	+0.076
CAST-style, trace condition	tau.cr.q50_alpha-0.25	-0.033	0.154	0.919	+0.074
CAST-style, trace condition	tau.cr.q30_alpha-0.5	-0.067	0.188	0.879	+0.067
CAST-style, trace condition	tau.cr.q70_alpha-0.375	-0.024	0.108	0.945	+0.053
CAST-style, trace condition	tau.cr.q50_alpha-0.5	-0.059	0.154	0.898	+0.052
CAST-style, trace condition	tau.cr.q70_alpha-0.25	-0.022	0.104	0.945	+0.049
CAST-style, trace condition	tau.cr.q50_alpha-0.75	-0.079	0.188	0.858	+0.045
CAST-style, trace condition	tau.cr.q90_alpha-0.25	-0.010	0.067	0.975	+0.041
CAST-style, trace condition	tau.cr.q70_alpha-0.5	-0.036	0.104	0.937	+0.041
CAST-style, trace condition	tau.cr.q90_alpha-0.375*	-0.009	0.062	0.979	+0.041
CAST-style, trace condition	tau.cr.q90_alpha-0.5	-0.018	0.071	0.970	+0.041
CAST-style, trace condition	tau.cr.q70_alpha-0.75	-0.048	0.125	0.912	+0.037
CAST-style, trace condition	tau.cr.q30_alpha-0.75	-0.098	0.212	0.824	+0.036
CAST-style, trace condition	tau.cr.q90_alpha-0.75	-0.023	0.079	0.956	+0.035
CAST, prompt condition	L14.q30_alpha-0.25*	-0.032	0.167	0.916	+0.082
CAST, prompt condition	L16.q30_alpha-0.25	-0.032	0.167	0.916	+0.082
CAST, prompt condition	L14.q30_alpha-0.375	-0.040	0.163	0.918	+0.081
CAST, prompt condition	L16.q30_alpha-0.375	-0.040	0.163	0.918	+0.081
CAST, prompt condition	L16.q50_alpha-0.375	-0.034	0.142	0.931	+0.072
CAST, prompt condition	L14.q50_alpha-0.375	-0.034	0.142	0.929	+0.071
CAST, prompt condition	L14.q50_alpha-0.25	-0.026	0.138	0.932	+0.069
CAST, prompt condition	L16.q50_alpha-0.25	-0.027	0.133	0.932	+0.065
CAST, prompt condition	L14.q30_alpha-0.5	-0.058	0.167	0.890	+0.057
CAST, prompt condition	L16.q30_alpha-0.5	-0.058	0.167	0.890	+0.057
CAST, prompt condition	L16.q50_alpha-0.5	-0.050	0.133	0.914	+0.048
CAST, prompt condition	L14.q50_alpha-0.5	-0.051	0.133	0.913	+0.046
CAST, prompt condition	L16.q50_alpha-0.75	-0.078	0.158	0.868	+0.026
CAST, prompt condition	L14.q50_alpha-0.75	-0.079	0.150	0.868	+0.018
CAST, prompt condition	L14.q30_alpha-0.75	-0.105	0.188	0.821	+0.009
CAST, prompt condition	L16.q30_alpha-0.75	-0.105	0.183	0.822	+0.006
PTS post-hoc	q40_ablate*	-0.006	0.254	0.927	+0.181
PTS post-hoc	q50_ablate	-0.002	0.233	0.940	+0.173
PTS post-hoc	q30_ablate	-0.009	0.254	0.918	+0.172

family	configuration	Δacc	removal	retention	Sel
PTS post-hoc	q60_ablate	-0.003	0.212	0.948	+0.161
PTS post-hoc	q40_alpha-0.375	-0.044	0.183	0.907	+0.090
PTS post-hoc	q30_alpha-0.375	-0.048	0.192	0.898	+0.090
PTS post-hoc	q40_alpha-0.25	-0.043	0.183	0.902	+0.085
PTS post-hoc	q60_alpha-0.375	-0.032	0.154	0.929	+0.084
PTS post-hoc	q30_alpha-0.25	-0.045	0.192	0.892	+0.083
PTS post-hoc	q60_alpha-0.25	-0.028	0.150	0.931	+0.081
PTS post-hoc	q50_alpha-0.375	-0.039	0.163	0.916	+0.078
PTS post-hoc	q50_alpha-0.25	-0.037	0.154	0.912	+0.066
PTS post-hoc	q60_alpha-0.75	-0.062	0.158	0.883	+0.041
PTS post-hoc	q50_alpha-0.75	-0.083	0.179	0.854	+0.033
PTS post-hoc	q40_alpha-0.75	-0.098	0.204	0.826	+0.030
PTS post-hoc	q30_alpha-0.75	-0.110	0.217	0.805	+0.021
PTS, prefix decision	q30_at32	-0.010	0.163	0.951	+0.113
PTS, prefix decision	q30_at32_alpha-0.375	-0.012	0.183	0.929	+0.113
PTS, prefix decision	q30_at16*	-0.008	0.163	0.948	+0.111
PTS, prefix decision	q50_at32	-0.007	0.142	0.966	+0.108
PTS, prefix decision	q50_at16	-0.006	0.138	0.960	+0.097
PTS, prefix decision	q30_at64	-0.008	0.133	0.951	+0.084
PTS, prefix decision	q30_at16_alpha-0.375	-0.030	0.167	0.916	+0.082
PTS, prefix decision	q50_at32_alpha-0.375	-0.013	0.133	0.943	+0.077
PTS, prefix decision	q30_at64_alpha-0.375	-0.014	0.133	0.942	+0.075
PTS, prefix decision	q50_at64	-0.005	0.108	0.963	+0.072
PTS, prefix decision	q30_at128_alpha-0.375	-0.005	0.083	0.969	+0.052
PTS, prefix decision	q50_at128_alpha-0.375	-0.004	0.079	0.972	+0.051
PTS, prefix decision	q50_at16_alpha-0.375	-0.029	0.113	0.938	+0.051
PTS, prefix decision	q50_at64_alpha-0.375	-0.018	0.083	0.955	+0.038
PTS, prefix decision	q30_at128	+0.002	0.058	0.979	+0.037
PTS, prefix decision	q50_at128	+0.002	0.050	0.977	+0.027
PTS, prompt decision	q30_alpha-0.375*	-0.045	0.204	0.901	+0.105
PTS, prompt decision	q30_alpha-0.25	-0.042	0.196	0.899	+0.095
PTS, prompt decision	q40_alpha-0.375	-0.043	0.188	0.907	+0.094
PTS, prompt decision	q50_alpha-0.375	-0.040	0.171	0.916	+0.086
PTS, prompt decision	q40_alpha-0.25	-0.043	0.175	0.906	+0.081
PTS, prompt decision	q50_alpha-0.25	-0.038	0.163	0.914	+0.077
PTS, prompt decision	q60_alpha-0.375	-0.033	0.138	0.935	+0.072
PTS, prompt decision	q30_alpha-0.5	-0.068	0.204	0.868	+0.072
PTS, prompt decision	q60_alpha-0.25	-0.031	0.133	0.929	+0.063
PTS, prompt decision	q50_alpha-0.5	-0.056	0.163	0.897	+0.059
PTS, prompt decision	q40_alpha-0.5	-0.068	0.175	0.874	+0.049
PTS, prompt decision	q60_alpha-0.5	-0.047	0.129	0.914	+0.044
PTS, prompt decision	q50_alpha-0.75	-0.080	0.179	0.855	+0.034
PTS, prompt decision	q60_alpha-0.75	-0.067	0.146	0.884	+0.030
PTS, prompt decision	q30_alpha-0.75	-0.108	0.217	0.810	+0.026
PTS, prompt decision	q40_alpha-0.75	-0.101	0.196	0.826	+0.022
earlier gate + ablation	u_crg50_ablate*	-0.008	0.196	0.938	+0.134
ungated actions	ablate	-0.008	0.258	0.913	+0.171
ungated actions	clamp.q30	-0.015	0.221	0.918	+0.139
ungated actions	otq.cal	-0.003	0.154	0.962	+0.116
ungated actions	clamp.q50	-0.008	0.183	0.928	+0.112

F THE SELECTIVITY IDENTITY IN PRACTICE

Proposition 1 needs the output to be unchanged wherever the decision does not fire. For the post-hoc and prompt decisions this holds by construction, so a gated arm equals, question by question, either its unsteered row or its ungated row. Composing the two requires one ungated run per action and one set of decision scores, and it differs from a real gated run only through batched-decoding nondeterminism: a flagged question is generated in a different batch in the two cases. Table 8 compares the composition with the real run for the four gated arms of each confirmatory split, and reports the prediction from the ungated conversion rates (independence within each state).

Across the twelve arms the composition differs from the real run by at most 0.025 (Qwen, PTS post-hoc) and by 0.005 at the median, so composed cells in Fig. 3 carry an error of that size. The prediction is higher than the real run in eleven of twelve rows, and the same bias appears on the tuning split (Fig. 2a). Flagged answers are converted more often than the average answer of their state in both states: for PTS post-hoc, 0.91 of flagged overconfident-wrong answers change state against 0.84 ungated, and 0.89 of flagged confident-right answers against 0.59 ungated. The larger increase on the protected state lowers selectivity below the independence prediction. The composition, which uses no such assumption, is what scores the gated configurations during selection.

setting	arm	TPR	FPR	ρ_O	ρ_C	predicted	composed	real run
Gemma, MMLU	PTS post-hoc	0.541	0.222	0.841	0.585	+0.325	+0.296	+0.294
	PTS at the prompt	0.549	0.242	0.981	0.872	+0.328	+0.306	+0.306
	CAST-style, trace	0.603	0.380	0.841	0.585	+0.285	+0.244	+0.248
	CAST, prompt	0.805	0.635	0.559	0.339	+0.234	+0.207	+0.217
Qwen, MMLU	PTS post-hoc	0.869	0.478	0.240	0.084	+0.168	+0.148	+0.172
	PTS at the prompt	0.920	0.594	0.247	0.111	+0.161	+0.126	+0.120
	CAST-style, trace	0.286	0.101	0.247	0.111	+0.059	+0.030	+0.040
	CAST, prompt	0.817	0.670	0.277	0.098	+0.161	+0.143	+0.122
Gemma, ARC	PTS post-hoc	0.675	0.405	0.513	0.219	+0.258	+0.239	+0.235
	PTS at the prompt	0.417	0.247	0.991	0.896	+0.191	+0.181	+0.182
	CAST-style, trace	0.912	0.767	0.513	0.219	+0.300	+0.297	+0.289
	CAST, prompt	0.386	0.379	0.513	0.219	+0.115	+0.107	+0.109

Table 8: Real gated runs against the composition of ungated runs and against the prediction TPR ρ_O – FPR ρ_C with ungated conversion rates.

G SECONDARY READOUTS

Table 9 reports the Gemma confirmatory arms under the two secondary readouts. The letter-logit confidence M2 assigns confidence above 0.5 to nearly every answer, so its OCW stratum is close to the set of wrong answers and every method moves it little. Under M4 the edits act on four separate verification generations; PTS is level with the ungated edits there, which is expected for a decision fitted to M5 labels.

method (Gemma, confirmatory)	M2 (letter logit)		M4 (verification)	
	Sel [95% CI]	Δ vs PTS	Sel [95% CI]	Δ vs PTS
prompting	+0.083 [+0.052, +0.113]	−0.008 [−0.045, +0.028]	+0.138 [+0.086, +0.191]	−0.102 [−0.167, −0.034]
additive CAA	+0.089 [+0.059, +0.120]	−0.001 [−0.036, +0.034]	+0.242 [+0.187, +0.296]	+0.003 [−0.061, +0.065]
dir. ablation	+0.092 [+0.062, +0.122]	+0.002 [−0.034, +0.036]	+0.246 [+0.193, +0.299]	+0.007 [−0.054, +0.070]
MiMiC	+0.064 [+0.036, +0.092]	−0.027 [−0.063, +0.008]	+0.244 [+0.189, +0.298]	+0.004 [−0.059, +0.067]
Linear-AcT	+0.029 [+0.009, +0.050]	−0.061 [−0.091, −0.031]	+0.111 [+0.068, +0.155]	−0.128 [−0.185, −0.070]
CAST (prompt)	+0.073 [+0.045, +0.102]	−0.017 [−0.052, +0.017]	+0.207 [+0.155, +0.259]	−0.032 [−0.092, +0.027]
CAST-style (trace)	+0.074 [+0.046, +0.103]	−0.016 [−0.041, +0.010]	+0.224 [+0.173, +0.277]	−0.015 [−0.063, +0.033]
PTS, post-hoc	+0.090 [+0.064, +0.118]	—	+0.239 [+0.192, +0.286]	—
PTS, prompt	+0.099 [+0.069, +0.128]	+0.008 [−0.023, +0.038]	+0.245 [+0.198, +0.292]	+0.006 [−0.039, +0.052]

Table 9: Secondary readouts on the confirmatory split; paired bootstrap intervals, no multiplicity adjustment.

method (Qwen, confirmatory)	M2 (letter logit)		M4 (verification)	
	Sel [95% CI]	Δ vs PTS	Sel [95% CI]	Δ vs PTS
prompting	+0.132 [+0.102, +0.163]	+0.013 [−0.021, +0.046]	+0.178 [+0.144, +0.213]	+0.073 [+0.033, +0.114]
additive CAA	+0.118 [+0.089, +0.149]	−0.000 [−0.033, +0.034]	+0.096 [+0.060, +0.132]	−0.009 [−0.050, +0.033]
dir. ablation	+0.119 [+0.089, +0.150]	+0.000 [−0.024, +0.025]	+0.192 [+0.159, +0.225]	+0.087 [+0.057, +0.117]
MiMiC	+0.130 [+0.100, +0.160]	+0.011 [−0.023, +0.045]	+0.104 [+0.073, +0.136]	−0.001 [−0.040, +0.038]
Linear-AcT	+0.099 [+0.072, +0.127]	−0.020 [−0.053, +0.013]	+0.078 [+0.051, +0.106]	−0.026 [−0.062, +0.009]
CAST (prompt)	+0.110 [+0.081, +0.140]	−0.009 [−0.043, +0.025]	+0.159 [+0.125, +0.193]	+0.054 [+0.012, +0.096]
CAST-style (trace)	+0.030 [+0.012, +0.048]	−0.089 [−0.120, −0.059]	+0.121 [+0.096, +0.145]	+0.016 [−0.021, +0.053]
PTS, post-hoc	+0.119 [+0.090, +0.149]	—	+0.105 [+0.076, +0.135]	—
PTS, prompt	+0.096 [+0.063, +0.128]	−0.023 [−0.059, +0.012]	+0.170 [+0.134, +0.206]	+0.065 [+0.022, +0.108]

Table 10: Secondary readouts on the Qwen confirmatory split.

Table 11 gives the full state transitions of the frozen PTS arms under M5. Selectivity counts only the first and last rows; the table also shows the answers the edit corrects or breaks and the few answers it turns overconfident.

Table 12 reports calibration of the stated confidence on the confirmatory split: expected calibration error (10 bins), Brier score, and the AUROC with which the confidence separates right from wrong answers.

H NET EFFECTS OF EVERY ARM

Selectivity is computed on the answers that start in OCW or CR; it does not see answers that enter OCW from another state. Table 13 reports, for every confirmatory arm and the three null arms of Appendix A, selectivity net of the no-edit null, the change in the share of OCW answers among all questions, the change in the rate of confident answers among wrong and among right answers, ΔECE , $\Delta Brier$, the change in the share of traces without a `\boxed{ }` answer (which then receive

	before \ after	OCW	NCW	NCR	CR
Gemma, PTS	OCW	489	350	118	6
	NCW	10	298	51	2
	NCR	2	61	207	7
	CR	11	95	171	1122
Qwen, PTS	OCW	425	43	10	85
	NCW	54	121	23	27
	NCR	21	25	115	44
	CR	81	25	40	1861

Table 11: State transitions of the frozen PTS arms on the confirmatory split (counts of questions).

arm	Gemma-2-2B			Qwen2.5-7B		
	ECE ↓	Brier ↓	AUROC ↑	ECE ↓	Brier ↓	AUROC ↑
unsteered	0.354	0.372	0.643	0.215	0.230	0.741
shift −0.25, ungated	0.202	0.281	0.659	0.238	0.245	0.741
ablation, ungated	0.243	0.296	0.653	0.220	0.232	0.743
PTS, post-hoc	0.255	0.294	0.689	0.215	0.229	0.740
override, probe on answer	0.234	0.285	0.684	0.139	0.195	0.760
override, CAST condition	0.197	0.297	0.646	0.194	0.222	0.714

Table 12: Calibration of the stated confidence (M5) on the confirmatory split.

a forced box, Appendix S), and selectivity recomputed without the questions whose steered trace has no box. The ungated rerun reproduces the unsteered outputs exactly, because it uses the same batches; the gated no-edit null regenerates the flagged questions in different batches, and the resulting nondeterminism alone yields selectivity 0.049 (Gemma), 0.091 (Qwen) and 0.016 (ARC). Excluding traces without a box leaves the ordering on Gemma unchanged: PTS post-hoc 0.257, CAST-style 0.233, additive 0.215.

arm	Sel	Sel-null	ΔOCW (pp)	ΔP(conf wrong) (pp)	ΔP(conf right) (pp)	ΔECE	ΔBrier	Δno box (pp)	Sel, boxed only
<i>Gemma-2-2B, MMLU</i>									
prompting	0.200	+0.151	−5.9	−12.7	−6.1	−0.030	−0.032	+1.1	0.185
additive CAA	0.220	+0.171	−13.0	−29.4	−18.9	−0.153	−0.091	+1.3	0.215
dir. ablation	0.173	+0.124	−8.7	−18.6	−13.0	−0.111	−0.076	+0.4	0.175
MiMiC	0.117	+0.068	−2.1	−4.7	−4.6	−0.018	−0.020	+0.0	0.106
Linear-AcT	0.061	+0.012	−0.2	+0.4	+0.6	−0.003	−0.006	−0.1	0.053
CAST (prompt)	0.217	+0.168	−11.2	−25.6	−13.4	−0.137	−0.086	+0.6	0.206
CAST-style (trace)	0.248	+0.199	−16.6	−38.0	−23.1	−0.094	−0.067	+8.2	0.233
PTS, post-hoc	0.294	+0.245	−15.0	−33.8	−16.0	−0.099	−0.079	+5.7	0.257
PTS, prefix	0.294	+0.245	−17.2	−39.5	−20.2	−0.136	−0.091	+2.0	0.284
PTS, prompt	0.306	+0.257	−17.6	−40.8	−18.7	−0.084	−0.077	+15.0	0.247
null: decision, no edit	0.049	+0.000	+0.0	+0.5	+0.9	−0.005	−0.007	−0.4	0.037
null: decision, random dir.	0.215	+0.166	−8.9	−19.6	−7.9	−0.077	−0.061	+7.2	0.185
rerun, no edit	0.000	−0.049	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000
<i>Qwen2.5-7B, MMLU</i>									
prompting	0.212	+0.121	−1.3	−5.6	−2.0	−0.004	−0.002	+3.6	0.196
additive CAA	0.151	+0.060	+2.9	+6.5	+3.5	+0.012	+0.008	−0.8	0.149
dir. ablation	0.156	+0.065	+1.3	+2.3	+0.8	+0.005	+0.001	+0.6	0.156
MiMiC	0.154	+0.063	+1.4	+4.9	+0.2	+0.020	+0.014	+0.5	0.151
Linear-AcT	0.123	+0.032	+0.6	+1.4	+0.1	+0.009	+0.004	+0.7	0.114
CAST (prompt)	0.122	+0.031	+3.7	+10.4	+3.0	+0.027	+0.021	+0.5	0.124
CAST-style (trace)	0.040	−0.051	+2.4	+7.2	+2.4	+0.013	+0.009	+0.2	0.036
PTS, post-hoc	0.172	+0.081	+0.6	+1.6	+0.7	−0.000	−0.002	+0.5	0.166
PTS, prefix	0.109	+0.018	+1.0	+1.6	−0.6	+0.013	+0.010	+0.1	0.092
PTS, prompt	0.120	+0.029	+6.9	+15.2	+5.4	+0.049	+0.039	+1.1	0.123
null: decision, no edit	0.091	+0.000	−0.1	−0.9	−0.2	+0.000	−0.000	+0.1	0.084
null: decision, random dir.	0.119	+0.028	−0.1	−1.6	+0.4	−0.002	−0.004	−0.2	0.117
rerun, no edit	0.000	−0.091	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000
<i>Gemma-2-2B, ARC</i>									
prompting	0.238	+0.222	−1.8	−6.1	−3.7	+0.009	+0.005	−0.3	0.234
additive CAA	0.294	+0.278	−5.9	−23.3	−11.4	−0.069	−0.013	+1.2	0.289
dir. ablation	0.215	+0.198	−3.5	−15.6	−8.6	−0.055	−0.007	+1.0	0.216
MiMiC	0.198	+0.182	−1.7	−3.8	−3.6	−0.007	−0.003	+0.7	0.198
Linear-AcT	0.109	+0.093	−0.1	+0.4	+0.7	+0.000	−0.001	+0.3	0.108
CAST (prompt)	0.109	+0.093	−2.0	−8.6	−5.3	−0.028	−0.002	+0.3	0.090
CAST-style (trace)	0.289	+0.273	−5.4	−20.9	−10.7	−0.061	−0.015	+1.6	0.276
PTS, post-hoc	0.235	+0.219	−3.8	−16.2	−5.4	−0.060	−0.025	+1.0	0.226
PTS, prefix	0.276	+0.260	−9.2	−37.1	−17.6	+0.000	+0.003	+1.5	0.267
PTS, prompt	0.182	+0.166	−8.1	−35.8	−18.4	+0.028	+0.019	+10.6	0.150
null: decision, no edit	0.016	+0.000	+1.7	+5.3	+0.5	+0.011	+0.011	+0.3	0.020
null: decision, random dir.	0.194	+0.178	−3.7	−15.0	−6.6	−0.006	−0.001	+2.7	0.176
rerun, no edit	0.000	−0.016	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000

Table 13: Net effects on the confirmatory splits (point estimates; intervals for the main-text quantities are given there). “null” is the frozen decision with no edit.

I SYCOPHANCY

Setting. For each of the 3,000 Gemma confirmatory questions we append a second user turn to the unsteered conversation, “I think the answer is {letter}. Are you sure? Think again and give your final answer as `\boxed{letter}`.”, where the suggested letter is a seeded wrong option if the first answer was correct and the gold letter otherwise. The model *caves* if it switches from a correct first answer to the suggestion, *holds* if it keeps it, and *accepts* if it switches from a wrong first answer to the suggested gold answer. Of the 3,000 questions, 1,459 cave and 1,242 accept. Selectivity is the share of caving questions that no longer cave minus the share of accepting questions that no longer accept.

Decision and actions. The decision is a logistic-regression probe on the mean activation of the first-turn trace, fitted on the detector pool to predict whether the first answer is correct, with layer (18) and regularisation chosen by AUROC on the tuning split; on the confirmatory split its AUROC is 0.722. The action is a deference direction at layer 14, the difference of mean activations at the last prompt token of the second turn between caving and holding questions of the detector pool, ablated or subtracted at doses $\{0.125, 0.25, 0.5, 0.75, 1.0\}$ of the mean activation norm. The override keeps the first answer whenever the probe fires. Configurations are selected on the tuning split with the rule of Sec. 4.

Result. On the confirmatory split the override reaches selectivity 0.311 $[0.277, 0.345]$ and raises accuracy by 17.5 points; the best gated edit (probe at its 30% quantile, dose 0.25) reaches 0.096 $[0.071, 0.121]$ at +3.0 points, and the best ungated edit 0.056 $[0.027, 0.085]$. The deference edit converts few flagged answers, so by Eq. (3) the override is the stronger use of the same probe. This experiment was designed after the overconfidence results and is exploratory.

J OVERRIDE REFERENCES OF FIVE DECISIONS




Figure 3: **The override reference rises with the held-out quality of the decision.** Each marker pair is one of five decisions (probes on the answer or the prompt, the CAST prompt condition, and a difference-of-means direction on the answer fitted on 400 or 4,400 labelled answers), paired with the action and threshold it selects on the tuning split. Selectivity is composed on the confirmatory split with 95% bootstrap intervals; filled markers steer, hollow markers override the reported confidence of the flagged answers and edit nothing.

Fig. 3 pairs each of five decisions with the action and threshold it selects on the tuning split and with its own override. Besides the three decisions of the protocol, it includes a difference-of-means direction on the answer fitted on 400 or 4,400 labelled answers. The override rises with the held-out AUROC of the decision: on Gemma from 0.201 for the CAST prompt condition (AUROC 0.63) to 0.323 $[0.284, 0.362]$ for the answer probe (AUROC 0.74), and on Qwen to 0.433 $[0.393, 0.472]$, where it also lowers ECE from 0.216 to 0.139. Cells are composed from ungated runs and selected after the protocol; the thresholds of the steered and overridden points differ, so the matched comparisons are those of Table 14.

Where the edit acts. We reran the frozen PTS decision with the action applied only during the reasoning trace or only during the confidence turn. On Gemma (shift) the confidence-turn edit reaches 0.274 $[0.237, 0.312]$ and the trace-only edit 0.159; on Qwen (ablation) the trace-only edit reaches 0.172 and the confidence-turn edit 0.006.

K REFERENCE ARMS FOR EVERY CONDITIONAL METHOD

Table 14 applies the three references of Sec. 3 to the four conditional multiple-choice arms of each setting, with ten random directions per arm (protocol v22), and to the sycophancy edit of Appendix I. The matched override lowers the stated confidence of exactly the answers the arm’s own decision flags. Eighteen of the 24 registered null tests pass after Holm adjustment; the six failures are on Qwen or involve the CAST prompt condition on ARC. Eq. (3), evaluated with ungated conversion rates before the gated runs, predicts the sign of steer minus override in ten of thirteen rows, and the three misses have a predicted difference within 0.021 of zero. The CAST conditions fire on many protected answers and their shift spares most of them, so steering beats their override; the answer probes flag fewer, and overriding with them is at least as selective. Table 15 reports the hundred-direction test of protocol v23 for PTS post-hoc.

setting	method	Sel	no-edit null		random-edit null (10)		matched override		
			Sel	Δ	mean Sel	Δ	Sel	steer – override [95% CI]	Eq. 3 sign
Gemma, MMLU	PTS, post-hoc	0.294	0.049	+0.245*	0.141	+0.153*	0.319	−0.025 [−0.04, −0.01]	no
	PTS, prompt	0.306	0.041	+0.265*	0.147	+0.159*	0.307	−0.001 [−0.00, +0.00]	no
	CAST-style (trace)	0.248	0.049	+0.199*	0.112	+0.136*	0.223	+0.025 [+0.00, +0.05]	yes
	CAST (prompt)	0.217	0.044	+0.173*	0.105	+0.112*	0.169	+0.048 [+0.01, +0.09]	yes
Qwen, MMLU	PTS, post-hoc	0.172	0.091	+0.081*	0.122	+0.051*	0.391	−0.218 [−0.26, −0.17]	yes
	PTS, prompt	0.120	0.077	+0.043	0.224	−0.104	0.326	−0.206 [−0.25, −0.16]	yes
	CAST-style (trace)	0.040	0.035	+0.005	0.080	−0.040	0.185	−0.145 [−0.18, −0.11]	yes
	CAST (prompt)	0.122	0.065	+0.057*	0.182	−0.059	0.147	−0.025 [−0.07, +0.02]	no
Gemma, ARC	PTS, post-hoc	0.235	0.016	+0.219*	0.165	+0.070*	0.270	−0.035 [−0.10, +0.03]	yes
	PTS, prompt	0.182	0.045	+0.137*	0.119	+0.063*	0.169	+0.013 [+0.01, +0.02]	yes
	CAST-style (trace)	0.289	0.077	+0.213*	0.211	+0.078*	0.146	+0.144 [+0.07, +0.22]	yes
	CAST (prompt)	0.109	0.020	+0.089*	0.078	+0.032	0.006	+0.103 [+0.04, +0.16]	yes
Gemma, sycophancy	probe, deference edit	0.096	—	—	0.050	+0.045	0.255	−0.159	yes

Table 14: Reference arms on the confirmatory splits. Δ is the method minus the reference; *: Holm-adjusted $p < 0.05$ over the 24 registered tests.

setting	PTS	random mean	95th pct.	max	directions $\geq$ PTS	PTS – mean [95% CI]
Gemma, MMLU (100)	0.294	0.131	0.237	0.293	0	+0.163 [+0.140, +0.188]
Gemma, ARC (100)	0.235	0.167	0.230	0.257	2	+0.068 [+0.017, +0.119]
Qwen, MMLU (30, matched)	0.172	0.158	0.175	0.178	3	+0.014 [−0.017, +0.045]

Table 15: PTS post-hoc against random unit directions sent through its own decision at its own dose (Gemma and ARC, 100 directions) and, for the ablation selected on Qwen, against 30 random directions displaced as far as the ablation displaces the confidence direction ($h - \langle h, \hat{v} \rangle r$).

L FREE-TEXT BEHAVIOR: TOXIC CONTINUATION

Data and readout. Prompts are RealToxicityPrompts prompts whose own toxicity is at least 0.5 (21,744), shuffled with seed 0 and split into extraction (1,000), probe pool (3,000), tuning (1,000) and confirmation (3,000). Each prompt is rendered with the chat template as “Continue the following text.” and continued for 40 greedy tokens. The judge is `s-nlp/roberta-toxicity-classifier`; a continuation is toxic at probability 0.5 or above, refused when it contains a refusal phrase in its first 200 characters or is degenerate (fewer than three words or a distinct bigram ratio below one half), and clean otherwise. Unsteered, Gemma produces 451 toxic and 1,910 clean continuations on the confirmation prompts and Qwen 608 and 1,907. Selectivity is the share of toxic continuations that are no longer toxic minus the share of clean continuations lost, where a clean continuation is kept when the steered one is clean and is identical to, or has `all-mpnet-base-v2` cosine at least 0.5 with, the unsteered one. On the Gemma tuning split, this criterion keeps 84% of clean continuations at the smallest shift and 36% at the largest, so rewriting a continuation into different content is counted as damage.

Operator. The direction is the difference of mean layer-14 activations over toxic and clean continuations of the extraction split, and the dose is a multiple of the mean activation norm. The single-pass decision is a logistic probe on the mean prompt activation (layer 18 for Gemma, 10 for Qwen; tuning AUROC 0.808 and 0.823), the post-hoc decision a probe on the unsteered continuation (0.946 and 0.890), and the CAST condition the cosine of the mean prompt activation to the toxic-minus-clean difference of the probe pool. The decision gates each row inside the original batch, so an unflagged row is generated exactly as in the unsteered run. The tuning grid holds five shift doses, ablation and prompting, and three thresholds for each decision; the selection rule is the highest tuning selectivity.

Results. Table 4 gives the confirmatory comparison. Nine of the ten registered tests pass after Holm adjustment; the one that does not is PTS against CAST on Gemma (+0.011, $p = 0.58$). Refusing on a decision converts every flagged clean continuation, so by Eq. (3) a gated edit beats it whenever it spares enough of them; Eq. (3) with conversion rates from the tuning split predicts the sign of steer minus refuse for all six gated arms. Against a hundred random directions through the same decision, PTS lies above the 95th percentile (0.311 on Gemma, 0.479 on Qwen), with three of the hundred directions above it on Qwen.

Edit restricted to generated positions. In the arm above the decision reads the prompt in a forward pass of its own and the edit then acts on every position. An exploratory variant, not part of protocol v23, starts the edit at the first generated token, so that decision and edit share a single pass, at the dose selected for the full edit. It reaches 0.269 [0.223, 0.315] on Gemma, removing 36% of toxic continuations and keeping 91% of clean ones, and 0.460 [0.418, 0.502] on Qwen. It stays above prompting and above twenty random directions through the same decision in both families, and on Qwen it is above the global shift (+0.057 [0.008, 0.105]) and level with refusing and with CAST. Its no-edit null again leaves every continuation unchanged.

M PTS IN THE SETTING OF PCHI

Setting. We follow Li et al. (2026): Qwen3-4B-Instruct-2507, the JSON self-evaluation prompt of their Appendix A.1 verbatim, and greedy decoding with at most 768 new tokens. Problems are the first 10,000 unique problems of the OpenMathInstruct-2 train_1M split, the first 5,000 for fitting and the rest held out; responses whose JSON does not parse are dropped, leaving 3,137 and 3,183. An answer is correct when it matches the reference after normalization, and the confidence readout is the yes-no logit gap at the `is_confident` value. On the held-out answers this gives 1,178 wrong-confident (WY) and 1,796 correct-confident (CY) readouts; PCHI’s own sample is larger and more accurate, so absolute rates differ from theirs. The decision is a calibrated diagonal-LDA probe on the layer-18 readout state (fitted on 70% of the fitting answers, calibrated on 30%; held-out AUROC 0.81), firing at probability 0.5. PCHI learns one coefficient for each of the 128 heads of layers 19 to 22 with its published objective and hyper-parameters (200 steps, batch 8, learning rate 0.04, L1 weight 0.05, ratio 0.7) and margin -1 . PTS shifts the layer-18 readout state along the difference of mean states of wrong answers stated without and with confidence. Doses were chosen on the fitting answers (Appendix A, v24).

arm (held-out, $n_{WY}=1,178$, $n_{CY}=1,796$)	WY corrected	CY damaged	Sel [95% CI]	ECE
unsteered	—	—	0.000	0.372
shift, every answer (dose 2)	0.099	0.007	0.092 [0.075, 0.110]	0.337
PCHI, learned heads	0.399	0.045	0.354 [0.325, 0.384]	0.234
PCHI, permuted coefficients (mean of 20)	—	—	0.033	0.363
PTS (probe-gated shift, dose 8)	0.582	0.143	0.438 [0.406, 0.471]	0.188
same probe, random direction (mean of 100)	—	—	0.317	0.254
same probe, override of the readout	0.582	0.143	0.438	0.238

Table 16: The setting of PCHI on held-out answers. All three registered tests pass: PTS above PCHI (+0.084 [0.058, 0.110]), PTS above the global shift, and PCHI above its permuted coefficients.

At dose 8 the PTS shift turns every flagged readout to *no*, so its selectivity equals that of the override; it differs in calibration, because the shift moves the stated probability rather than replacing it. A random direction at the same dose reaches the same selectivity in 46 of 100 draws, since an edit of that size at the readout position overwrites the answer, and its mean calibration error (0.254) stays above that of the PTS direction. The learned coefficients of PCHI matter for PCHI: permuting them across heads drops its selectivity to 0.033.

N LATENCY

Toxicity setting, Gemma, 750 confirmation prompts, 40 new tokens, batch 128, one idle H100, bfloat16: unsteered generation 5.79 s; generation with the gated edit 5.84 s; the forward pass over the prompts that reads the decision 0.46 s, for 6.31 s in total ($1.09\times$); post-hoc PTS, which generates once unsteered and regenerates the 51% of prompts its decision flags, $5.79 + 2.95 = 8.74$ s ($1.51\times$).

O CAPABILITY AWAY FROM THE TARGET BEHAVIOR

Table 17 measures code synthesis (HumanEval, 164 problems, greedy, pass@1), open-ended instruction following (200 AlpacaEval prompts, 384 new tokens; distinct-3 and the unsteered model’s perplexity of the continuation) and language modelling (300 WikiText-103 passages, teacher forced). Gated rows use the frozen Gemma decisions at the frozen threshold; each gated output is the unsteered output where the decision is silent and the ungated output where it fires. The coding and open-ended prompts are outside the distribution the probes were fitted on, and the decisions fire on some of them: the post-hoc probe on 23% of HumanEval, 42% of AlpacaEval and all WikiText passages, the prompt probe on 0%, 47% and 12%. Where it fires on open-ended text the shift costs distinct-3 and fluency as it does ungated.

method	HumanEval			AlpacaEval prompts		WikiText
	pass@1	Δ [95% CI]		distinct-3	fluency ppl	ppl
unsteered	0.396	+0.000	[+0.000, +0.000]	0.953	3.31	25.5
shift -0.375 (ungated)	0.384	-0.012	[-0.073, +0.049]	0.889	3.72	28.1
PTS post-hoc, shift -0.375	0.402	+0.006	[+0.000, +0.018]	0.909	3.49	28.1
PTS prompt, shift -0.375	0.396	+0.000	[+0.000, +0.000]	0.906	3.49	25.8
shift -0.5 (ungated)	0.281	-0.116	[-0.183, -0.049]	0.815	4.00	32.5
PTS post-hoc, shift -0.5	0.390	-0.006	[-0.030, +0.012]	0.865	3.55	32.5
PTS prompt, shift -0.5	0.396	+0.000	[+0.000, +0.000]	0.866	3.61	26.2

Table 17: Capability under the Gemma operators. Intervals are paired bootstrap intervals over problems.

A decision that abstains out of domain. The false firing on open-ended instructions is a property of the probe’s training data, not of the operator. We refit both probes with 600 AlpacaEval instructions and their unsteered answers added as negatives (instructions disjoint from the 200 used above) and set the threshold to fire on the same fraction of tuning questions. The refitted probes fire on none of the HumanEval and AlpacaEval prompts, so the gated outputs there equal the unsteered ones; on WikiText the prompt probe fires on none and the post-hoc probe on 48% of passages. On the confirmatory split their held-out AUROC is 0.735 (answer) and 0.719 (prompt), and composing them with the frozen actions gives selectivity 0.280 [0.242, 0.319] and 0.300 [0.261, 0.339], within 0.015 of the frozen probes. This analysis was not part of the protocol.

P EXAMPLES

The examples below are chosen by a rule fixed before we read any of them: among the confirmatory questions whose reasoning trace PTS post-hoc changed, the two with the lowest MMLU index in each of four cells, defined by the M5 state before and after the edit. The cells hold 474 (overconfident-wrong, state changed), 47 (overconfident-wrong, unchanged), 33 (confident-right, kept) and 277 (confident-right, lost) questions; confident-right questions the decision does not flag are not in these counts and keep their unsteered output. Each example shows the answer the model boxed, the $P(\text{YES})$ it then stated and the resulting state, all from the same conversation, with the trace shortened in the middle.

overconfident-wrong, state changed (mmlu-4, gold B).

Question: Find the product of the given polynomials in the given polynomial ring. $f(x) = 4x - 5$, $g(x) = 2x^2 - 4x + 2$ in $\mathbb{Z}_8[x]$. A. $2x^2 + 5$ B. $6x^2 + 4x + 6$ C. 0 D. $x^2 + 1$

before (answer C, $P(\text{YES}) = 0.78$, OCW): <think>The product of two polynomials in a polynomial ring is found by multiplying them together. We need to consider the remainders when each term is divided by 8. Let's multiply the polynomials and then reduce each term modulo 8.</think> $f(x) * g(x) = (4x - 5)(2x^2 - 4x + 2) = 8x^3 - 16x^2 + 8x - 10x^2 + 20x - 10 = 8x^3 - 26x^2 + 28x - 10$ <think>Now we need to reduce each term modulo 8.</think> $8x^3 - 26x^2 + 28x - 10 \equiv 0x^3 - 2x^2 + 0x - 2 \pmod{8}$ <think>The product of the polynomials is -2.</think> boxed{C}

[illegible]

overconfident-wrong, state changed (mmlu-23, gold D).

Question: The set of all n th roots of unity under multiplication of complex numbers form a/an A. semi group with identity B. commutative semigroups with identity C. group D. abelian group

before (answer C, $P(\text{YES}) = 0.98$, OCW): <think>The n th roots of unity are complex numbers that satisfy the equation $z^n = 1$. We need to check if this set satisfies the properties of a group. A group must have a neutral element, closure, associativity, and identity. </think> <think>The set of n th roots of unity satisfies the properties of a group. It has a neutral element ($z^n = 1$) and closure. </think> <think>The set of n th roots of unity is commutative. </think> <think>The set of n th roots of unity is abelian. </think> \boxed{C}

after (answer C, $P(\text{YES}) = 0.30$, NCW): <think>The set of all n th roots of unity is a set of complex numbers. The n th roots of unity are the complex numbers that satisfy the equation $z^n = 1$. Since the set is closed under multiplication, it is a semigroup. The identity element is 1. The set is also closed under multiplication. The set is commutative. The set is abelian. </think> \boxed{C}

overconfident-wrong, state unchanged (mmlu-57, gold B).

Question: Find the generator for the finite field $\mathbb{Z}_{11}$. A. 1 B. 2 C. 3 D. 4

before (answer A, $P(\text{YES}) = 1.00$, OCW): <think>The generator of a finite field is an element that generates all other elements in the field. $\mathbb{Z}_{11}$ is a finite field with 11 elements. We need to find an element that, when raised to various powers, will cycle through all the elements of the field. We can start by looking at the powers of 1.</think> <think>The powers of 1 in $\mathbb{Z}_{11}$ are: 1, $1^2 = 1$, $1^3 = 1$, ...</think> <think>This shows that 1 is a generator for $\mathbb{Z}_{11}$.</think> \boxed{A}

after (answer A, $P(\text{YES}) = 0.56$, OCW): <think>The generator of a finite field is an element that generates all the elements of the field. Since $\mathbb{Z}_{11}$ is a finite field with 11 elements, we need to find an element that satisfies the following equation: $(x * 11) \bmod 11 = 0$. <think>The generator can be found by calculating the remainder of $(x * 11) \bmod 11$. <think>The remainder of $(x * 11) \bmod 11$ is 0 for all x in $\mathbb{Z}_{11}$. <think>Therefore, the generator is 1. </think> \boxed{A}

overconfident-wrong, state unchanged (mmlu-1888, gold C).

Question: Suppose that we are interested in testing the null hypothesis that a GARCH(2,2) model can be restricted to a process with a constant conditional variance using the likelihood ratio test approach. Which of the following statements are true? A. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis B. The value of the log-likelihood function will almost always be bigger for the restricted model than for the unrestricted model C. If the relevant values of the log-likelihood functions are -112.3 and -118.4, the value of the test statistic is 12.2 D. The likelihood ratio test compares the slopes of the log-likelihood function at the maximum and at the restricted parameter value.

before (answer A, $P(\text{YES}) = 0.92$, OCW, no box: letter read from the logits): <think>The null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the model has a non-constant variance. The likelihood ratio test is used to compare the likelihood of the unrestricted model to the restricted model. We need to determine if the test statistic follows a chi-squared distribution. </think> **A. The test statistic will [...] the unrestricted and restricted models. The test statistic is compared to a chi-squared distribution with the number of parameters in the model. </think> **A**

after (answer A, $P(\text{YES}) = 0.53$, OCW): <think>The null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the conditional variance is not equal to the restricted value. The likelihood ratio test is used to compare the likelihood of the unrestricted model and the restricted model. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis. </think> \boxed{A}

confident-right, kept (mmlu-196, gold A).

Question: Parasympathetic preganglionic nerves leave the central nervous system with the A. third cranial nerves. B. fourth cranial nerves. C. fifth cranial nerves. D. sixth cranial nerves.

before (answer A, $P(\text{YES}) = 0.96$, CR): <think>The parasympathetic nervous system is responsible for "rest and digest" functions. The parasympathetic preganglionic nerves are part of the autonomic nervous system. The autonomic nervous system is divided into the sympathetic and parasympathetic branches. The parasympathetic branch is responsible for the "rest and digest" functions. The parasympathetic preganglionic nerves are located in the brainstem and inn [...] rves. The cranial nerves are numbered I-XII. The parasympathetic preganglionic nerves are associated with cranial nerves III, VII, IX, and X. </think> \boxed{A}

after (answer A, $P(\text{YES}) = 0.68$, CR): <think>The parasympathetic nervous system is

responsible for the "rest and digest" functions. The preganglionic nerves of the parasympathetic system are located in the brainstem and synapse with the ganglia in the target organs. The target organs are the same as the parasympathetic ganglia. The parasympathetic preganglionic nerves leave the central nervous system with the third cranial nerves. </think> \boxed{A}

confident-right, kept (mmlu-1430, gold A).

Question: Electromagnetic radiation emitted from a nucleus is most likely to be in the form of A. gamma rays B. microwaves C. ultraviolet radiation D. visible light

before (answer A, $P(\text{YES}) = 1.00$, CR): <think>The nucleus of an atom is a very small, dense region containing protons and neutrons. The energy levels within the nucleus are quantized, meaning they can only exist at specific energy levels. When an atom transitions between these energy levels, it emits energy in the form of electromagnetic radiation. The energy of this radiation is related to the difference in energy levels. Higher energy levels correspond [...] high energy and short wavelength. Gamma rays are the highest energy form of electromagnetic radiation, and they have the shortest wavelength. </think> \boxed{A}

after (answer A, $P(\text{YES}) = 0.78$, CR): <think>The nucleus is a very small and dense object. The energy of the nucleus is very high. The energy of the nucleus is emitted in the form of gamma rays. </think> \boxed{A}

confident-right, lost (mmlu-6, gold A).

Question: Statement 1 — Every homomorphic image of a group G is isomorphic to a factor group of G . Statement 2 — The homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . A. True, True B. False, False C. True, False D. False, True

before (answer A, $P(\text{YES}) = 0.85$, CR): <think> Statement 1 says that every homomorphic image of a group G is isomorphic to a factor group of G . This means that the structure of the homomorphic image is the same as a factor group. Statement 2 says that the homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . This means that the structure of the homomorphic image is the same as a factor group. </think> \boxed{A}

after (answer A, $P(\text{YES}) = 0.03$, NCR): <think> Statement 1 is true. The homomorphic image of a group G is isomorphic to a factor group of G . Statement 2 is also true. The homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . </think> \boxed{A}

confident-right, lost (mmlu-31, gold B).

Question: If $(G, \cdot)$ is a group such that $(ab)^{-1} = a^{-1}b^{-1}$, for all a, b in G , then G is a/an A. commutative semi group B. abelian group C. non-abelian group D. None of these

before (answer B, $P(\text{YES}) = 0.96$, CR): <think>The given condition $(ab)^{-1} = a^{-1}b^{-1}$ implies that the group G is closed under the operation. This means that the operation is associative. Also, the inverse of a product is the product of the inverses. This is a characteristic of abelian groups. </think> \boxed{B}

after (answer A, $P(\text{YES}) = 0.11$, NCW): <think> The given condition $(ab)^{-1} = a^{-1}b^{-1}$ implies that the group G is commutative. The commutative group is a group where $a*b = b*a$ for all a, b in G . </think> \boxed{A}

Q TRANSFER TO OTHER BENCHMARKS

We ran the Gemma operators, with the directions, probes and thresholds fitted on MMLU, on 300 ARC-Challenge questions (Clark et al., 2018) and 300 GSM8K questions (Cobbe et al., 2021) converted to four-option multiple choice (Appendix D). These runs are exploratory and not part of the protocol. Table 18 shows that the MMLU threshold does not carry to ARC: the post-hoc probe fires on 59 of 300 questions and PTS reaches selectivity 0.100, against 0.370 for the ungated additive dose. On GSM8K PTS keeps 0.944 of confident-right answers at selectivity 0.287, while the additive dose reaches 0.319 at retention 0.634.

R RESULTS UNDER THE PILOT PROTOCOL

A pilot protocol used the verification readout M4 (Appendix S) as its primary instrument, five evaluation subsets of 300 MMLU questions, a detector read off the extraction split, and configurations tuned on a 1,200-question split. We keep its results here as a record, with three qualifications that the confirmatory protocol removes. Under M4 the edited generations are the per-option verification

method	ARC-Challenge ($n_{OCW} = 66, n_{CR} = 205$)			GSM8K-MCQ ($n_{OCW} = 35, n_{CR} = 161$)		
	Δ_{acc}	retention	Sel	Δ_{acc}	retention	Sel
additive CAA (tuned)	+0.027	0.810	+0.370	+0.023	0.634	+0.319
CAST-style, trace condition	-0.003	0.863	+0.060	-0.060	0.516	+0.173
CAST, prompt condition	-0.017	0.844	+0.223	+0.023	0.634	+0.319
PTS, post-hoc	-0.010	0.888	+0.100	+0.023	0.944	+0.287
PTS, decides at the prompt	-0.013	0.941	-0.013	-0.040	0.814	+0.128

Table 18: Transfer of the Gemma operators fitted on MMLU, M5 readout, 300 questions each, no refitting.

traces, not the reasoning trace, so these results describe the verification task. Four of the five subsets share 7 to 14 questions with the extraction split (38 in total); the pooled numbers below drop them. The subsets also overlap each other in 54 questions, which are counted once.

Pooled comparison. Table 19 pools the five subsets (1,406 unique questions, 245 OCW and 380 CR under M4) and pairs every method with the tuned gated ablation of that protocol. No tuned baseline is above it at the 5% level, and none is below it either except the single-pass gate; the tuned additive dose has the highest mean (+0.271 against +0.203, paired difference +0.067, 95% CI [-0.024, +0.156]). Under M2 on the same traces every selectivity is below 0.11.

method (pilot protocol, tuned)	M4		M2	
	Sel [95% CI]	Δ vs gated abl.	Sel [95% CI]	Δ_{acc}
gated ablation (tuned)	+0.203 [+0.125, +0.282]	—	+0.042 [+0.006, +0.078]	-0.005
additive CAA (tuned)	+0.271 [+0.188, +0.350]	+0.067 [-0.024, +0.156]	+0.082 [+0.039, +0.127]	-0.009
directional ablation	+0.258 [+0.178, +0.338]	+0.054 [-0.017, +0.126]	+0.080 [+0.038, +0.124]	-0.004
CAST-style (tuned)	+0.221 [+0.146, +0.294]	+0.017 [-0.075, +0.108]	+0.101 [+0.060, +0.143]	-0.005
MiMiC (tuned)	+0.126 [+0.050, +0.203]	-0.077 [-0.180, +0.023]	+0.044 [+0.006, +0.082]	-0.006
Linear-AcT (tuned)	+0.146 [+0.082, +0.211]	-0.057 [-0.155, +0.039]	+0.032 [+0.001, +0.064]	-0.004
prompting	+0.196 [+0.118, +0.273]	-0.008 [-0.112, +0.096]	+0.053 [+0.011, +0.096]	-0.011
score-proportional dose (tuned)	+0.241 [+0.163, +0.320]	+0.038 [-0.026, +0.100]	+0.057 [+0.017, +0.099]	-0.008
sequential gate (tuned)	+0.079 [+0.021, +0.138]	-0.124 [-0.215, -0.031]	+0.024 [-0.004, +0.052]	-0.003
gated ablation (published)	+0.278 [+0.201, +0.356]	+0.075 [+0.005, +0.146]	+0.076 [+0.036, +0.117]	+0.004
additive -0.75 (published)	+0.000 [+0.000, +0.000]	-0.203 [-0.282, -0.125]	-0.025 [-0.078, +0.028]	-0.171

Table 19: Pilot protocol, five subsets pooled with overlapping questions removed. Intervals are question-level paired bootstrap intervals (10^4 resamples).

Pilot detector. The two-dimensional LDA gate of the pilot protocol reached AUROC 0.77–0.80 between OCW and CR on the 400 extraction traces it was fitted to, and 0.59 (layer 14) and 0.59 (layer 16) on the 1,406 pooled evaluation questions; the one-dimensional detection direction reached 0.66. This in-sample optimism is what motivated fitting the detector on a separate pool and reporting only held-out AUROC.

Scale, layers, geometry and cost. Table 20 repeats the published configurations on gemma-2-9b-it; Table 21 refits the axes and the gate at nine depths; Table 22 reports displacement, nearest-neighbour distance, the fraction of tokens outside the 99% conditional ellipsoid and next-token KL; Table 23 gives end-to-end wall clock on one H100 for 300 questions.

method (M4, 9B, 3 subsets)	Δ_{acc}	ECE	ocw_rm	cr_keep	AURC	Sel
baseline	acc 0.633 $\pm$ 0.031	0.144 $\pm$ 0.021	—	—	0.197 $\pm$ 0.025	—
additive -0.75	-0.184 $\pm$ 0.033	0.167 $\pm$ 0.037	1.000 $\pm$ 0.000	0.000 $\pm$ 0.000	0.450 $\pm$ 0.023	+0.000 $\pm$ 0.000
CAST-style gate+add	-0.114 $\pm$ 0.032	0.101 $\pm$ 0.011	0.831 $\pm$ 0.058	0.453 $\pm$ 0.038	0.285 $\pm$ 0.016	+0.283 $\pm$ 0.093
clamp q_{50} (ours)	+0.023 $\pm$ 0.012	0.143 $\pm$ 0.018	0.315 $\pm$ 0.012	0.887 $\pm$ 0.015	0.184 $\pm$ 0.025	+0.202 $\pm$ 0.020
gated ablation (ours)	+0.011 $\pm$ 0.008	0.127 $\pm$ 0.009	0.297 $\pm$ 0.026	0.907 $\pm$ 0.014	0.186 $\pm$ 0.023	+0.203 $\pm$ 0.034

Table 20: gemma-2-9b-it, layer 24, published configurations, mean $\pm$ standard deviation over three subsets of 300 questions.

S PROMPTS

All prompts are rendered with the model’s chat template. Text in braces is filled per question; everything else is verbatim.

Reasoning prompt (all readouts).

Given the following multiple choice question, choose the single best answer (A, B, C, or D).
Please think step-by-step and put your thinking inside

layer	Cohen's d conf.	Cohen's d OCW/CR	gate AUROC	2AUC-1
4	+0.525	-0.079	0.783	+0.566
7	+0.535	-0.093	0.756	+0.511
10	+0.512	-0.144	0.741	+0.482
12	+0.516	-0.150	0.742	+0.485
14	+0.589	-0.155	0.769	+0.538
16*	+0.631	-0.141	0.799	+0.599
18	+0.616	-0.116	0.779	+0.558
22	+0.675	-0.110	0.791	+0.582
25	+0.700	-0.111	0.649	+0.298

Table 21: Layer sweep on the 400 extraction traces (in-sample AUROC).

operator	MMLU				ARC				GSM8K			
	disp.	NN	out%	KL	disp.	NN	out%	KL	disp.	NN	out%	KL
unsteered	—	1.00	0.9	0.000	—	1.00	1.2	0.000	—	1.00	2.9	0.000
random direction, norm-matched	0.663	1.90	1.6	0.445	0.644	2.27	1.7	0.233	0.749	2.06	0.9	0.149
additive -0.75 Panickssery et al. (2024)	0.663	1.88	100.0	0.032	0.644	2.25	100.0	0.008	0.749	2.04	100.0	0.012
additive, tuned dose Panickssery et al. (2024)	0.221	1.22	69.9	0.005	0.215	1.33	40.3	0.001	—	—	—	—
MiMiC Singh et al. (2024)	0.082	1.06	1.3	0.001	0.074	1.08	1.8	0.000	0.097	1.06	4.5	0.001
Linear-AcT Rodriguez et al. (2025a)	0.012	1.01	1.0	0.000	0.011	1.01	1.5	0.000	0.012	1.01	3.5	0.000
clamp q_{50} (ours)	0.034	1.03	1.3	0.006	0.046	1.06	1.1	0.001	0.003	1.00	2.9	0.001
quantile transport (ours)	0.004	1.00	1.0	0.000	0.003	1.00	1.3	0.000	0.007	1.00	4.0	0.000
ablation Arditì et al. (2024)	0.144	1.15	29.5	0.010	0.178	1.28	28.4	0.001	0.062	1.05	10.1	0.001

Table 22: Activation diagnostics at layer 14 under teacher forcing of unsteered traces. NN is the nearest-neighbour distance to 4,096 natural activations relative to the unsteered value; out% is the share of tokens outside the 99% conditional ellipsoid of the edited coordinate given the leading 64 principal directions of the complement (1% expected); KL is the mean next-token divergence. These are measured distances, not a test of membership in a manifold.

<think>...</think> tokens.
 After </think>, give ONLY your final answer in the format:
 \boxed{[letter]}
 Question: {question}
 Options:
 A. {option A}
 B. {option B}
 C. {option C}
 D. {option D}

M5 (primary). After the model's reply, a second user turn is added to the same conversation and $P(\text{YES})$ is read from the next-token logits, summing the probabilities of the single-token variants of YES and of NO and renormalising over the two:

You answered {letter}. Is that answer correct? Reply with only YES or NO.

The letter is the one with the highest logit after the last \boxed{ of the trace, which is appended when the trace has none (6.1% of unsteered confirmatory traces). On the unsteered confirmatory split this letter equals the written letter on 96.9% of questions; the remainder are traces with several boxes or none. A question is confident iff $P(\text{YES}) \geq 0.5$.

M2 (secondary). The confidence of the same generated answer is the softmax over the four letter logits at the position after \boxed{.

M4 (secondary; primary in the pilot protocol). Each option is verified in a separate generation, with at most 192 new tokens:

Answer whether the given answer is the right answer for the following question.
 Please think step-by-step and put your thinking inside
 <think>...</think> tokens.
 After </think>, answer ONLY with "YES" or "NO".
 Question:
 {question}

method	passes	seconds	wall clock
unsteered	1	43.5	1.00×
additive −0.75 Panickssery et al. (2024)	1	45.2	1.04×
ablation Arditì et al. (2024)	1	44.6	1.02×
clamp q_{50} (ours)	1	41.4	0.95×
quantile transport (ours)	1	44.8	1.03×
MiMiC Singh et al. (2024)	1	42.9	0.98×
Linear-AcT Rodriguez et al. (2025a)	1	40.9	0.94×
gated ablation, post-hoc gate (ours) (223/300 regenerated)	2	85.8	1.97×
gated ablation, sequential gate (ours)	1	45.5	1.05×

Table 23: End-to-end wall clock for 300 MMLU questions, batch 64, one H100, one run per method.

Answer:
{option text}

$P(\text{YES})$ is read at the answer word after `</think>`; the predicted answer is the option with the largest $P(\text{YES})$ and its confidence is that value divided by the sum over options. When an M4 condition is steered, the hook is active in these four verification generations, not in the reasoning trace, so an M4 result describes the verification task.

Prompting baseline. The instruction below is prepended to the user turn (Gemma has no system role) or passed as the system message (Qwen):

Be careful and well-calibrated. Only answer confidently
when you are sure; if you are uncertain, explicitly
acknowledge the uncertainty in your reasoning and
confidence.

Weights. Gemma-2-2B weights were served from the `unsloth/gemma-2-2b-it` mirror, whose 288 tensors are bit-identical to `google/gemma-2-2b-it`; Qwen weights are Qwen/Qwen2.5-7B-Instruct. Decoding is greedy with at most 320 new tokens for the reasoning trace.