

PROJECTION-TRANSPORT STEERING: AN EXACT ACCOUNT OF SELECTIVE CONDITIONAL STEERING

Anonymous authors

Paper under double-blind review

ABSTRACT

Activation steering edits a language model’s residual stream at inference time, but most edits act on every input, so an edit strong enough to remove a behavior also changes inputs that never showed it. Conditional methods such as CAST and DSAS gate the edit, yet their evaluations cannot say how much of the resulting selectivity comes from the gate and how much from the edit. We propose Projection-Transport Steering (PTS), which composes a decision, a likelihood-ratio test on a learned statistic, with an action on a low-dimensional behavioral projection, and accounts exactly for the pair. For any gated edit that leaves unflagged outputs unchanged, selectivity equals the decision’s true-positive rate times the action’s conversion rate on flagged targets, minus its false-positive rate times the conversion rate on flagged protected outputs. The identity predicts gated selectivity from one ungated run ($r = 0.97$ and 0.91), bounds a decision from unsteered answers alone, and fixes three reference arms that share the decision. For the action, optimal transport gives the unique minimum-displacement edit toward a chosen target law. On overconfident multiple-choice answers of Gemma-2-2B, PTS is the most selective of eight methods on 3,000 held-out questions and cuts calibration error from 0.354 to 0.255 at unchanged accuracy. On free-text toxicity in three model families and two held-out prompt sets, PTS beats refusal on its own decision and a hundred random directions, and it is level with a tuned DSAS-style per-token gate in selectivity while keeping 80–85% of clean continuations against 67% on two of the three families.

1 INTRODUCTION

Many high-level behaviors of large language models appear to be carried by directions in their activations. Activation steering methods (Turner et al., 2023; Panickssery et al., 2024; Zou et al., 2023; Li et al., 2023) read such a direction off contrastive examples and add it to the residual stream at inference time, at the cost of one vector addition per token and no retraining. However, these methods apply the same edit to every input, so the strength needed to remove a behavior where it occurs also changes the inputs where it does not: a refusal vector at a useful strength raises refusal on harmful and harmless prompts alike (Arditi et al., 2024).

The difficulty is sharpest for a behavior that a model shows on only some inputs, such as overconfidence in multiple-choice reasoning, where the answers to change and the answers to protect are both confident, one wrong and one right, and the two states overlap on the model’s confidence direction. On Gemma-2-2B, our main model, a global shift along it strong enough to remove every overconfident-wrong answer also removes every confident-right one (Figure 1a).

In this paper we introduce Projection-Transport Steering (PTS), which splits a steering operator into an *action* that sets how to edit and a *decision* that sets which inputs to edit. The action is a transfer function on the projection of the residual stream onto a behavioral subspace, a few directions read off the model’s own labeled outputs. The decision is a likelihood-ratio test on a detection statistic fitted to tell the target state from the protected one, and the action runs only where the test fires. Our method is based on the observation that this statistic, read from the full activation, separates the

method	transfer function $g(p)$	decision	unflagged	account of Sel
ActAdd/CAA, ITI, RepE	$p + \alpha, p + \alpha(x)$	none	edited	—
ablation, SAE clamp	$0; c$	none	edited	—
MiMiC	full-space affine map	none	edited	—
Linear-AcT	$(1-\lambda)p + \lambda(\omega p + \beta)$	none	edited	—
CAST	$p + \alpha$	prompt similarity	unchanged	—
DSAS	$(1-\gamma_t)p + \gamma_t T(p)$, any T	per-token gate γ_t	edited by γ_t	—
PCHI	learned head rescaling	probe at readout	unchanged	—
PTS (ours)	$S(p)$, transport map to ν	likelihood-ratio test	unchanged	exact, Prop. 1

Table 1: Steering operators differ in their transfer function g on a k -dimensional projection p , in whether a decision gates it, and in what happens to the inputs it does not flag. PTS derives g from a target law ν (Theorem 1); its selectivity account (Proposition 1) covers the CAST, DSAS and PCHI rows too.

two confident states better than the coordinate the action moves (Sec. 5). Table 1 places published operators in this frame.

Separating the decision from the action is not new in itself: CAST (Lee et al., 2025) gates a vector on the prompt, and DSAS (Ferrando et al., 2026) learns a per-token gate that interpolates between the original and the steered activation of any steering map. These methods are judged by aggregate trade-off curves, so a reader cannot tell whether a gain comes from the gate, from the edit, or from regenerating the input at all. PTS supplies that account, and the account applies to CAST and DSAS as much as to PTS. It also exposes a design choice: a per-token gate edits every token a little, while a decision on the input leaves unflagged outputs untouched.

The selectivity of a conditional edit, the share of target outputs it removes minus the share of protected outputs it loses, equals the decision’s true-positive rate times the action’s conversion rate on flagged targets, minus its false-positive rate times the conversion rate on flagged protected outputs (Proposition 1). The identity predicts gated selectivity from one ungated run, bounds a decision from unsteered answers (Theorem 2), and states when steering beats overwriting every flagged output (Eq. (3)).

Our main contributions are:

- An exact decomposition of conditional-steering selectivity into decision and action terms, with a ceiling for any threshold decision and the reference arms it implies (Sec. 3). Over 240 tuning-split configurations per model it predicts gated selectivity with correlations of 0.97 and 0.91 (Sec. 5.2).
- For a subspace and a target law, the transport map is the unique minimum-displacement projection-local edit (Theorem 1).
- Held-out evidence on overconfident reasoning, with every configuration committed before the held-out run: on 3,000 MMLU questions for the main model, PTS is the most selective of eight methods, and its direction is more selective than each of a hundred random directions with the same decision and dose (Sec. 5).
- Free-text control of toxic continuation in three model families, where the PTS operator chosen by one fixed rule beats refusal on its own decision in every family and on two disjoint prompt sets. Against a tuned DSAS-style per-token gate it is level in selectivity and keeps 85% and 80% of clean continuations on Gemma and Qwen, against 67% (Sec. 6).
- An attribution of each gain with the reference arms. The edit direction carries the gain on Gemma and on free text, the decision carries it on Qwen and in the setting of PCHI, where PCHI’s own probe with a shift is more selective than its learned head rescaling, and a shift is as selective as the transport map on fresh prompts (Sec. 6).

2 RELATED WORK

Activation steering. ActAdd (Turner et al., 2023) adds the activation difference of one contrast prompt pair to the residual stream, and CAA (Panickssery et al., 2024) averages it over many pairs. ITI (Li et al., 2023) shifts attention heads selected by linear probes, RepE (Zou et al., 2023) reads and writes representation directions, directional ablation (Arditi et al., 2024) removes one direction from every activation, and SAE clamping (Templeton et al., 2024) fixes one sparse feature. AxBench (Wu




Figure 1: PTS removes more of the target behavior than the global shift at the same retention. Held-out Gemma results; up and right is better, dotted lines mark equal selectivity. Orange: global shift at increasing dose; gray: random directions through the PTS decision; star: main PTS arm (post-hoc shift in a, tuning-rule choice in b).

et al., 2025) finds prompting stronger than such edits, and reliability studies (Tan et al., 2024; Braun et al., 2025) report that their effect varies widely across inputs. These operators act on every input with a transfer function chosen in advance; PTS derives the transfer function from a target law (Theorem 1).

Transport maps. MiMiC (Singh et al., 2024) applies to the full activation the affine map that matches the mean and covariance of two activation distributions, the Gaussian optimal-transport map. Linear-Act (Rodriguez et al., 2025a) fits a univariate affine transport map to each neuron, interpolated by a strength $\lambda \in [0, 1]$, and writes earlier edits as its special cases; LinEAS (Rodriguez et al., 2025b) learns such maps across layers end to end. COAST (Nguyen et al., 2026) moves each activation, at fixed norm, to a fixed cosine with a target direction at the smallest second-moment-weighted displacement, which in our terms transports that cosine onto a point mass. A-LQR (Skifstad et al., 2026) casts steering as linear-quadratic control. Unlike these methods, PTS transports only a k -dimensional projection and leaves its orthogonal complement fixed, so the cost-minimal edit toward a target law is optimal transport in k dimensions, with a closed form for $k = 1$.

Deciding when to act. Activations predict whether an answer is right (Kadavath et al., 2022; Azaria & Mitchell, 2023), and Zhao et al. (2026) trace inflated verbalized confidence to components that write at late token positions. The first generated tokens predict whether steering will succeed (Fan et al., 2026), while probes of a finished behavior forecast it poorly from a prefix (Kortukov et al., 2026). Within one forward pass, a steering write can degrade a direction read after it (Luo et al., 2026). SADI (Wang et al., 2025), dynamic activation composition (Scalena et al., 2024) and GAPS (Fereidouni et al., 2026) adapt the edit per input, per generation step and per dimension, respectively. Closer to our work, three methods separate the decision from the action. CAST (Lee et al., 2025) compares the prompt’s hidden state with a condition vector at a grid-searched threshold and adds a behavior vector when the comparison holds. DSAS (Ferrando et al., 2026) decouples when to steer from how, with a logistic gate that interpolates per token between the original and the steered activation. PCHI (Li et al., 2026) conditions a learned rescaling of downstream attention heads on a frozen probe read at a confidence-template token, and we run PTS in its setting (Sec. 6). To these methods we add an account of the separation: an identity that splits selectivity into decision and action terms (Prop. 1), a ceiling known before any steered rollout (Theorem 2), and reference arms that share the decision, which we run for every conditional method we compare.

3 PROJECTION-TRANSPORT STEERING

To steer a behavior only where it occurs, an operator has to settle what edit to apply and to which inputs. PTS answers the first question with a transport map on a behavioral subspace and the second

with a likelihood-ratio test on a detection statistic, and it accounts exactly for how the two compose. Proofs are in Appendix B.

The action. Let $h \in \mathbb{R}^d$ be the residual-stream activation at a fixed decoder layer, and let $V \in \mathbb{R}^{k \times d}$ have orthonormal rows spanning what we call the behavioral subspace; $q = Vh$ is the projection. In our experiments $k = 1$, and the row of V is the unit difference of mean activations between confident and hesitant answers, which we call the confidence direction, with q the confidence coordinate; on free text the direction separates toxic from clean continuations. We write μ for the law of h under unsteered generation, μ_V for the induced law of q , and ν for a target law on $\mathbb{R}^k$. Our goal is an edit that gives the projection the law ν , leaves the orthogonal complement untouched, and moves h as little as possible, since large edits degrade generation (Sec. 5).

Definition 1 (admissible steering). $T : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is *admissible* for (V, ν) if (i) $(I - V^\top V) T(h) = (I - V^\top V) h$ μ -a.s. (locality: the orthogonal complement is untouched), and (ii) $(V \circ T)_\# \mu = \nu$ (goal: steered projections follow the target law). Its cost is $C(T) = \mathbb{E}_\mu \|T(h) - h\|^2$.

Theorem 1 (steering is optimal transport of the projection). Let μ_V be absolutely continuous with finite second moment and let ν have a finite second moment. Then $\min_T C(T) = W_2^2(\mu_V, \nu)$ over admissible T , attained by

$$T^*(h) = h + V^\top (S(Vh) - Vh), \quad (1)$$

where S is the $(\mu_V$ -a.s. unique) Brenier map from μ_V to ν . Every optimal T satisfies $VT(h) = S(Vh)$, hence $T = T^*$ by locality, μ -almost surely.

In words, once the subspace and the target law are fixed, the edit is determined: it moves only the projection, by the transfer function S (g in Table 1). For $k = 1$, $S = F_\nu^{-1} \circ F_{\mu_V}$ is the monotone rearrangement, which we store as a 41-knot piecewise-linear lookup; here it suffices that μ_V has no atoms, and ν may have atoms. Directional ablation is then the map onto δ_0 , a clamp at c is the map onto the law of $\min(q, c)$, and the shift $S(q) = q + \alpha$ is the map onto a translate of μ_V ; for Gaussian marginals S is affine (Knott & Smith, 1984; Gelbrich, 1990). Since every nondecreasing transfer function is the optimal map onto its own image law, Theorem 1 turns the choice of edit into a choice of target law. Which law works best is an empirical question, which we answer by selection on a tuning split (Sec. 4). For the shift, the theorem gives the following corollary.

Corollary 1 (when additive steering is optimal, and its limit). (a) $S(q) = q + \alpha$ is the optimal map iff ν is the translate of μ_V by α . (b) Let two subpopulations have the same law of q . If, under an edit, the probability that an input leaves its state depends on h only through q , the edit changes the state of both at the same rate.

The decision. A decision is needed because the two states we must tell apart look alike to the action. Overconfident-wrong (OCW) and confident-right (CR) answers are both confident, and their laws on the confidence coordinate overlap (Sec. 5), so Corollary 1(b) suggests that a transfer function of that coordinate changes both states at comparable rates, so the global shift removes confident-right answers along with overconfident-wrong ones (Fig. 1a). The decision therefore reads a second statistic s of the input and runs the action on an event Γ determined by s , in practice $\{s > \tau\}$ for a scalar score; off Γ the output is exactly the unsteered one. We measure the result with selectivity, $\text{Sel} = P(\text{leave OCW} \mid \text{OCW}) - P(\text{leave CR} \mid \text{CR})$, the share of OCW answers removed minus the share of CR answers lost; on free text, toxic and clean continuations take their places. We write TPR and FPR for the rates at which the decision fires on OCW and CR, and ρ_O, ρ_C for the rates at which the action changes the state of a flagged OCW and a flagged CR answer.

Proposition 1 (selectivity decomposes). For any decision and action that leave unflagged outputs unchanged,

$$\text{Sel} = \text{TPR} \cdot \rho_O - \text{FPR} \cdot \rho_C. \quad (2)$$

The proof is one line of total probability. The decision sets (TPR, FPR) and the action sets (ρ_O, ρ_C) on the flagged answers. A gated operator with a post-hoc or prompt decision returns, input by input, either the unsteered or the ungated output, so one ungated run per action composes every such configuration; the ungated rates also predict its selectivity when the action converts flagged and unflagged answers of a state alike (Sec. 5.2).

The identity also fixes three reference arms on the same decision. A no-edit arm measures what regeneration alone changes, and a random direction at the same dose measures what any perturbation

of that size changes. An override of the reported confidence of every flagged answer ($\rho_O = \rho_C = 1$; on free text, a refusal) has selectivity $\text{TPR} - \text{FPR}$, and the action is more selective than this override exactly when

$$\text{TPR}(1 - \rho_O) < \text{FPR}(1 - \rho_C), \quad (3)$$

that is, when the share of protected answers flagged but spared exceeds the share of target answers flagged but not converted.

Theorem 2 (optimal decision and selectivity ceiling). *Let p_O, p_C be the densities of s on OCW and CR answers, and $r_O(s), r_C(s)$ the probabilities that the action converts such answers at s . (a) If r_O and r_C do not depend on s , then among decisions with equal FPR the likelihood-ratio test on s maximizes Sel (Neyman–Pearson); under equal-covariance Gaussian class conditionals it is the linear discriminant. In general, among decisions with equal $\text{FPR} \cdot \rho_C$, the maximizer thresholds $p_O(s) r_O(s) / (p_C(s) r_C(s))$. (b) If s is scalar and the concave hull of its ROC has area AUC, every decision $\{s > \tau\}$ has $\text{TPR} - \text{FPR} \leq 2 \text{AUC} - 1$, which caps the override on it and any action with $\rho_O = \rho_C$.*

Part (b) bounds, from unsteered answers alone, what a threshold decision contributes when the action converts both states alike, so the ceiling of a decision is known before any steered rollout; Appendix B treats the general case in which conversion depends on s . In the multiple-choice and free-text experiments the decision is a logistic-regression probe, which estimates the log-likelihood ratio up to a constant that τ absorbs. The probe reads the activations of the finished answer, of its first t tokens or of the prompt, and we fit it on a labeled pool disjoint from every tuning and evaluation split. Once the decision is fixed, Theorem 1 carries over to the inputs it flags.

Theorem 3 (transport under a fixed decision). *Let $P(\Gamma) > 0$, let $\mu_V(\cdot | \Gamma)$ be absolutely continuous, and let ν_Γ have a finite second moment. Among interventions that are local to V , the identity off Γ , and whose projections on Γ follow ν_Γ , the cost-minimal one applies on Γ the Brenier map from $\mu_V(\cdot | \Gamma)$ to ν_Γ , at cost $P(\Gamma) W_2^2(\mu_V(\cdot | \Gamma), \nu_\Gamma)$.*

Proposition 4 replaces the switch by a graded dose, the form of the per-token interpolation of DSAS, and shows that under a linear gain and a price on displacement the threshold gate is optimal exactly when the optimal dose is almost surely 0 or 1. A post-hoc decision scores the finished answer and regenerates the inputs it flags; Proposition 2 states when the decision and the action share one forward pass.

Proposition 2 (single-pass decision). *Let z_t be computed from the activations of the prompt, or of the prompt and the first t generated tokens. Among decisions measurable with respect to z_t , Theorem 2 holds with z_t in place of s , so the ceiling of an early decision is read from the ROC of z_t on unsteered answers. If the action starts after the last position z_t reads, the operator runs in one forward pass: the decision reads only unedited activations, earlier positions keep their cached keys and values, and for $k = 1$ each token costs at most two extra d -dimensional dot products, one lookup and one vector update.*

4 EXPERIMENTAL SETUP

The main model is Gemma-2-2B-it (Gemma Team, 2024) and the main behavior is overconfidence on MMLU (Hendrycks et al., 2021). The multiple-choice comparison is repeated on Qwen2.5-7B-Instruct, a second model family on which every component is refitted, and on ARC-Challenge (Clark et al., 2018). On multiple-choice questions both models reason inside `<think>...</think>` and answer `\boxed{letter}` under greedy decoding. Section 6 adds toxic continuation as a free-text behavior, with Mistral-7B-Instruct-v0.3 as a third family, and the setting of PCHI (Li et al., 2026), the closest conditional method. In both multiple-choice models the action runs at layer 14 (of 26 and 28 layers), where the confidence direction separates confident from hesitant answers of the tuning split within 0.016 AUROC of the best layer (Appendix L).

The MMLU splits are disjoint subsets of its test set: 400 questions to extract the directions and transport targets, a pool of 4,000 to fit the detection probe, 1,200 on which every configuration is tuned, and 3,000 used once for confirmation. On ARC every decision is refitted with 1,012 questions from the ARC training and validation sets added to its pool, 400 further questions from those sets tune the configurations, and the 1,165 test questions confirm them.

The primary readout (M5) scores the answer the model writes. Correctness compares the boxed letter with the gold answer. Confidence is $P(\text{YES})$ when the model is then asked, in the same conversation, whether that answer is correct, and a steered question is steered in both turns. Correctness and whether $P(\text{YES}) \geq 0.5$ define four states; the two confident ones are overconfident-wrong (OCW, wrong) and confident-right (CR, right). Each question is paired with itself before and after the edit. Removal is $P(\text{not OCW} \mid \text{OCW})$, retention is $P(\text{CR} \mid \text{CR})$, and Sel = removal $- (1 - \text{retention})$, which is 1 for an edit that removes every OCW and keeps every CR and 0 for one that treats the two states alike. Because selectivity conditions on the starting state, we also report, over all questions, the change in the share of OCW answers, the 10-bin expected calibration error (ECE) of the stated confidence, and the change in accuracy (secondary readouts and prompts in Appendices E and O).

PTS decides post-hoc, from the finished unsteered answer, and then regenerates the flagged questions; at the prompt, in one forward pass over it before generation; or at a prefix of $t \in \{16, 32, 64, 128\}$ generated tokens, which keeps decision and edit in one pass. Its action moves only the projection. On multiple-choice questions it is a shift at one of several doses or an ablation, both transport maps of Theorem 1; on free text it can also be the displacement of the quantile map onto the projection law of clean continuations, scaled by $\lambda \in \{1, 2, 3, 4, 6\}$, which at $\lambda = 1$ is the map itself.

We compare PTS with seven baselines: prompting with a calibration instruction, additive CAA (Panickssery et al., 2024), directional ablation (Arditi et al., 2024), MiMiC (Singh et al., 2024), Linear-Act (Rodriguez et al., 2025a), CAST with its prompt condition (Lee et al., 2025) fitted on the same labeled pool as our probe, and a CAST-style trace condition, a shift conditioned on the finished trace that, like post-hoc PTS, regenerates the questions it flags. On free text the baselines are prompting, the global shift and CAST. All methods are reimplemented on our hooks and scored with the same readout. One rule selects every configuration on the tuning split: the highest selectivity among configurations whose accuracy change is at least -0.01 , or the highest overall when none meets that floor; every confirmatory number is a real run. Budgets range from one configuration for prompting and ablation to sixteen for each conditional method, ours and both CAST variants alike (Appendix C).

Every conditional arm is run against three references that share its decision: a no-edit arm, which regenerates the flagged inputs with the identity action; random unit directions at the same dose (ten per conditional multiple-choice arm; one hundred for PTS, or thirty displaced as far as the ablation selected on the second family); and an override that marks every flagged answer as not confident or, on free text, replaces every flagged continuation by a refusal. Intervals are 95% paired bootstrap intervals over questions (prompts on free text) with 10^4 resamples. The seven comparisons of PTS with the baselines are two-sided and Holm-adjusted within each setting, each later protocol Holm-adjusts its own tests, and an accuracy change is non-inferior when the lower end of its interval is above -0.02 . Every protocol was committed before its held-out rollouts, the multiple-choice one with the hashes of its selected configurations and fitted detectors (Appendix A).

5 SELECTIVE REMOVAL OF OVERCONFIDENT ANSWERS

The PTS probe raises the selectivity ceiling of a threshold decision (Theorem 2b) above that of the confidence coordinate on both models. On the held-out questions the two confident states overlap on the coordinate, with AUROC 0.635 on Gemma and 0.53 on Qwen, in line with the frontier of the global shift in Fig. 1a. The best threshold on the coordinate separates them by $\max(\text{TPR} - \text{FPR}) = 0.216$ and 0.072. With the PTS probe, fitted on a separate pool of 4,000 labeled answers, the ceiling is 0.370 and 0.449 (AUROC 0.741 and 0.785; Fig. 2b).

5.1 PTS ON HELD-OUT QUESTIONS

Post-hoc PTS is the most selective of the eight methods on the 3,000 held-out questions (Table 2). It removes 49% of overconfident-wrong answers and keeps 80% of confident-right ones, a selectivity of 0.294 [0.256, 0.334]. The closest baseline, the CAST-style trace condition, reaches 0.248, and the other six trail PTS by 0.074 to 0.233; all seven differences survive Holm adjustment. The lead over the CAST-style condition also holds when the confidence threshold that defines the states moves from 0.5 to 0.4 (+0.044 [0.007, 0.080]) or to 0.6 (+0.068 [0.028, 0.107]).




Figure 2: Proposition 1 predicts gated selectivity, and the PTS probe raises the ceiling. (a) 480 tuning-split configurations against their prediction; (b) ceiling $\max(\text{TPR} - \text{FPR})$ from unsteered held-out answers, next to held-out PTS selectivity; (c) held-out AUROC of the probe reading the prompt, a prefix or the whole answer.

method	Δacc	removal	retention	Sel $\uparrow$ [95% CI]	$\Delta\text{Sel} (p)$	$\Delta\text{OCW} (\text{pp}) \downarrow$	$\Delta\text{ECE} \downarrow$
prompting	+0.005	0.42	0.78	0.200 [0.16, 0.24]	-0.094 (<0.001)	-5.9*	-0.030*
additive CAA	-0.001	0.56	0.66	0.220 [0.18, 0.26]	-0.074 (0.008)	-13.0*	-0.153*
dir. ablation	+0.009	0.46	0.72	0.173 [0.13, 0.21]	-0.121 (<0.001)	-8.7*	-0.111*
MiMiC	+0.000	0.31	0.80	0.117 [0.08, 0.15]	-0.177 (<0.001)	-2.1*	-0.018*
Linear-AcT	+0.005	0.14	0.92	0.061 [0.04, 0.09]	-0.233 (<0.001)	-0.2	-0.003
CAST (prompt)	-0.002	0.46	0.76	0.217 [0.18, 0.26]	-0.077 (0.004)	-11.2*	-0.137*
CAST-style (trace)	-0.005	0.54	0.70	0.248 [0.21, 0.29]	-0.046 (0.020)	-16.6*	-0.094*
PTS, post-hoc	+0.003	0.49	0.80	0.294 [0.26, 0.33]	—	-15.0*	-0.099*
PTS, prefix	-0.008	0.56	0.73	0.294 [0.26, 0.33]	+0.000	-17.2*	-0.136*
PTS, prompt	-0.012 [†]	0.55	0.76	0.306 [0.27, 0.35]	+0.012	-17.6*	-0.084*
same decision, no edit	+0.003	—	—	0.049 [0.03, 0.07]	-0.245 (<0.001)	—	-0.005
same decision, random direction	—	—	—	0.131	-0.163 (<0.001)	—	—
same decision, override readout	+0.000	—	—	0.319 [0.28, 0.36]	+0.025	—	—

Table 2: PTS is the most selective of eight methods on held-out questions. Gemma-2-2B, 3,000 MMLU questions (unsteered accuracy 0.559, ECE 0.354; $n_{OCW}=963$, $n_{CR}=1,399$), each method at the configuration selected on the tuning split; ΔSel is the row minus PTS post-hoc (Holm-adjusted p). The middle block decides before the answer is finished; the last rows keep the PTS decision with no edit, a random direction (mean of 100) and an override (Eq. (3)). *: interval excludes 0; [†]: accuracy interval reaches -0.02.

Over all questions, PTS lowers the share of overconfident-wrong answers by 15.0 points [13.7, 16.4] and the calibration error from 0.354 to 0.255, with a non-inferior accuracy change of +0.003 [-0.009, 0.015]. PTS thus corrects what the model says about its answer and leaves the answer as it was, which is what calibration asks for.

The decision can also be taken before the answer is finished. The probe separates the two states from the prompt almost as well as from the finished answer (Fig. 2c), and a decision read from the first 32 generated tokens, in the same pass as the edit, gives the same selectivity, 0.294.

The edit along the confidence direction carries this selectivity. Under the PTS decision, regenerating the flagged questions without an edit gives 0.049, and one hundred random unit directions at the PTS dose average 0.131; PTS lies above all hundred and above their mean by 0.163 [0.140, 0.188] (gray points in Fig. 1a). On code synthesis the prompt decision keeps the edit off: it fires on none of the 164 HumanEval problems, so pass@1 stays at its unsteered 0.396, whereas the same shift applied to every problem lowers it to 0.281 (Appendix M).

5.2 GATED SELECTIVITY FROM ONE UNGATED RUN

Proposition 1 predicts the selectivity of every decision placed in front of an action from one ungated run of that action. On the tuning split we measured the conversion rates of six actions once, composed each with two decisions at 20 thresholds, and compared the prediction $\text{TPR}_{\rho_O} - \text{FPR}_{\rho_C}$ with the gated selectivity composed from the same run (Fig. 2a). The correlation is 0.97 over 240 configurations on Gemma and 0.91 over 240 on Qwen. The composition that scores gated configurations during selection also matches real gated runs: on the three held-out splits, the twelve real

method	Qwen2.5-7B, MMLU ($n_{OCW}=563$, $n_{CR}=2,007$)				Gemma-2-2B, ARC ($n_{OCW}=228$, $n_{CR}=780$)			
	Δ_{acc}	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	$\Delta ECE \downarrow$	Δ_{acc}	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	$\Delta ECE \downarrow$
prompting	-0.003	0.212 [0.17, 0.25]	+0.040 (0.313)	-0.004	+0.003	0.238 [0.17, 0.31]	+0.003 (1.000)	+0.009
additive CAA	-0.015 [†]	0.151 [0.12, 0.19]	-0.021 (0.998)	+0.012	-0.001 [†]	0.294 [0.22, 0.37]	+0.059 (0.356)	-0.069*
dir. ablation	-0.009 [†]	0.156 [0.12, 0.19]	-0.016 (0.998)	+0.005	-0.008 [†]	0.215 [0.14, 0.29]	-0.020 (1.000)	-0.055*
MiMiC	-0.002	0.154 [0.12, 0.19]	-0.018 (0.998)	+0.020*	+0.010	0.198 [0.13, 0.26]	-0.037 (1.000)	-0.007
Linear-AcT	-0.003	0.123 [0.09, 0.16]	-0.049 (0.133)	+0.009	+0.003	0.109 [0.06, 0.16]	-0.126 (0.008)	+0.000
CAST (prompt)	-0.011 [†]	0.122 [0.09, 0.16]	-0.050 (0.133)	+0.027*	-0.003	0.109 [0.05, 0.17]	-0.126 (0.003)	-0.028*
CAST-style (trace)	-0.007	0.040 [0.02, 0.06]	-0.132 (<0.001)	+0.013*	+0.001	0.289 [0.22, 0.36]	+0.054 (0.356)	-0.061*
PTS, post-hoc	-0.002	0.172 [0.14, 0.21]	—	-0.000	-0.006 [†]	0.235 [0.17, 0.30]	—	-0.060*
PTS, prefix	-0.008	0.109 [0.08, 0.14]	-0.063	+0.013*	-0.008 [†]	0.276 [0.21, 0.35]	+0.041	+0.000
PTS, prompt	-0.034 [†]	0.120 [0.08, 0.16]	-0.052	+0.049*	-0.027 [†]	0.182 [0.11, 0.25]	-0.053	+0.028*
same decision, no edit	-0.002	0.091 [0.06, 0.12]	-0.081 (<0.001)	+0.000	-0.004	0.016 [-0.01, 0.05]	-0.219 (<0.001)	+0.011
same decision, random direction	—	0.158	-0.014 (0.175)	—	—	0.167	-0.068 (0.004)	—
same decision, override readout	+0.000	0.391 [0.36, 0.43]	+0.218	—	+0.000	0.270 [0.20, 0.34]	+0.035	—

Table 3: PTS is above its no-edit reference arm on a second model family (left) and a second benchmark (right). Notation as in Table 2; unsteered accuracy 0.737 (Qwen) and 0.744 (ARC); random directions: 100 on ARC, 30 displaced as far as the ablation on Qwen.

method	Gemma-2-2B (toxic 15.0%)			Qwen2.5-7B (toxic 20.3%)			Mistral-7B (toxic 3.4%)		
	rem./kept	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	rem./kept	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)	rem./kept	Sel $\uparrow$ [95% CI]	Δ_{Sel} (p)
prompting	42/68	0.104 [0.06, 0.15]	-0.422 (<0.001)	59/79	0.386 [0.34, 0.43]	-0.148 (<0.001)	72/70	0.419 [0.33, 0.51]	-0.101 (0.084)
global shift	90/58	0.474 [0.44, 0.51]	-0.051 (0.010)	84/57	0.403 [0.37, 0.44]	-0.131 (<0.001)	71/72	0.429 [0.34, 0.52]	-0.091 (0.134)
CAST	81/70	0.515 [0.47, 0.56]	-0.011 (0.585)	74/70	0.443 [0.40, 0.48]	-0.091 (<0.001)	64/81	0.446 [0.35, 0.54]	-0.074 (0.220)
PTS, shift	73/80	0.526 [0.48, 0.57]	—	82/72	0.534 [0.50, 0.57]	—	74/78	0.520 [0.43, 0.61]	—
PTS, scaled transport	79/79	0.584 [0.54, 0.62]	+0.058 (<0.001)	—	—	—	—	—	—
PTS, post-hoc	91/81	0.723 [0.69, 0.75]	+0.198 (<0.001)	76/89	0.648 [0.61, 0.68]	+0.114 (<0.001)	88/83	0.712 [0.65, 0.77]	+0.192 (<0.001)
same decision, random direction	—	0.099	-0.427 (<0.001)	—	0.184	-0.349 (<0.001)	—	0.247	-0.273 (<0.001)
same decision, refusal	82/60	0.427	-0.099 (<0.001)	86/59	0.457	-0.077 (<0.001)	81/49	0.305	-0.214 (<0.001)

Table 4: On the first held-out prompt set (3,000 per model), PTS is more selective than prompting, the global shift, CAST and refusal on its decision in all three families. rem./kept: % of toxic removed / clean kept; Δ_{Sel} : row minus PTS with the shift.

gated arms lie within 0.025 of their composition, with a median gap of 0.005. For post-hoc and prompt decisions the tuning search thus needs one ungated generation per action (Appendix D).

6 PTS BEYOND THE MAIN SETTING

PTS stays selective on a second model family and on a second benchmark (Table 3). The tuning rule selects ablation as the action on Qwen, where PTS reaches selectivity 0.172 [0.136, 0.210], 0.081 above regenerating the flagged questions without an edit and 0.132 above the CAST-style trace condition. On ARC-Challenge PTS reaches 0.235 [0.168, 0.303] and lowers the calibration error by 0.060. There it is above Linear-AcT and the CAST prompt condition after Holm adjustment, and above the 95th percentile of a hundred random directions sent through its decision. Prompting on Qwen (0.212) and CAA and the CAST-style condition on ARC (0.294, 0.289) have higher point estimates, none separable from PTS after Holm adjustment. The reference arms credit the Qwen result to the decision: thirty displacement-matched random directions reach 0.158 through it.

6.1 THE SETTING OF PCHI

In the setting of PCHI (Li et al., 2026), PTS is more selective than the learned head rescaling under the same probe. PCHI gates a rescaling of 128 attention heads with a probe read where the model states its confidence. We rebuild this setting: Qwen3-4B-Instruct answers OpenMathInstruct-2 problems in a JSON format that ends with a yes/no confidence field, and a probe on the layer-18 state at that field flags overconfident-wrong answers. PTS keeps the probe and replaces the learned rescaling with a shift of the same state. On 3,183 held-out answers it reaches selectivity 0.438, against 0.354 for PCHI (+0.084 [0.058, 0.110]), and a calibration error of 0.188 against 0.234 (Appendix J). At the selected dose the shift converts every flagged answer, as an override would, and 45 of 100 random directions reach the same selectivity, so the reference arms credit the gain to the probe paired with a simple edit; along the confidence direction it also has the lowest calibration error of all arms (random directions: 0.254).

6.2 FREE TEXT: TOXIC CONTINUATION

We test the action on free text, where a continuation has no single readout that a decision could overwrite. We steer toxic continuation of RealToxicityPrompts (Gehman et al., 2020), as Linear-AcT and DSAS do, in all three model families (Appendix H). A RoBERTa classifier labels a continuation

fresh prompts: Sel (removed / kept, %)	Gemma-2-2B	Qwen2.5-7B	Mistral-7B
PTS, selected by the tuning rule	0.549 (70/85)	0.602 (81/80)	0.496 (70/79)
DSAS-style token gate	0.597 (93/67)	0.594 (92/67)	0.410 (53/88)
PTS — gate [95% CI]	−0.048 [−0.10, 0.00]	+0.008 [−0.02, 0.04]	+0.087 [−0.03, 0.21]

Table 5: Fresh prompts (3,000 per model; 429, 604, 88 toxic), protocol pushed before any steered rollout; all arms in Appendix I.

toxic at score 0.5. A clean continuation counts as kept only if the steered one is clean and within sentence-embedding cosine 0.5 of the unsteered one, so rewriting a clean text counts as damage. The direction is the difference of mean activations between toxic and clean continuations, and the decision is a probe on the prompt, read before generation.

PTS is more selective than refusal on its own decision in every family, and it lies above the 95th percentile of a hundred random directions sent through that decision (Table 4). For every gated arm that we compared with refusal, Eq. (3) evaluated with tuning-split conversion rates predicts the held-out winner. On Gemma and Qwen, PTS is also more selective than the global shift and prompting, and on Qwen it is above CAST, 0.534 against 0.443. With Mistral, whose continuations are rarely toxic, PTS lowers the toxic rate from 3.4% to 1.4% at selectivity 0.520 [0.430, 0.605].

On Gemma the tuning rule selects three times the displacement of the quantile map from the toxic to the clean projection law: 0.584 [0.542, 0.624] against 0.540 for a gated shift of equal mean displacement. This scaled edit lies outside Theorem 1, and the theorem’s own map ties its matched shift on fresh prompts (0.538, 0.535), so a shift, the transport map onto a translate, suffices on free text; the selectivity comes from the direction and the decision in front of it.

A second, previously unused set of 3,000 prompts per model tests the method under the same selection rule, which puts the prompt decision in front of a token-gated shift on Gemma and Qwen and in front of a plain shift on Mistral (Table 5). PTS again beats refusal on its decision by 0.15 to 0.17 and all hundred random directions in every family. Against a tuned DSAS-style token gate without the input decision it is level in selectivity after Holm adjustment, but the gate removes more toxic text and rewrites more clean text, while PTS leaves every unflagged continuation identical to the unsteered one and keeps 85% and 80% of clean continuations on Gemma and Qwen, against 67%. The input decision thus buys a guarantee on the inputs it leaves alone, which a per-token gate cannot give, for 9% more generation time (Appendix K).

7 LIMITATIONS

Where a behavior has a single readout, overriding it on every flagged answer is at least as selective as the action (Appendix G); the action matters where no such readout exists, as in free text. Free-text evidence rests on 40-token greedy continuations, one toxicity classifier and one embedding threshold, in 2B to 7B models; the multiple-choice protocols became public with their results.

8 CONCLUSION

PTS accounts exactly for how a decision and a projection-local action compose, for any conditional steerer, CAST and DSAS included: it predicts gated selectivity from one ungated run, bounds a decision before any steered rollout, and attributes each gain to the decision or the edit. PTS leads eight methods on overconfident answers and matches a per-token gate on free text while leaving unflagged text untouched.

AI USE STATEMENT

A large language model (LLM) was employed solely to support improvements in the language, grammar, and clarity of this manuscript. All AI-generated suggestions based on the text originally written by the authors were carefully reviewed, edited, and approved by the authors, who accept full responsibility for the final version of the text.

REFERENCES

- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda, *Refusal in language models is mediated by a single direction*, Advances in Neural Information Processing Systems 37, 2024.
- Amos Azaria and Tom Mitchell, *The internal state of an LLM knows when it’s lying*, Findings of the Association for Computational Linguistics: EMNLP, 2023.
- J. Braun, C. Eickhoff, D. Krueger, S. A. Bahrainian, D. Krasheninnikov. *Understanding (un)reliability of steering vectors in language models*. ICLR Workshop on Foundation Models in the Wild, 2025.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord, *Think you have solved question answering? Try ARC, the AI2 reasoning challenge*, arXiv:1803.05457, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, et al., *Training verifiers to solve math word problems*, arXiv:2110.14168, 2021.
- Chenrui Fan, Yize Cheng, Ming Li, Soheil Feizi, and Tianyi Zhou, *When is your LLM steerable? Predicting steering success from internal states*, arXiv:2606.11599, 2026.
- Alex Ferrando de las Morenas, Xavier Suau, Jordi Gonzàlez, and Pau Rodríguez, *Dynamically scaled activation steering*, arXiv:2512.03661, 2026.
- Gemma Team, *Gemma 2: Improving open language models at a practical size*, arXiv:2408.00118, 2024.
- Moghis Fereidouni, Muhammad Umair Haider, Hassan Sajjad, and A. B. Siddique, *GAPS: Dimension-level gates for conditional activation steering*, arXiv:2609.01878, 2026.
- Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith, *RealToxicityPrompts: Evaluating neural toxic degeneration in language models*, Findings of the Association for Computational Linguistics: EMNLP, 2020.
- Matthias Gelbrich, *On a formula for the L^2 Wasserstein metric between measures on Euclidean and Hilbert spaces*, Mathematische Nachrichten **147** (1990), no. 1, 185–203.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt, *Measuring massive multitask language understanding*, International Conference on Learning Representations, 2021.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, et al., *Language models (mostly) know what they know*, arXiv:2207.05221, 2022.
- Martin Knott and Cyril S. Smith, *On the optimal mapping of distributions*, Journal of Optimization Theory and Applications **43** (1984), 39–49.
- Evgenii Kortukov, Piotr Komorowski, Florian Klein, Paula Engl, Gabriele Sarti, Seong Joon Oh, Sebastian Lapuschkin, and Wojciech Samek, *Predicting future behaviors in reasoning models enables better steering*, arXiv:2606.11172, 2026.
- Bruce W. Lee, Inkit Padhi, Karthikeyan Natesan Ramamurthy, Erik Miehl, Pierre Dognin, Manish Nagireddy, and Amit Dhurandhar, *Programming refusal with conditional activation steering*, International Conference on Learning Representations, 2025.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg, *Inference-time intervention: Eliciting truthful answers from a language model*, Advances in Neural Information Processing Systems 36, 2023.
- Ke Li, Chongzhe Zhang, Zifan Zeng, Feng Liu, Qunli Zhang, and Zheng Hu, *Calibrating overconfidence without sacrificing confidence: Probe-conditioned head intervention for LLMs*, arXiv:2606.09876, 2026.

- Ziyang Luo, Zhongyao Chu, Xinjie He, Youting Wang, Xukui Qin, Runxiong Wu, and Yan-Syuan Chen, *You only pass once: answering and abstaining together in a single forward pass of a frozen language model*, arXiv:2608.14465, 2026.
- Tam Nguyen, Tu Anh Nguyen, Sina Alemohammad, and Richard G. Baraniuk, *Minimizing collateral damage in activation steering*, arXiv:2605.01167, 2026.
- Nina Panickssery, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner, *Steering Llama 2 via contrastive activation addition*, Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics, 2024.
- Pau Rodriguez, Arno Blaas, Michal Klein, Luca Zappella, Nicholas Apostoloff, Marco Cuturi, and Xavier Suau, *Controlling language and diffusion models by transporting activations*, International Conference on Learning Representations, 2025.
- Pau Rodriguez, Michal Klein, Eleonora Gualdoni, Arno Blaas, Luca Zappella, Marco Cuturi, and Xavier Suau, *LinEAS: End-to-end learning of activation steering with a distributional loss*, arXiv:2503.10679, 2025.
- Daniel Scalena, Gabriele Sarti, and Malvina Nissim, *Multi-property steering of language models with dynamic activation composition*, Proceedings of BlackboxNLP, 2024.
- Shashwat Singh, Shauli Ravfogel, Jonathan Herzig, Roei Aharoni, Ryan Cotterell, and Ponnurangam Kumaraguru, *Representation surgery: Theory and practice of affine steering*, Proceedings of the 41st International Conference on Machine Learning, 2024.
- Julian Skifstad, Xinyue Annie Yang, and Glen Chou, *Local Linearity of LLMs Enables Activation Steering via Model-Based Linear Optimal Control*, arXiv:2604.19018, 2026.
- D. Tan, D. Chanin, A. Lynch, B. Paige, D. Kanoulas, A. Garriga-Alonso, R. Kirk, *Analysing the generalisation and reliability of steering vectors*. NeurIPS, 2024.
- Adly Templeton, Tom Conerly, Jonathan Marcus, et al., *Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet*, Transformer Circuits Thread, 2024.
- Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J. Vazquez, Ulisse Mini, and Monte MacDiarmid, *Steering language models with activation engineering*, arXiv:2308.10248, 2023.
- Weixuan Wang, Jingyuan Yang, and Wei Peng, *Semantics-adaptive activation intervention for LLMs via dynamic steering vectors*, International Conference on Learning Representations, 2025.
- Zhengxuan Wu, Aryaman Arora, Atticus Geiger, Zheng Wang, Jing Huang, Dan Jurafsky, Christopher D. Manning, and Christopher Potts, *AxBench: Steering LLMs? Even simple baselines outperform sparse autoencoders*, Proceedings of the 42nd International Conference on Machine Learning, 2025.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, et al., *Representation engineering: A top-down approach to AI transparency*, arXiv:2310.01405, 2023.
- Tianyi Zhao, Yinhan He, Wendy Zheng, Yujie Zhang, and Chen Chen, *Wired for overconfidence: A mechanistic perspective on inflated verbalized confidence in LLMs*, arXiv:2604.01457, 2026.

A CONFIRMATORY PROTOCOL

The multiple-choice protocols (v19 for MMLU, v20 for ARC) were committed to the local repository before any steered rollout on their confirmatory splits was generated and first pushed together with the results, so their ordering rests on local commit times; the later protocols v21 to v24 were pushed to a public repository, and so timestamped by the hosting service, before their rollouts existed (rebuttal/protocols/confirmatory-v19.md, with two amendments made after the tuning pass and before launch), together with the SHA-256 hashes of the selected configurations and of every fitted detector (confirmatory-v19-gemma.sha256).

Splits. All splits are subsets of the 14,042-question MMLU test set, identified by index and constructed by `label_pool.split_records`. The extraction split (400) excludes the first evaluation subset of the pilot protocol; the tuning split (1,200) excludes the extraction split and the five pilot evaluation subsets; the detector pool (4,000) and the confirmatory split (3,000) are consecutive blocks of a seeded permutation of the remaining questions. No question appears in two of the four splits.

Selection rule. For every method the configuration with the highest tuning-split selectivity among those with $\Delta\text{acc} \geq -0.01$ is carried unchanged to the confirmatory split. Gated configurations are scored on the tuning split by combining the ungated rollouts of the same split with the gate’s flags (Prop. 1 makes this exact up to batched-decoding nondeterminism; Appendix D); single-pass configurations are real runs. Every confirmatory number is a real run.

Tests. Primary: PTS post-hoc against each of the seven baselines on pooled M5 selectivity, question-level paired bootstrap with 10^4 resamples, two-sided, Holm-adjusted over the seven comparisons. Accuracy: an arm is non-inferior if the lower end of the 95% interval of its accuracy change is above -0.02 . The single-pass and prompt-decision variants, M4 and M2 selectivities, removal, retention and ECE are secondary and reported without a claim.

Released records. Every unsteered and steered answer of the tuning and confirmatory splits is released with its state, confidence and correctness (`results/release/records.csv.gz`); `results/recompute_from_records.py` recomputes each selectivity and paired interval in Tables 2 and 3, including the null arms, from that file alone. `results/release/toxicity.csv.gz` holds every confirmatory continuation of the free-text setting with its judge score and state (texts omitted for the random-direction arms), `results/release/toxicity_fresh.csv.gz` the same for the fresh split, analysed by `behaviour_specific/toxicity/fresh.py`, and `results/release/pchi_val.csv.gz` the held-out readout gaps of every arm of the PCHI setting. `behaviour_specific/overconfidence/ceiling.py` writes the ceilings of Fig. 2b to `results/v4-pooled/ceiling_{gemma,qwen}.json`, and `behaviour_specific/toxicity/latency.py` times Appendix K.

ARC replication. A second protocol (`confirmatory-v20-arc.md`), frozen before any steered ARC rollout, repeats the Gemma comparison on ARC-Challenge: a detector pool of 1,012 and a tuning split of 400 from the ARC training and validation questions, and the 1,165-question test split as the confirmatory split. The directions, moments and action layer are those fitted on MMLU; every decision is refitted on the ARC pool together with the MMLU pool, and every configuration is selected on the ARC tuning split with the same rule and grids.

Null arms. A third protocol (`confirmatory-v21-null.md`), with the job list and the hash of the runner, was committed and pushed to a public repository before any null rollout existed; the commit is timestamped by the hosting service. On each confirmatory split it adds the frozen PTS decision with an identity action, the frozen decision with the selected action on a unit random direction (seed 0), and an ungated rerun with an identity action. It registers two one-sided hypotheses per setting, PTS above each gated null, with Holm adjustment over the six tests, and the net measures of Appendix F for every arm. Five of the six tests pass; PTS against the single random direction on ARC gives $+0.041$ [$-0.029, 0.109$], Holm $p = 0.12$.

Reference arms. A fourth protocol (`confirmatory-v22-references.md`), pushed before any of its rollouts existed, extends the no-edit null to all four conditional methods, replaces the single random direction by ten (seeds 0 to 9), and adds random directions for the sycophancy edit. It registers 24 one-sided tests (four methods, three settings, no-edit and random-edit null), Holm-adjusted together, of which 18 pass (Table 14); a count of the cases in which Eq. (3) with ungated conversion rates predicts the sign of steer minus matched override (10 of 13); and paired intervals for selectivity at cutoffs 0.4 and 0.6 and over boxed traces for PTS against the CAST-style condition on Gemma. The v22 random-edit null supersedes the single-direction test of v21.

Free text and dense nulls. A fifth protocol (`confirmatory-v23-freetext.md`), pushed with the hashes of the tuning selections, fitted decisions and code before any confirmatory rollout of the toxicity setting existed, fixes the data, judge, retention criterion, arms and ten primary hypotheses of Appendix H, and extends the random directions of PTS post-hoc to one hundred on Gemma-MMLU and Gemma-ARC and to thirty displacement-matched directions on Qwen (Table 15).

Transport map and a third family. A seventh protocol (`confirmatory-v25-transport-mistral.md`), pushed before its rollouts, adds the quantile transport map to the action family on free text, records its selection on the Gemma tuning split, and registers five tests: the transport map against a shift of equal mean displacement on Gemma, and PTS against the global shift, prompting, refusal on its decision and a hundred random directions on Mistral-7B. Three of the five pass after Holm adjustment, including the transport-map test ($p = 0.003$).

PCHI setting. A sixth protocol (`confirmatory-v24-pchi.md`), pushed after the train-split fitting and tuning and before any validation rollout, fixes the arms and three hypotheses of Appendix J.

Analyses added after the protocol. The detect-and-override reference, the difference-of-means decision fitted on 4,400 labels, the latency measurement (Appendix K), the layer analysis (Appendix L), calibration and the out-of-domain negatives were run after the confirmatory results were known. They use the same splits and the frozen decisions, and we report them as exploratory.

B PROOFS

Theorem 1. By locality, $T(h) - h \in \text{span}(V^\top) \mu$ -a.s., hence $\|T(h) - h\|^2 = \|VT(h) - Vh\|^2$ because the rows of V are orthonormal, and $C(T) = \mathbb{E}\|VT(h) - Vh\|^2$. The joint law γ of $(Vh, VT(h))$ is a coupling of (μ_V, ν) by the goal constraint, so $C(T) = \int \|y - x\|^2 d\gamma \geq W_2^2(\mu_V, \nu)$. Because μ_V is absolutely continuous and both laws have finite second moments, Brenier’s theorem gives a unique optimal coupling, induced by $S = \nabla\varphi$ with φ convex. $T^*(h) = h + V^\top(S(Vh) - Vh)$ is admissible ($VV^\top = I_k$ gives $VT^*(h) = S(Vh)$, and the complement is untouched) and attains the bound. For any minimiser the induced coupling is optimal, hence equal to the unique one, so $VT(h) = S(Vh)$ a.s. $\square$

For $k = 1$ absolute continuity can be weakened to atomlessness of μ_V : the monotone rearrangement $F_V^{-1} \circ F_{\mu_V}$ is then an optimal map for every ν with a finite second moment, including targets with atoms such as δ_0 (ablation) or a law truncated at c (clamp); uniqueness of the map can fail only on a μ_V -null set.

Proposition 3 (metric covariance). *If the cost is $(\Delta h)^\top M(\Delta h)$ with $M \succ 0$, the admissible problem is optimal transport of μ_V to ν in $\mathbb{R}^k$ under the cost $\|\cdot\|_{M_V}^2$, $M_V = VMV^\top \succ 0$, and the optimal operator is $T^*(h) = h + V^\top(S_M(Vh) - Vh)$ with $S_M = M_V^{-1/2} \circ \nabla\varphi \circ M_V^{1/2}$, φ convex.*

Proof. Locality writes $\Delta h = V^\top a$ with $a \in \mathbb{R}^k$, and the goal constraint with $VV^\top = I_k$ fixes $a = S(Vh) - Vh$ for the map S in use. Then $\|\Delta h\|_M^2 = a^\top M_V a$. The cost is a latent-space Mahalanobis cost on $\mathbb{R}^k$. It exceeds, in general, the least M -cost of a full-space displacement with the same projection, $a^\top (VM^{-1}V^\top)^{-1}a$, because locality confines Δh to the row space of V . In the coordinates $z = M_V^{1/2}q$ it is Euclidean, the change of variables is internal to $\mathbb{R}^k$, and Brenier’s theorem in z (the pushforward of an absolutely continuous law under an invertible linear map is absolutely continuous) gives $\nabla\varphi$; mapping back gives S_M . $\square$

Corollary 1. (a) If $\nu = \mu_V * \delta_\alpha$ the map $x \mapsto x + \alpha$ is the gradient of a convex function pushing μ_V to ν , hence the Brenier map. Conversely, if the optimal map is $x \mapsto x + \alpha$ then ν is its pushforward, the shift. For one-dimensional Gaussians the monotone map is $x \mapsto m_t + (\sigma_t/\sigma_s)(x - m_s)$, with slope σ_t/σ_s . (b) If two groups of inputs have the same law of p and the probability that an input changes state depends on h only through p , every transfer function changes both groups at the same rate, since both rates are the same integral against the law of p . A decision that reads p alone therefore adds little, which is why the detection statistic reads the full activation. $\square$

Proposition 1. Let R be the event that an OCW trace leaves OCW and L the event that a CR trace leaves CR. Off Γ the output equals the unsteered output, which is in its original state by definition, so $R \subseteq \Gamma$ and $L \subseteq \Gamma$ up to null sets. Then $P(R \mid OCW) = P(\Gamma \mid OCW)P(R \mid OCW, \Gamma) = \text{TPR } \rho_O$, and likewise $P(L \mid CR) = \text{FPR } \rho_C$; subtracting gives Eq. (2). If, within OCW, the indicator that the action converts the trace is independent of Γ , then $\rho_O = P(R^{\text{ungated}} \mid OCW)$, and likewise for CR; the ungated selectivity is $\rho_O - \rho_C$, and the difference is $\rho_C(1 - \text{FPR}) - \rho_O(1 - \text{TPR})$. $\square$

The identity requires that the unflagged output be exactly the unsteered one. The post-hoc gate satisfies this by construction. The single-pass gate satisfies it for traces on which it never fires, since the action hook is the identity until the decision token.

Theorem 2. Write p_O, p_C for the densities of s on OCW and CR answers and $r_O(s), r_C(s)$ for the probabilities that the action converts an answer of each state with statistic s . For a decision region Γ , $\text{TPR } \rho_O = \int_{\Gamma} p_O r_O$ and $\text{FPR } \rho_C = \int_{\Gamma} p_C r_C$, so $\text{Sel}(\Gamma) = \int_{\Gamma} (p_O r_O - p_C r_C)$. Among regions with a given value of $\int_{\Gamma} p_C r_C$, the Neyman–Pearson lemma applied to the nonnegative measures $p_O r_O$ and $p_C r_C$ shows that $\int_{\Gamma} p_O r_O$ is maximised by a superlevel set of their ratio, which is the general statement in (a). If r_O and r_C are constant in s , the ratio is a constant multiple of p_O/p_C , so every objective increasing in TPR and decreasing in FPR is maximised on the likelihood-ratio ROC curve, which is (a). The assumption matters: with $s \in \{a, b\}$, $p_O(a) = p_O(b) = \frac{1}{2}$, $p_C(a) = \frac{1}{4}$, $p_C(b) = \frac{3}{4}$, $r_O(a) = 0.2$, $r_O(b) = 1$, $r_C(a) = 1$ and $r_C(b) = 0$, the likelihood ratio ranks a above b , yet $\Gamma = \{b\}$ gives selectivity 0.5, against -0.15 for $\{a\}$ and 0.35 for $\{a, b\}$. On our data r does depend on s : flagged answers are converted more often than unflagged answers of the same state (Appendix D). Our detector is a logistic-regression probe, a discriminative estimate of the log-likelihood ratio up to the prior term; (a) states what the probe approximates, not that the fitted probe is optimal. (b) The concavified ROC contains the path $(0, 0) \rightarrow (\text{FPR}^*, \text{TPR}^*) \rightarrow (1, 1)$ through the Youden-optimal point, whose area is $\frac{1}{2}(1 + J)$, so $J \leq 2 \text{AUC} - 1$. The override reference has selectivity $\text{TPR} - \text{FPR} \leq J$, and with $\rho_O = \rho_C = \rho$, Eq. (2) gives $\text{Sel} = \rho(\text{TPR} - \text{FPR}) \leq J$. $\square$

Theorem 3. Split $C(T)$ over Γ and Γ^c . The second term is zero by the identity constraint; the first, with the conditional goal constraint, is Theorem 1 for $\mu(\cdot \mid \Gamma)$, which applies when $\mu_V(\cdot \mid \Gamma)$ is absolutely continuous. $\square$

This result fixes the decision. It does not say that a threshold gate is the best decision for the displacement-priced objective; the next statement shows that it generally is not, and when it is.

Proposition 4 (graded dose). *Let $T_{\lambda}(h) = h + \lambda(s) V^{\top}(S(Vh) - Vh)$ with a dose $\lambda : \mathbb{R} \rightarrow [0, 1]$ measurable in the detector statistic s . Write $\pi(s) = P(\text{OCW} \mid s)$, $d(s) = \mathbb{E}[\|S(Vh) - Vh\|^2 \mid s] > 0$, and suppose the behavioral gain is linear in the dose, with slope $\beta_1 > 0$ on OCW and $-\beta_0 < 0$ on CR. The dose maximising gain minus displacement priced at $1/(2c)$ is*

$$\lambda^*(s) = \left[\frac{c(\beta_0 + \beta_1)(\pi(s) - \tau)}{d(s)} \right]_0^1, \quad \tau = \frac{\beta_0}{\beta_0 + \beta_1}.$$

It is zero exactly on $\{\pi \leq \tau\}$ and nondecreasing in the ratio $(\pi(s) - \tau)/d(s)$. It is nondecreasing in π when d is constant, but not in general. The threshold gate $\mathbf{1}[\pi > \tau]$ is optimal iff $\lambda^ \in \{0, 1\}$ almost surely, and strictly suboptimal iff $P(0 < \lambda^* < 1) > 0$.*

Proof. The displacement of T_{λ} is $\lambda(s)^2 \|S(Vh) - Vh\|^2$, with expectation $\mathbb{E}_s[\lambda(s)^2 d(s)]$; the gain is $\mathbb{E}_s[\lambda(s)((\beta_0 + \beta_1)\pi(s) - \beta_0)]$. The objective is an integral over s of a strictly concave quadratic in λ , maximised pointwise on $[0, 1]$ by the clipped stationary point. It is the unique maximiser a.s., so the indicator is optimal iff it agrees with λ^* a.s., i.e. iff $\lambda^* \in \{0, 1\}$ a.s.; otherwise the objective is strictly larger at λ^* on a set of positive measure. For non-monotonicity in π , take s_1, s_2 with $\pi(s_1) = 0.6 < \pi(s_2) = 0.7$, $\tau = 0.5$, $c(\beta_0 + \beta_1) = 1$, $d(s_1) = 0.2$ and $d(s_2) = 1$: then $\lambda^*(s_1) = 0.5 > \lambda^*(s_2) = 0.2$. For a counterexample to strict suboptimality at every finite price, let π take only the values 0.1 and 0.9 with $\tau = 0.5$ and $d \leq 0.4c(\beta_0 + \beta_1)$: then $\lambda^* \in \{0, 1\}$ a.s. and the threshold gate is optimal. $\square$

Proposition 2. Let $\mathcal{F}_t$ be generated by the first t generated tokens and z_t the mean detector feature over them. A decision measurable with respect to $\sigma(z_t)$ is a test between the class-conditional laws of z_t , so Theorem 2 applies with z_t in place of s and nothing is assumed about how z_t relates to the finished trace. The reader hook sits at a layer at or above the actor; within one forward pass the actor at position t sees the decision accumulated over positions $< t$, the running mean needs one sum and one counter per sequence, and an edit at position t changes only the keys and values written at t .

Proposition 5 (computational cost). *Any algorithm implementing a non-degenerate intervention local to a subspace spanned by a dense direction reads all d coordinates of h , so it costs $\Omega(d)$ per token. The edit of Eq. (1) costs $\Theta(kd + k \log K)$ with a K -knot lookup, and the detector adds one d -dimensional dot product per token.*

Proof. If a coordinate i with $(\hat{v})_i \neq 0$ is never read, an adversary perturbs h_i , which changes $p = \langle h, \hat{v} \rangle$ and, on a set of positive measure, the correct output, while the algorithm's output is unchanged. The upper bound counts Vh , a binary search in the lookup, and $h + V^\top(\cdot)$. $\square$

C TUNING GRIDS

Every configuration searched on the tuning split, with the selected one marked $\star$. Gated rows are scored by composing the ungated rollouts of the same split with the decision's flags (Appendix D); single-pass rows are real runs. The rule is the highest selectivity among configurations with $\Delta\text{acc} \geq -0.01$, and, when no configuration of a family meets the floor, the highest selectivity overall; this fallback applies to the additive, CAST and prompt-decision families on Qwen, where every shift in the grid costs more than one accuracy point.

Table 6: Tuning split, gemma-2-2b-it ($n = 1,200$, M5).

family	configuration	Δacc	removal	retention	Sel
additive CAA	alpha-0.375	-0.022	0.831	0.424	+0.255
additive CAA	alpha-0.25 $\star$	-0.002	0.536	0.661	+0.197
additive CAA	alpha-0.5	-0.042	0.974	0.136	+0.110
additive CAA	alpha-0.125	+0.010	0.315	0.763	+0.078
additive CAA	alpha-0.75	-0.153	1.000	0.000	+0.000
additive CAA	alpha-1.0	-0.205	1.000	0.000	+0.000
additive CAA	alpha-1.25	-0.270	1.000	0.000	+0.000
additive CAA	alpha-1.5	-0.299	1.000	0.000	+0.000
MiMiC	L10.reg0.0001 $\star$	+0.022	0.349	0.823	+0.172
MiMiC	L10.reg0.1	+0.012	0.349	0.809	+0.158
MiMiC	L10.reg0.01	+0.007	0.320	0.814	+0.135
MiMiC	L18.reg0.1	+0.022	0.260	0.860	+0.121
MiMiC	L18.reg0.01	+0.026	0.242	0.869	+0.111
MiMiC	L18.reg0.0001	+0.023	0.250	0.857	+0.107
MiMiC	L14.reg0.01	+0.001	0.224	0.857	+0.081
MiMiC	L14.reg0.1	+0.004	0.216	0.850	+0.066
MiMiC	L14.reg0.0001	-0.013	0.214	0.841	+0.055
Linear-AcT	L10.lam0.5 $\star$	+0.016	0.133	0.936	+0.069
Linear-AcT	L10.lam1.0	+0.015	0.138	0.919	+0.057
Linear-AcT	L18.lam1.0	+0.012	0.130	0.926	+0.056
Linear-AcT	L14.lam0.25	+0.007	0.128	0.924	+0.052
Linear-AcT	L14.lam1.0	+0.009	0.138	0.913	+0.051
Linear-AcT	L14.lam0.5	+0.010	0.125	0.919	+0.044
Linear-AcT	L18.lam0.5	+0.004	0.117	0.926	+0.043
Linear-AcT	L10.lam0.25	+0.010	0.112	0.919	+0.031
Linear-AcT	L18.lam0.25	+0.005	0.096	0.922	+0.019
CAST-style, trace condition	tau.cr.q30.alpha-0.375	-0.020	0.815	0.458	+0.273
CAST-style, trace condition	tau.cr.q50.alpha-0.375	-0.015	0.732	0.532	+0.264
CAST-style, trace condition	tau.cr.q70.alpha-0.5	-0.014	0.638	0.611	+0.249
CAST-style, trace condition	tau.cr.q70.alpha-0.375 $\star$	-0.006	0.560	0.682	+0.242
CAST-style, trace condition	tau.cr.q70.alpha-0.75	-0.066	0.646	0.595	+0.241
CAST-style, trace condition	tau.cr.q50.alpha-0.5	-0.028	0.859	0.380	+0.239
CAST-style, trace condition	tau.cr.q50.alpha-0.75	-0.110	0.875	0.353	+0.228
CAST-style, trace condition	tau.cr.q90.alpha-0.5	-0.001	0.411	0.802	+0.214
CAST-style, trace condition	tau.cr.q90.alpha-0.75	-0.022	0.414	0.795	+0.209
CAST-style, trace condition	tau.cr.q90.alpha-0.375	+0.000	0.370	0.832	+0.202
CAST-style, trace condition	tau.cr.q30.alpha-0.25	-0.001	0.521	0.670	+0.190
CAST-style, trace condition	tau.cr.q50.alpha-0.25	+0.000	0.474	0.716	+0.190
CAST-style, trace condition	tau.cr.q30.alpha-0.5	-0.038	0.945	0.244	+0.189
CAST-style, trace condition	tau.cr.q70.alpha-0.25	+0.000	0.352	0.809	+0.161
CAST-style, trace condition	tau.cr.q30.alpha-0.75	-0.139	0.969	0.178	+0.147
CAST-style, trace condition	tau.cr.q90.alpha-0.25	+0.001	0.245	0.894	+0.139
CAST, prompt condition	L16.q50.alpha-0.375	-0.018	0.607	0.673	+0.280
CAST, prompt condition	L16.q30.alpha-0.375	-0.017	0.711	0.565	+0.276
CAST, prompt condition	L14.q30.alpha-0.375	-0.017	0.708	0.567	+0.275
CAST, prompt condition	L14.q50.alpha-0.375	-0.018	0.594	0.680	+0.274
CAST, prompt condition	L16.q50.alpha-0.5	-0.022	0.682	0.580	+0.262
CAST, prompt condition	L14.q50.alpha-0.5	-0.021	0.669	0.588	+0.258
CAST, prompt condition	L16.q50.alpha-0.75	-0.069	0.690	0.562	+0.252
CAST, prompt condition	L14.q50.alpha-0.75	-0.070	0.674	0.571	+0.245
CAST, prompt condition	L14.q30.alpha-0.5	-0.034	0.810	0.415	+0.225
CAST, prompt condition	L16.q30.alpha-0.5	-0.034	0.815	0.408	+0.223
CAST, prompt condition	L14.q30.alpha-0.75	-0.106	0.820	0.378	+0.198
CAST, prompt condition	L16.q30.alpha-0.75	-0.104	0.826	0.369	+0.195
CAST, prompt condition	L16.q30.alpha-0.25 $\star$	-0.007	0.443	0.730	+0.172
CAST, prompt condition	L14.q30.alpha-0.25	-0.008	0.440	0.728	+0.168
CAST, prompt condition	L16.q50.alpha-0.25	-0.009	0.362	0.792	+0.153
CAST, prompt condition	L14.q50.alpha-0.25	-0.010	0.354	0.792	+0.146
PTS post-hoc	q60.alpha-0.75	-0.031	0.570	0.763	+0.334

family	configuration	Δ acc	removal	retention	Sel
PTS post-hoc	q50.alpha-0.75	-0.050	0.674	0.654	+0.328
PTS post-hoc	q40.alpha-0.75	-0.073	0.771	0.555	+0.326
PTS post-hoc	q30.alpha-0.75	-0.093	0.862	0.454	+0.316
PTS post-hoc	q30.alpha-0.375	-0.013	0.737	0.578	+0.315
PTS post-hoc	q60.alpha-0.375*	-0.003	0.510	0.790	+0.300
PTS post-hoc	q40.alpha-0.375	-0.015	0.654	0.645	+0.299
PTS post-hoc	q50.alpha-0.375	-0.012	0.586	0.707	+0.293
PTS post-hoc	q30.alpha-0.25	+0.002	0.477	0.735	+0.212
PTS post-hoc	q40.alpha-0.25	+0.002	0.424	0.767	+0.191
PTS post-hoc	q50.alpha-0.25	+0.002	0.380	0.806	+0.186
PTS post-hoc	q60.alpha-0.25	+0.005	0.320	0.864	+0.184
PTS post-hoc	q30.ablate	+0.012	0.388	0.765	+0.153
PTS post-hoc	q40.ablate	+0.014	0.336	0.795	+0.131
PTS post-hoc	q50.ablate	+0.016	0.281	0.841	+0.122
PTS post-hoc	q60.ablate	+0.012	0.232	0.887	+0.119
PTS, prefix decision	q50.at32.alpha-0.375*	+0.004	0.573	0.737	+0.310
PTS, prefix decision	q50.at16.alpha-0.375	-0.002	0.573	0.731	+0.304
PTS, prefix decision	q30.at16.alpha-0.375	+0.002	0.729	0.572	+0.302
PTS, prefix decision	q50.at64.alpha-0.375	+0.017	0.562	0.739	+0.301
PTS, prefix decision	q30.at32.alpha-0.375	+0.002	0.690	0.602	+0.293
PTS, prefix decision	q50.at128.alpha-0.375	-0.001	0.547	0.708	+0.255
PTS, prefix decision	q30.at64.alpha-0.375	+0.010	0.656	0.595	+0.252
PTS, prefix decision	q30.at128.alpha-0.375	+0.000	0.651	0.574	+0.225
PTS, prefix decision	q50.at32	+0.019	0.245	0.894	+0.139
PTS, prefix decision	q30.at32	+0.018	0.320	0.816	+0.137
PTS, prefix decision	q50.at64	+0.002	0.214	0.898	+0.111
PTS, prefix decision	q30.at16	-0.002	0.326	0.781	+0.106
PTS, prefix decision	q50.at16	+0.003	0.224	0.878	+0.102
PTS, prefix decision	q50.at128	+0.000	0.182	0.912	+0.094
PTS, prefix decision	q30.at64	-0.001	0.250	0.836	+0.086
PTS, prefix decision	q30.at128	-0.003	0.214	0.871	+0.085
PTS, prompt decision	q60.alpha-0.75	-0.041	0.570	0.767	+0.337
PTS, prompt decision	q60.alpha-0.5*	-0.009	0.568	0.769	+0.336
PTS, prompt decision	q50.alpha-0.75	-0.061	0.656	0.661	+0.317
PTS, prompt decision	q50.alpha-0.5	-0.018	0.654	0.663	+0.316
PTS, prompt decision	q60.alpha-0.375	-0.015	0.518	0.790	+0.308
PTS, prompt decision	q50.alpha-0.375	-0.010	0.589	0.712	+0.301
PTS, prompt decision	q30.alpha-0.375	-0.013	0.724	0.569	+0.293
PTS, prompt decision	q40.alpha-0.5	-0.024	0.724	0.567	+0.291
PTS, prompt decision	q40.alpha-0.75	-0.084	0.729	0.562	+0.291
PTS, prompt decision	q40.alpha-0.375	-0.012	0.646	0.643	+0.289
PTS, prompt decision	q30.alpha-0.5	-0.020	0.826	0.463	+0.288
PTS, prompt decision	q30.alpha-0.75	-0.099	0.839	0.445	+0.284
PTS, prompt decision	q30.alpha-0.25	+0.007	0.464	0.735	+0.199
PTS, prompt decision	q60.alpha-0.25	-0.002	0.320	0.871	+0.191
PTS, prompt decision	q40.alpha-0.25	+0.004	0.409	0.781	+0.190
PTS, prompt decision	q50.alpha-0.25	+0.002	0.365	0.825	+0.190
earlier gate + ablation	u.crg50.ablate*	+0.007	0.378	0.740	+0.118
ungated actions	ablate	+0.007	0.451	0.691	+0.141
ungated actions	clamp.q50	+0.017	0.307	0.814	+0.122
ungated actions	clamp.q30	+0.012	0.339	0.763	+0.102
ungated actions	otq.cal	+0.004	0.115	0.926	+0.040

Table 7: Tuning split, Qwen2.5-7B-Instruct ($n = 1,200$, M5).

family	configuration	Δ acc	removal	retention	Sel
additive CAA	alpha-0.125*	-0.026	0.212	0.911	+0.123
additive CAA	alpha-0.375	-0.050	0.208	0.892	+0.100
additive CAA	alpha-0.25	-0.044	0.200	0.887	+0.087
additive CAA	alpha-0.5	-0.076	0.204	0.855	+0.059
additive CAA	alpha-0.75	-0.138	0.233	0.759	-0.007
additive CAA	alpha-1.0	-0.217	0.412	0.544	-0.043
additive CAA	alpha-1.25	-0.334	0.567	0.272	-0.161
additive CAA	alpha-1.5	-0.473	0.375	0.230	-0.395
MiMiC	L18.reg0.0001	-0.013	0.279	0.906	+0.185
MiMiC	L18.reg0.01*	-0.007	0.263	0.911	+0.173
MiMiC	L14.reg0.01	-0.018	0.254	0.912	+0.166
MiMiC	L14.reg0.0001	-0.023	0.233	0.912	+0.145
MiMiC	L10.reg0.1	-0.019	0.250	0.890	+0.140
MiMiC	L14.reg0.1	-0.021	0.217	0.923	+0.140
MiMiC	L18.reg0.1	-0.013	0.233	0.906	+0.139
MiMiC	L10.reg0.01	-0.022	0.237	0.892	+0.129
MiMiC	L10.reg0.0001	-0.031	0.233	0.868	+0.101
Linear-AcT	L18.lam1.0*	-0.004	0.192	0.948	+0.140
Linear-AcT	L14.lam1.0	+0.002	0.175	0.955	+0.130
Linear-AcT	L14.lam0.25	+0.004	0.138	0.969	+0.106
Linear-AcT	L10.lam1.0	-0.008	0.158	0.947	+0.105
Linear-AcT	L18.lam0.25	-0.006	0.138	0.957	+0.095

	family	configuration	Δ acc	removal	retention	Sel
864						
865	Linear-AcT	L10.lam0.25	-0.013	0.133	0.958	+0.092
866	Linear-AcT	L18.lam0.5	-0.006	0.142	0.950	+0.091
867	Linear-AcT	L10.lam0.5	-0.013	0.125	0.951	+0.076
868	Linear-AcT	L14.lam0.5	-0.005	0.113	0.961	+0.073
869	CAST-style, trace condition	tau.cr.q30.alpha-0.375	-0.043	0.188	0.907	+0.094
870	CAST-style, trace condition	tau.cr.q30.alpha-0.25	-0.035	0.183	0.906	+0.089
871	CAST-style, trace condition	tau.cr.q50.alpha-0.375	-0.040	0.158	0.918	+0.076
872	CAST-style, trace condition	tau.cr.q50.alpha-0.25	-0.033	0.154	0.919	+0.074
873	CAST-style, trace condition	tau.cr.q30.alpha-0.5	-0.067	0.188	0.879	+0.067
874	CAST-style, trace condition	tau.cr.q70.alpha-0.375	-0.024	0.108	0.945	+0.053
875	CAST-style, trace condition	tau.cr.q50.alpha-0.5	-0.059	0.154	0.898	+0.052
876	CAST-style, trace condition	tau.cr.q70.alpha-0.25	-0.022	0.104	0.945	+0.049
877	CAST-style, trace condition	tau.cr.q50.alpha-0.75	-0.079	0.188	0.858	+0.045
878	CAST-style, trace condition	tau.cr.q90.alpha-0.25	-0.010	0.067	0.975	+0.041
879	CAST-style, trace condition	tau.cr.q70.alpha-0.5	-0.036	0.104	0.937	+0.041
880	CAST-style, trace condition	tau.cr.q90.alpha-0.375*	-0.009	0.062	0.979	+0.041
881	CAST-style, trace condition	tau.cr.q90.alpha-0.5	-0.018	0.071	0.970	+0.041
882	CAST-style, trace condition	tau.cr.q70.alpha-0.75	-0.048	0.125	0.912	+0.037
883	CAST-style, trace condition	tau.cr.q30.alpha-0.75	-0.098	0.212	0.824	+0.036
884	CAST-style, trace condition	tau.cr.q90.alpha-0.75	-0.023	0.079	0.956	+0.035
885	CAST, prompt condition	L14.q30.alpha-0.25*	-0.032	0.167	0.916	+0.082
886	CAST, prompt condition	L16.q30.alpha-0.25	-0.032	0.167	0.916	+0.082
887	CAST, prompt condition	L14.q30.alpha-0.375	-0.040	0.163	0.918	+0.081
888	CAST, prompt condition	L16.q30.alpha-0.375	-0.040	0.163	0.918	+0.081
889	CAST, prompt condition	L16.q50.alpha-0.375	-0.034	0.142	0.931	+0.072
890	CAST, prompt condition	L14.q50.alpha-0.375	-0.034	0.142	0.929	+0.071
891	CAST, prompt condition	L14.q50.alpha-0.25	-0.026	0.138	0.932	+0.069
892	CAST, prompt condition	L16.q50.alpha-0.25	-0.027	0.133	0.932	+0.065
893	CAST, prompt condition	L14.q30.alpha-0.5	-0.058	0.167	0.890	+0.057
894	CAST, prompt condition	L16.q30.alpha-0.5	-0.058	0.167	0.890	+0.057
895	CAST, prompt condition	L16.q50.alpha-0.5	-0.050	0.133	0.914	+0.048
896	CAST, prompt condition	L14.q50.alpha-0.5	-0.051	0.133	0.913	+0.046
897	CAST, prompt condition	L16.q50.alpha-0.75	-0.078	0.158	0.868	+0.026
898	CAST, prompt condition	L14.q50.alpha-0.75	-0.079	0.150	0.868	+0.018
899	CAST, prompt condition	L14.q30.alpha-0.75	-0.105	0.188	0.821	+0.009
900	CAST, prompt condition	L16.q30.alpha-0.75	-0.105	0.183	0.822	+0.006
901	PTS post-hoc	q40.ablate*	-0.006	0.254	0.927	+0.181
902	PTS post-hoc	q50.ablate	-0.002	0.233	0.940	+0.173
903	PTS post-hoc	q30.ablate	-0.009	0.254	0.918	+0.172
904	PTS post-hoc	q60.ablate	-0.003	0.212	0.948	+0.161
905	PTS post-hoc	q40.alpha-0.375	-0.044	0.183	0.907	+0.090
906	PTS post-hoc	q30.alpha-0.375	-0.048	0.192	0.898	+0.090
907	PTS post-hoc	q40.alpha-0.25	-0.043	0.183	0.902	+0.085
908	PTS post-hoc	q60.alpha-0.375	-0.032	0.154	0.929	+0.084
909	PTS post-hoc	q30.alpha-0.25	-0.045	0.192	0.892	+0.083
910	PTS post-hoc	q60.alpha-0.25	-0.028	0.150	0.931	+0.081
911	PTS post-hoc	q50.alpha-0.375	-0.039	0.163	0.916	+0.078
912	PTS post-hoc	q50.alpha-0.25	-0.037	0.154	0.912	+0.066
913	PTS post-hoc	q60.alpha-0.75	-0.062	0.158	0.883	+0.041
914	PTS post-hoc	q50.alpha-0.75	-0.083	0.179	0.854	+0.033
915	PTS post-hoc	q40.alpha-0.75	-0.098	0.204	0.826	+0.030
916	PTS post-hoc	q30.alpha-0.75	-0.110	0.217	0.805	+0.021
917	PTS, prefix decision	q30.at32	-0.010	0.163	0.951	+0.113
918	PTS, prefix decision	q30.at32.alpha-0.375	-0.012	0.183	0.929	+0.113
919	PTS, prefix decision	q30.at16*	-0.008	0.163	0.948	+0.111
920	PTS, prefix decision	q50.at32	-0.007	0.142	0.966	+0.108
921	PTS, prefix decision	q50.at16	-0.006	0.138	0.960	+0.097
922	PTS, prefix decision	q30.at64	-0.008	0.133	0.951	+0.084
923	PTS, prefix decision	q30.at16.alpha-0.375	-0.030	0.167	0.916	+0.082
924	PTS, prefix decision	q50.at32.alpha-0.375	-0.013	0.133	0.943	+0.077
925	PTS, prefix decision	q30.at64.alpha-0.375	-0.014	0.133	0.942	+0.075
926	PTS, prefix decision	q50.at64	-0.005	0.108	0.963	+0.072
927	PTS, prefix decision	q30.at128.alpha-0.375	-0.005	0.083	0.969	+0.052
928	PTS, prefix decision	q50.at128.alpha-0.375	-0.004	0.079	0.972	+0.051
929	PTS, prefix decision	q50.at16.alpha-0.375	-0.029	0.113	0.938	+0.051
930	PTS, prefix decision	q50.at64.alpha-0.375	-0.018	0.083	0.955	+0.038
931	PTS, prefix decision	q30.at128	+0.002	0.058	0.979	+0.037
932	PTS, prefix decision	q50.at128	+0.002	0.050	0.977	+0.027
933	PTS, prompt decision	q30.alpha-0.375*	-0.045	0.204	0.901	+0.105
934	PTS, prompt decision	q30.alpha-0.25	-0.042	0.196	0.899	+0.095
935	PTS, prompt decision	q40.alpha-0.375	-0.043	0.188	0.907	+0.094
936	PTS, prompt decision	q50.alpha-0.375	-0.040	0.171	0.916	+0.086
937	PTS, prompt decision	q40.alpha-0.25	-0.043	0.175	0.906	+0.081
938	PTS, prompt decision	q50.alpha-0.25	-0.038	0.163	0.914	+0.077
939	PTS, prompt decision	q60.alpha-0.375	-0.033	0.138	0.935	+0.072
940	PTS, prompt decision	q30.alpha-0.5	-0.068	0.204	0.868	+0.072
941	PTS, prompt decision	q60.alpha-0.25	-0.031	0.133	0.929	+0.063
942	PTS, prompt decision	q50.alpha-0.5	-0.056	0.163	0.897	+0.059
943	PTS, prompt decision	q40.alpha-0.5	-0.068	0.175	0.874	+0.049
944	PTS, prompt decision	q60.alpha-0.5	-0.047	0.129	0.914	+0.044

family	configuration	Δacc	removal	retention	Sel
PTS, prompt decision	q50_alpha-0.75	-0.080	0.179	0.855	+0.034
PTS, prompt decision	q60_alpha-0.75	-0.067	0.146	0.884	+0.030
PTS, prompt decision	q30_alpha-0.75	-0.108	0.217	0.810	+0.026
PTS, prompt decision	q40_alpha-0.75	-0.101	0.196	0.826	+0.022
earlier gate + ablation	u.crq50_ablate*	-0.008	0.196	0.938	+0.134
ungated actions	ablate	-0.008	0.258	0.913	+0.171
ungated actions	clamp.q30	-0.015	0.221	0.918	+0.139
ungated actions	otq.cal	-0.003	0.154	0.962	+0.116
ungated actions	clamp.q50	-0.008	0.183	0.928	+0.112

D THE SELECTIVITY IDENTITY IN PRACTICE

Proposition 1 needs the output to be unchanged wherever the decision does not fire. For the post-hoc and prompt decisions this holds by construction, so a gated arm equals, question by question, either its unsteered row or its ungated row. Composing the two requires one ungated run per action and one set of decision scores, and it differs from a real gated run only through batched-decoding nondeterminism: a flagged question is generated in a different batch in the two cases. Table 8 compares the composition with the real run for the four gated arms of each confirmatory split, and reports the prediction from the ungated conversion rates (independence within each state).

setting	arm	TPR	FPR	ρ_O	ρ_C	predicted	composed	real run
Gemma, MMLU	PTS post-hoc	0.541	0.222	0.841	0.585	+0.325	+0.296	+0.294
	PTS at the prompt	0.549	0.242	0.981	0.872	+0.328	+0.306	+0.306
	CAST-style, trace	0.603	0.380	0.841	0.585	+0.285	+0.244	+0.248
	CAST, prompt	0.805	0.635	0.559	0.339	+0.234	+0.207	+0.217
Qwen, MMLU	PTS post-hoc	0.869	0.478	0.240	0.084	+0.168	+0.148	+0.172
	PTS at the prompt	0.920	0.594	0.247	0.111	+0.161	+0.126	+0.120
	CAST-style, trace	0.286	0.101	0.247	0.111	+0.059	+0.030	+0.040
	CAST, prompt	0.817	0.670	0.277	0.098	+0.161	+0.143	+0.122
Gemma, ARC	PTS post-hoc	0.675	0.405	0.513	0.219	+0.258	+0.239	+0.235
	PTS at the prompt	0.417	0.247	0.991	0.896	+0.191	+0.181	+0.182
	CAST-style, trace	0.912	0.767	0.513	0.219	+0.300	+0.297	+0.289
	CAST, prompt	0.386	0.379	0.513	0.219	+0.115	+0.107	+0.109

Table 8: Real gated runs against the composition of ungated runs and against the prediction TPR ρ_O – FPR ρ_C with ungated conversion rates.

Across the twelve arms the composition differs from the real run by at most 0.025 (Qwen, PTS post-hoc) and by 0.005 at the median. The prediction is higher than the real run in eleven of twelve rows, and the same bias appears on the tuning split (Fig. 2a). Flagged answers are converted more often than the average answer of their state in both states: for PTS post-hoc, 0.91 of flagged overconfident-wrong answers change state against 0.84 ungated, and 0.89 of flagged confident-right answers against 0.59 ungated. The larger increase on the protected state lowers selectivity below the independence prediction. The composition, which uses no such assumption, is what scores the gated configurations during selection.

E SECONDARY READOUTS

Table 9 reports the Gemma confirmatory arms under the two secondary readouts. The letter-logit confidence M2 assigns confidence above 0.5 to nearly every answer, so its OCW stratum is close to the set of wrong answers and every method moves it little. Under M4 the edits act on four separate verification generations; PTS is level with the ungated edits there, which is expected for a decision fitted to M5 labels.

Table 11 gives the full state transitions of the frozen PTS arms under M5. Selectivity counts only the first and last rows; the table also shows the answers the edit corrects or breaks and the few answers it turns overconfident.

Table 12 reports calibration of the stated confidence on the confirmatory split: expected calibration error (10 bins), Brier score, and the AUROC with which the confidence separates right from wrong answers.

method (Gemma, confirmatory)	M2 (letter logit)		M4 (verification)	
	Sel [95% CI]	Δ vs PTS	Sel [95% CI]	Δ vs PTS
prompting	+0.083 [+0.052, +0.113]	-0.008 [-0.045, +0.028]	+0.138 [+0.086, +0.191]	-0.102 [-0.167, -0.034]
additive CAA	+0.089 [+0.059, +0.120]	-0.001 [-0.036, +0.034]	+0.242 [+0.187, +0.296]	+0.003 [-0.061, +0.065]
dir. ablation	+0.092 [+0.062, +0.122]	+0.002 [-0.034, +0.036]	+0.246 [+0.193, +0.299]	+0.007 [-0.054, +0.070]
MiMiC	+0.064 [+0.036, +0.092]	-0.027 [-0.063, +0.008]	+0.244 [+0.189, +0.298]	+0.004 [-0.059, +0.067]
Linear-AcT	+0.029 [+0.009, +0.050]	-0.061 [-0.091, -0.031]	+0.111 [+0.068, +0.155]	-0.128 [-0.185, -0.070]
CAST (prompt)	+0.073 [+0.045, +0.102]	-0.017 [-0.052, +0.017]	+0.207 [+0.155, +0.259]	-0.032 [-0.092, +0.027]
CAST-style (trace)	+0.074 [+0.046, +0.103]	-0.016 [-0.041, +0.010]	+0.224 [+0.173, +0.277]	-0.015 [-0.063, +0.033]
PTS, post-hoc	+0.090 [+0.064, +0.118]	—	+0.239 [+0.192, +0.286]	—
PTS, prompt	+0.099 [+0.069, +0.128]	+0.008 [-0.023, +0.038]	+0.245 [+0.198, +0.292]	+0.006 [-0.039, +0.052]

Table 9: Secondary readouts on the confirmatory split; paired bootstrap intervals, no multiplicity adjustment.

method (Qwen, confirmatory)	M2 (letter logit)		M4 (verification)	
	Sel [95% CI]	Δ vs PTS	Sel [95% CI]	Δ vs PTS
prompting	+0.132 [+0.102, +0.163]	+0.013 [-0.021, +0.046]	+0.178 [+0.144, +0.213]	+0.073 [+0.033, +0.114]
additive CAA	+0.118 [+0.089, +0.149]	-0.000 [-0.033, +0.034]	+0.096 [+0.060, +0.132]	-0.009 [-0.050, +0.033]
dir. ablation	+0.119 [+0.089, +0.150]	+0.000 [-0.024, +0.025]	+0.192 [+0.159, +0.225]	+0.087 [+0.057, +0.117]
MiMiC	+0.130 [+0.100, +0.160]	+0.011 [-0.023, +0.045]	+0.104 [+0.073, +0.136]	-0.001 [-0.040, +0.038]
Linear-AcT	+0.099 [+0.072, +0.127]	-0.020 [-0.053, +0.013]	+0.078 [+0.051, +0.106]	-0.026 [-0.062, +0.009]
CAST (prompt)	+0.110 [+0.081, +0.140]	-0.009 [-0.043, +0.025]	+0.159 [+0.125, +0.193]	+0.054 [+0.012, +0.096]
CAST-style (trace)	+0.030 [+0.012, +0.048]	-0.089 [-0.120, -0.059]	+0.121 [+0.096, +0.145]	+0.016 [-0.021, +0.053]
PTS, post-hoc	+0.119 [+0.090, +0.149]	—	+0.105 [+0.076, +0.135]	—
PTS, prompt	+0.096 [+0.063, +0.128]	-0.023 [-0.059, +0.012]	+0.170 [+0.134, +0.206]	+0.065 [+0.022, +0.108]

Table 10: Secondary readouts on the Qwen confirmatory split.

Where the edit acts. Rerunning the frozen PTS decision with the action applied only during the reasoning trace or only during the confidence turn gives, on Gemma (shift), 0.159 and 0.274 [0.237, 0.312] against 0.294 for both turns, and on Qwen (ablation) 0.172 and 0.006 against 0.172 (exploratory).

F NET EFFECTS OF EVERY ARM

Selectivity is computed on the answers that start in OCW or CR; it does not see answers that enter OCW from another state. Table 13 reports, for every confirmatory arm and the three null arms of Appendix A, selectivity net of the no-edit null, the change in the share of OCW answers among all questions, the change in the rate of confident answers among wrong and among right answers, Δ ECE, Δ Brier, the change in the share of traces without a `\boxed{ }` answer (which then receive a forced box, Appendix O), and selectivity recomputed without the questions whose steered trace has no box. The ungated rerun reproduces the unsteered outputs exactly, because it uses the same batches; the gated no-edit null regenerates the flagged questions in different batches, and the resulting nondeterminism alone yields selectivity 0.049 (Gemma), 0.091 (Qwen) and 0.016 (ARC). Excluding traces without a box leaves the ordering on Gemma unchanged: PTS post-hoc 0.257, CAST-style 0.233, additive 0.215.

G REFERENCE ARMS FOR EVERY CONDITIONAL METHOD

Table 14 applies the three references of Sec. 3 to the four conditional multiple-choice arms of each setting, with ten random directions per arm (protocol v22). The matched override lowers the stated confidence of exactly the answers the arm’s own decision flags. PTS post-hoc passes all six of its registered null tests, and 18 of the 24 tests pass after Holm adjustment. Eq. (3), evaluated with ungated conversion rates before the gated runs, predicts the sign of steer minus override in all rows where the predicted difference exceeds 0.021. The CAST conditions fire on many protected answers and their shift spares most of them, so steering is more selective than overriding with them. Table 15 reports the hundred-direction test of protocol v23 for PTS post-hoc.

H FREE-TEXT BEHAVIOR: TOXIC CONTINUATION

Data and readout. Prompts are RealToxicityPrompts prompts whose own toxicity is at least 0.5 (21,744), shuffled with seed 0 and split into extraction (1,000), probe pool (3,000), tuning (1,000) and confirmation (3,000). Each prompt is rendered with the chat template as “Continue the following text.” and continued for 40 greedy tokens. The judge is

	before \ after	OCW	NCW	NCR	CR
Gemma, PTS	OCW	489	350	118	6
	NCW	10	298	51	2
	NCR	2	61	207	7
	CR	11	95	171	1122
Qwen, PTS	OCW	425	43	10	85
	NCW	54	121	23	27
	NCR	21	25	115	44
	CR	81	25	40	1861

Table 11: State transitions of the frozen PTS arms on the confirmatory split (counts of questions).

arm	Gemma-2-2B			Qwen2.5-7B		
	ECE ↓	Brier ↓	AUROC ↑	ECE ↓	Brier ↓	AUROC ↑
unsteered	0.354	0.372	0.643	0.215	0.230	0.741
shift −0.25, ungated	0.202	0.281	0.659	0.238	0.245	0.741
ablation, ungated	0.243	0.296	0.653	0.220	0.232	0.743
PTS, post-hoc	0.255	0.294	0.689	0.215	0.229	0.740
override, probe on answer	0.234	0.285	0.684	0.139	0.195	0.760
override, CAST condition	0.197	0.297	0.646	0.194	0.222	0.714

Table 12: Calibration of the stated confidence (M5) on the confirmatory split.

`s-nlp/roberta-toxicity-classifier`; a continuation is toxic at probability 0.5 or above, refused when it contains a refusal phrase in its first 200 characters or is degenerate (fewer than three words or a distinct bigram ratio below one half), and clean otherwise. Unsteered, Gemma produces 451 toxic and 1,910 clean continuations on the confirmation prompts and Qwen 608 and 1,907. Selectivity is the share of toxic continuations that are no longer toxic minus the share of clean continuations lost, where a clean continuation is kept when the steered one is clean and is identical to, or has `all-mpnet-base-v2` cosine at least 0.5 with, the unsteered one. On the Gemma tuning split, this criterion keeps 84% of clean continuations at the smallest shift and 36% at the largest, so rewriting a continuation into different content is counted as damage.

Operator. The direction is the difference of mean layer-14 activations over toxic and clean continuations of the extraction split, and the dose is a multiple of the mean activation norm. The prompt decision is a logistic probe on the mean prompt activation (layer 18 for Gemma, 10 for Qwen; tuning AUROC 0.808 and 0.823), the post-hoc decision a probe on the unsteered continuation (0.946 and 0.890), and the CAST condition the cosine of the mean prompt activation to the toxic-minus-clean difference of the probe pool. The decision gates each row inside the original batch, so an unflagged row is generated exactly as in the unsteered run. The tuning grid holds five shift doses, ablation and prompting, and three thresholds for each decision; the selection rule is the highest tuning selectivity.

Results. Table 4 gives the confirmatory comparison. Nine of the ten registered tests pass after Holm adjustment. Refusing on a decision converts every flagged clean continuation, so by Eq. (3) a gated edit beats it whenever it spares enough of them; Eq. (3) with conversion rates from the tuning split predicts the sign of steer minus refuse for all six gated arms. Against a hundred random directions through the same decision, PTS lies above the 95th percentile (0.311 on Gemma, 0.479 on Qwen).

Changes after the v23 freeze. Two edits to the analysis code followed the freeze and precede none of the registered quantities: the conversion rates that feed the Eq. (3) prediction are taken from the tuning split, because the confirmatory split has no ungated run of the post-hoc and CAST actions, and a loading routine was rewritten for speed without changing any number. Data, judge, retention criterion, arms and the ten hypotheses are as frozen.

Quantile transport map. The map sends the per-token projection on the toxicity direction from its law over toxic continuations to its law over clean continuations of the extraction split (41 quantiles), and PTS applies $h + \lambda V^\top (S(Vh) - Vh)$ on the rows its prompt decision flags. On the Gemma tuning split $\lambda = 1, 2, 3, 4, 6$ gives selectivity 0.373, 0.577, 0.605, 0.586 and 0.552, against 0.594 for the tuned shift, so the selection rule picks $\lambda = 3$. For $\lambda > 1$ the edit extrapolates the displacement of the map and is not itself monotone on the whole source law, so it is an edit along the transport displacement rather than the map of Theorem 1, which ($\lambda = 1$) reaches 0.373. Its mean displacement on toxic tokens equals a shift of $-0.432\|h\|$. On the held-out prompts it reaches 0.584 [0.542, 0.624]

arm	Sel	Sel-null	Δ OCW (pp)	ΔP (conf wrong) (pp)	ΔP (conf right) (pp)	Δ ECE	Δ Brier	Δ no box (pp)	Sel, boxed only
<i>Gemma-2-2B, MMLU</i>									
prompting	0.200	+0.151	-5.9	-12.7	-6.1	-0.030	-0.032	+1.1	0.185
additive CAA	0.220	+0.171	-13.0	-29.4	-18.9	-0.153	-0.091	+1.3	0.215
dir. ablation	0.173	+0.124	-8.7	-18.6	-13.0	-0.111	-0.076	+0.4	0.175
MiMiC	0.117	+0.068	-2.1	-4.7	-4.6	-0.018	-0.020	+0.0	0.106
Linear-AcT	0.061	+0.012	-0.2	+0.4	+0.6	-0.003	-0.006	-0.1	0.053
CAST (prompt)	0.217	+0.168	-11.2	-25.6	-13.4	-0.137	-0.086	+0.6	0.206
CAST-style (trace)	0.248	+0.199	-16.6	-38.0	-23.1	-0.094	-0.067	+8.2	0.233
PTS, post-hoc	0.294	+0.245	-15.0	-33.8	-16.0	-0.099	-0.079	+5.7	0.257
PTS, prefix	0.294	+0.245	-17.2	-39.5	-20.2	-0.136	-0.091	+2.0	0.284
PTS, prompt	0.306	+0.257	-17.6	-40.8	-18.7	-0.084	-0.077	+15.0	0.247
null: decision, no edit	0.049	+0.000	+0.0	+0.5	+0.9	-0.005	-0.007	-0.4	0.037
null: decision, random dir.	0.215	+0.166	-8.9	-19.6	-7.9	-0.077	-0.061	+7.2	0.185
rerun, no edit	0.000	-0.049	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000
<i>Qwen2.5-7B, MMLU</i>									
prompting	0.212	+0.121	-1.3	-5.6	-2.0	-0.004	-0.002	+3.6	0.196
additive CAA	0.151	+0.060	+2.9	+6.5	+3.5	+0.012	+0.008	-0.8	0.149
dir. ablation	0.156	+0.065	+1.3	+2.3	+0.8	+0.005	+0.001	+0.6	0.156
MiMiC	0.154	+0.063	+1.4	+4.9	+0.2	+0.020	+0.014	+0.5	0.151
Linear-AcT	0.123	+0.032	+0.6	+1.4	+0.1	+0.009	+0.004	+0.7	0.114
CAST (prompt)	0.122	+0.031	+3.7	+10.4	+3.0	+0.027	+0.021	+0.5	0.124
CAST-style (trace)	0.040	-0.051	+2.4	+7.2	+2.4	+0.013	+0.009	+0.2	0.036
PTS, post-hoc	0.172	+0.081	+0.6	+1.6	+0.7	-0.000	-0.002	+0.5	0.166
PTS, prefix	0.109	+0.018	+1.0	+1.6	-0.6	+0.013	+0.010	+0.1	0.092
PTS, prompt	0.120	+0.029	+6.9	+15.2	+5.4	+0.049	+0.039	+1.1	0.123
null: decision, no edit	0.091	+0.000	-0.1	-0.9	-0.2	+0.000	-0.000	+0.1	0.084
null: decision, random dir.	0.119	+0.028	-0.1	-1.6	+0.4	-0.002	-0.004	-0.2	0.117
rerun, no edit	0.000	-0.091	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000
<i>Gemma-2-2B, ARC</i>									
prompting	0.238	+0.222	-1.8	-6.1	-3.7	+0.009	+0.005	-0.3	0.234
additive CAA	0.294	+0.278	-5.9	-23.3	-11.4	-0.069	-0.013	+1.2	0.289
dir. ablation	0.215	+0.198	-3.5	-15.6	-8.6	-0.055	-0.007	+1.0	0.216
MiMiC	0.198	+0.182	-1.7	-3.8	-3.6	-0.007	-0.003	+0.7	0.198
Linear-AcT	0.109	+0.093	-0.1	+0.4	+0.7	+0.000	-0.001	+0.3	0.108
CAST (prompt)	0.109	+0.093	-2.0	-8.6	-5.3	-0.028	-0.002	+0.3	0.090
CAST-style (trace)	0.289	+0.273	-5.4	-20.9	-10.7	-0.061	-0.015	+1.6	0.276
PTS, post-hoc	0.235	+0.219	-3.8	-16.2	-5.4	-0.060	-0.025	+1.0	0.226
PTS, prefix	0.276	+0.260	-9.2	-37.1	-17.6	+0.000	+0.003	+1.5	0.267
PTS, prompt	0.182	+0.166	-8.1	-35.8	-18.4	+0.028	+0.019	+10.6	0.150
null: decision, no edit	0.016	+0.000	+1.7	+5.3	+0.5	+0.011	+0.011	+0.3	0.020
null: decision, random dir.	0.194	+0.178	-3.7	-15.0	-6.6	-0.006	-0.001	+2.7	0.176
rerun, no edit	0.000	-0.016	+0.0	+0.0	+0.0	+0.000	+0.000	+0.0	0.000

Table 13: Net effects on the confirmatory splits (point estimates; intervals for the main-text quantities are given there). “null” is the frozen decision with no edit.

setting	method	Sel	no-edit null		random-edit null (10)		matched override		
			Sel	Δ	mean Sel	Δ	Sel	steer	override [95% CI] Eq. 3 sign
Gemma, MMLU	PTS, post-hoc	0.294	0.049	+0.245*	0.141	+0.153*	0.319	-0.025 [-0.04, -0.01]	no
	PTS, prompt	0.306	0.041	+0.265*	0.147	+0.159*	0.307	-0.001 [-0.00, +0.00]	no
	CAST-style (trace)	0.248	0.049	+0.199*	0.112	+0.136*	0.223	+0.025 [+0.00, +0.05]	yes
	CAST (prompt)	0.217	0.044	+0.173*	0.105	+0.112*	0.169	+0.048 [+0.01, +0.09]	yes
Qwen, MMLU	PTS, post-hoc	0.172	0.091	+0.081*	0.122	+0.051*	0.391	-0.218 [-0.26, -0.17]	yes
	PTS, prompt	0.120	0.077	+0.043	0.224	-0.104	0.326	-0.206 [-0.25, -0.16]	yes
	CAST-style (trace)	0.040	0.035	+0.005	0.080	-0.040	0.185	-0.145 [-0.18, -0.11]	yes
	CAST (prompt)	0.122	0.065	+0.057*	0.182	-0.059	0.147	-0.025 [-0.07, +0.02]	no
Gemma, ARC	PTS, post-hoc	0.235	0.016	+0.219*	0.165	+0.070*	0.270	-0.035 [-0.10, +0.03]	yes
	PTS, prompt	0.182	0.045	+0.137*	0.119	+0.063*	0.169	+0.013 [+0.01, +0.02]	yes
	CAST-style (trace)	0.289	0.077	+0.213*	0.211	+0.078*	0.146	+0.144 [+0.07, +0.22]	yes
	CAST (prompt)	0.109	0.020	+0.089*	0.078	+0.032	0.006	+0.103 [+0.04, +0.16]	yes

Table 14: Reference arms on the confirmatory splits. Δ is the method minus the reference; *: Holm-adjusted $p < 0.05$ over the 24 registered tests.

against 0.540 [0.495, 0.584] for that matched shift, removes 79% of toxic continuations and keeps 79% of clean ones. On Qwen the rule keeps the shift.

Third family. Mistral-7B-Instruct-v0.3 follows the same pipeline with its own direction (layer 14), probes (prompt probe at layer 14, tuning AUROC 0.787) and tuning. Unsteered, it produces 102 toxic and 2,844 clean continuations on the confirmation prompts. PTS with the prompt decision reaches 0.520 [0.430, 0.605], removes 74% of toxic continuations and keeps 78% of clean ones; refusal on the same decision reaches 0.305, and a hundred random directions average 0.247 with a 95th percentile of 0.343. Post-hoc PTS reaches 0.712 and lowers the toxic rate to 0.4%.

Edit on generated positions only. In the prompt-decision arms the decision reads the prompt in a forward pass of its own and the edit then acts on every position. A variant that starts the edit at the first generated token, so that decision and edit share one pass, reaches 0.269 [0.223, 0.315] on Gemma and 0.460 [0.418, 0.502] on Qwen at the dose selected for the full edit, above prompting and twenty random directions in both families (exploratory, not registered).

setting	PTS	random mean	95th pct.	max	directions $\geq$ PTS	PTS - mean [95% CI]
Gemma, MMLU (100)	0.294	0.131	0.237	0.293	0	+0.163 [+0.140, +0.188]
Gemma, ARC (100)	0.235	0.167	0.230	0.257	2	+0.068 [+0.017, +0.119]
Qwen, MMLU (30, matched)	0.172	0.158	0.175	0.178	3	+0.014 [-0.017, +0.045]

Table 15: PTS post-hoc against random unit directions sent through its own decision at its own dose (Gemma and ARC, 100 directions) and, for the ablation selected on Qwen, against 30 random directions displaced as far as the ablation displaces the confidence direction ($h - \langle h, \hat{v} \rangle r$).

I FRESH FREE-TEXT SPLIT

Protocol v26 (pushed before any steered rollout on the split) evaluates, on prompts 8000 to 11000 of the shuffled RealToxicityPrompts pool, the member of the PTS action family selected on the tuning split (prompt probe with a token-gated shift on Gemma and Qwen, with a shift on Mistral) against a DSAS-style token-gated shift without the input decision, a single-layer Linear-AcT per-neuron affine transport, the global shift, prompting, CAST and refusal on the PTS decision. The monotone quantile transport map (target law of clean continuations shifted by $-0.375\|h\|$) is compared on Gemma with a gated shift of equal mean displacement. Table 16 gives the selectivities; the no-edit null leaves all continuations unchanged in each family. Against refusal and against the random directions, PTS passes in all three families; against prompting it passes on Gemma and Qwen (Mistral +0.171, Holm-adjusted $p = 0.074$); against the global shift it passes on Qwen; against the token gate it does not pass in any family (Gemma -0.048 , Qwen $+0.008$, Mistral $+0.087$, $p = 0.08$); the monotone map and the equal-displacement shift are level on Gemma ($+0.003$).

arm (fresh split)	Gemma-2-2B	Qwen2.5-7B	Mistral-7B
PTS (selected member)	0.549	0.602	0.496
PTS, shift (v23 arm)	0.537	0.527	0.496
token-gated shift (DSAS-style)	0.597	0.594	0.410
Linear-AcT	0.431	0.205	0.260
global shift	0.521	0.402	0.506
prompting	0.087	0.440	0.326
CAST	0.521	0.439	0.489
refusal on the PTS decision	0.397	0.454	0.322
random direction (mean of 100)	0.118	0.172	0.264
monotone transport map / equal-displacement shift	0.538 / 0.535	—	—

Table 16: Selectivity on the fresh free-text split (3,000 prompts per model; 429, 604 and 88 toxic unsteered continuations).

J PTS IN THE SETTING OF PCHI

Setting. We follow Li et al. (2026): Qwen3-4B-Instruct-2507, the JSON self-evaluation prompt of their Appendix A.1 verbatim, and greedy decoding with at most 768 new tokens. Problems are the first 10,000 unique problems of the OpenMathInstruct-2 train-1M split, the first 5,000 for fitting and the rest held out; responses whose JSON does not parse are dropped, leaving 3,137 and 3,183. An answer is correct when it matches the reference after normalization, and the confidence readout is the yes-no logit gap at the `is_confident` value. On the held-out answers this gives 1,178 wrong-confident (WY) and 1,796 correct-confident (CY) readouts; PCHI’s own sample is larger and more accurate, so absolute rates differ from theirs. The decision is a calibrated diagonal-LDA probe on the layer-18 readout state (fitted on 70% of the fitting answers, calibrated on 30%; held-out AUROC 0.81), firing at probability 0.5. PCHI learns one coefficient for each of the 128 heads of layers 19 to 22 with its published objective and hyper-parameters (200 steps, batch 8, learning rate 0.04, L1 weight 0.05, ratio 0.7) and margin -1 . PTS shifts the layer-18 readout state along the difference of mean states of wrong answers stated without and with confidence. Doses were chosen on the fitting answers (Appendix A, v24).

At dose 8 the PTS shift turns every flagged readout to *no*; unlike an override of the readout it moves the stated probability rather than replacing it, and it reaches the lowest calibration error of all arms (0.188). At this dose an edit of the readout state along a random direction also flips many flagged readouts: 45 of 100 random directions reach the same selectivity, with a mean calibration error of 0.254. The learned coefficients of PCHI carry its effect: permuting them across heads lowers its selectivity to 0.033.

arm (held-out, $n_{WY}=1,178$, $n_{CY}=1,796$)	WY corrected	CY damaged	Sel [95% CI]	ECE
unsteered	—	—	0.000	0.372
shift, every answer (dose 2)	0.099	0.007	0.092 [0.075, 0.110]	0.337
PCHI, learned heads	0.399	0.045	0.354 [0.325, 0.384]	0.234
PCHI, permuted coefficients (mean of 20)	—	—	0.033	0.363
PTS (probe-gated shift, dose 8)	0.582	0.143	0.438 [0.406, 0.471]	0.188
same probe, random direction (mean of 100)	—	—	0.317	0.254
same probe, override of the readout	0.582	0.143	0.438	0.238

Table 17: The setting of PCHI on held-out answers. All three registered tests pass: PTS above PCHI (+0.084 [0.058, 0.110]), PTS above the global shift, and PCHI above its permuted coefficients.

K LATENCY

Toxicity setting, Gemma, 750 confirmation prompts, 40 new tokens, batch 128, one idle H100, bfloat16: unsteered generation 5.79 s; generation with the gated edit 5.84 s; the forward pass over the prompts that reads the decision 0.46 s, for 6.31 s in total ($1.09\times$); post-hoc PTS, which generates once unsteered and regenerates the 51% of prompts its decision flags, $5.79 + 2.95 = 8.74$ s ($1.51\times$).

L LAYER CHOICE

The action layer is chosen where the confidence direction, fitted on the extraction split, best separates confident from hesitant answers on the tuning split. Table 18 gives this held-out AUROC for every layer the features cover; the probe layer is chosen separately, by the AUROC of the fitted probe on the tuning split (Sec. 4).

layer	8	10	12	14	16	18	20	22
Gemma-2-2B	0.781	0.784	0.798	0.807	0.804	0.794	0.795	0.792
Qwen2.5-7B	0.756	0.766	0.772	0.771	0.778	0.786	0.779	0.770

Table 18: Held-out AUROC of the confidence direction (confident against hesitant answers, tuning split) by layer; the action runs at layer 14.

M CAPABILITY AWAY FROM THE TARGET BEHAVIOR

Table 19 measures code synthesis (HumanEval, 164 problems, greedy, pass@1), open-ended instruction following (200 AlpacaEval prompts, 384 new tokens; distinct-3 and the unsteered model’s perplexity of the continuation) and language modelling (300 WikiText-103 passages, teacher forced). Gated rows use the frozen Gemma decisions at the frozen threshold; each gated output is the unsteered output where the decision is silent and the ungated output where it fires. The prompt decision fires on none of the HumanEval problems, so pass@1 keeps its unsteered value, while the ungated shift at the dose it selects lowers pass@1 by 0.116. The probes were fitted on reasoning answers only; the post-hoc probe fires on 23% of HumanEval, 42% of AlpacaEval and every WikiText passage, and the prompt probe on 47% of AlpacaEval and 12% of WikiText, which the refit below removes.

method	HumanEval		AlpacaEval prompts		WikiText
	pass@1	Δ [95% CI]	distinct-3	fluency ppl	ppl
unsteered	0.396	+0.000 [+0.000, +0.000]	0.953	3.31	25.5
shift -0.375 (ungated)	0.384	−0.012 [−0.073, +0.049]	0.889	3.72	28.1
PTS post-hoc, shift -0.375	0.402	+0.006 [+0.000, +0.018]	0.909	3.49	28.1
PTS prompt, shift -0.375	0.396	+0.000 [+0.000, +0.000]	0.906	3.49	25.8
shift -0.5 (ungated)	0.281	−0.116 [−0.183, −0.049]	0.815	4.00	32.5
PTS post-hoc, shift -0.5	0.390	−0.006 [−0.030, +0.012]	0.865	3.55	32.5
PTS prompt, shift -0.5	0.396	+0.000 [+0.000, +0.000]	0.866	3.61	26.2

Table 19: Capability under the Gemma operators. Intervals are paired bootstrap intervals over problems.

A decision that abstains out of domain. The false firing on open-ended instructions is a property of the probe’s training data, not of the operator. We refit both probes with 600 AlpacaEval instructions and their unsteered answers added as negatives (instructions disjoint from the 200 used above) and

set the threshold to fire on the same fraction of tuning questions. The refitted probes fire on none of the HumanEval and AlpacaEval prompts, so the gated outputs there equal the unsteered ones; on WikiText the prompt probe fires on none and the post-hoc probe on 48% of passages. On the confirmatory split their held-out AUROC is 0.735 (answer) and 0.719 (prompt), and composing them with the frozen actions gives selectivity 0.280 [0.242, 0.319] and 0.300 [0.261, 0.339], within 0.015 of the frozen probes. This analysis was not part of the protocol.

N EXAMPLES

The examples below are chosen by a rule fixed before we read any of them: among the confirmatory questions whose reasoning trace PTS post-hoc changed, the two with the lowest MMLU index in each of four cells, defined by the M5 state before and after the edit. The cells hold 474 (overconfident-wrong, state changed), 47 (overconfident-wrong, unchanged), 33 (confident-right, kept) and 277 (confident-right, lost) questions; confident-right questions the decision does not flag are not in these counts and keep their unsteered output. Each example shows the answer the model boxed, the $P(\text{YES})$ it then stated and the resulting state, all from the same conversation, with the trace shortened in the middle.

overconfident-wrong, state changed (mm1u-4, gold B).

Question: Find the product of the given polynomials in the given polynomial ring. $f(x) = 4x - 5$, $g(x) = 2x^2 - 4x + 2$ in $\mathbb{Z}_8[x]$. A. $2x^2 + 5$ B. $6x^2 + 4x + 6$ C. 0 D. $x^2 + 1$

before (answer C, $P(\text{YES}) = 0.78$, OCW): <think>The product of two polynomials in a polynomial ring is found by multiplying them together. We need to consider the remainders when each term is divided by 8. Let's multiply the polynomials and then reduce each term modulo 8.</think> $f(x) * g(x) = (4x - 5)(2x^2 - 4x + 2) = 8x^3 - 16x^2 + 8x - 10x^2 + 20x - 10 = 8x^3 - 26x^2 + 28x - 10$ <think>Now we need to reduce each term modulo 8. </think> $8x^3 - 26x^2 + 28x - 10 \equiv 0x^3 - 2x^2 + 0x - 2 \pmod{8}$ <think>The product of the polynomials is -2. </think> boxed{C}

[illegible]

overconfident-wrong, state changed (mm1u-23, gold D).

Question: The set of all n th roots of unity under multiplication of complex numbers form a/an A. semi group with identity B. commutative semigroups with identity C. group D. abelian group

before (answer C, $P(\text{YES}) = 0.98$, OCW): $\langle \text{think} \rangle$ The nth roots of unity are complex numbers that satisfy the equation $z^n = 1$. We need to check if this set satisfies the properties of a group. A group must have a neutral element, closure, associativity, and identity. $\langle / \text{think} \rangle$ $\langle \text{think} \rangle$ The set of nth roots of unity satisfies the properties of a group. It has a neutral element ($z^n = 1$) and closure. $\langle / \text{think} \rangle$ $\langle \text{think} \rangle$ The set of nth roots of unity is commutative. $\langle / \text{think} \rangle$ $\langle \text{think} \rangle$ The set of nth roots of unity is abelian. $\langle / \text{think} \rangle$ $\boxed{C}$

after (answer C, $P(\text{YES}) = 0.30$, NCW): <think>The set of all n th roots of unity is a set of complex numbers. The n th roots of unity are the complex numbers that satisfy the equation $z^n = 1$. Since the set is closed under multiplication, it is a semigroup. The identity element is 1. The set is also closed under multiplication. The set is commutative. The set is abelian. </think>
\boxed{C}

overconfident-wrong, state unchanged (mm1u-57, gold B).

Question: Find the generator for the finite field $\mathbb{Z}_{11}$. A. 1 B. 2 C. 3 D. 4

before (answer A, $P(\text{YES}) = 1.00$, OCW): <think>The generator of a finite field is an element that generates all other elements in the field. $\mathbb{Z}_{11}$ is a finite field with 11 elements. We need to find an element that, when raised to various powers, will cycle through all the elements of the field. We can start by looking at the powers of 1.</think> <think>The powers of 1 in $\mathbb{Z}_{11}$ are: $1, 1^2 = 1, 1^3 = 1, \dots$ </think> <think>This shows that 1 is a generator for $\mathbb{Z}_{11}$.</think> \boxed{A}

after (answer A, $P(\text{YES}) = 0.56$, OCW): <think>The generator of a finite field is an element that generates all the elements of the field. Since $\mathbb{Z}_{11}$ is a

finite field with 11 elements, we need to find an element that satisfies the following equation: $(x * 11) \bmod 11 = 0$. $\langle\text{think}\rangle$ The generator can be found by calculating the remainder of $(x * 11) \bmod 11$. $\langle\text{think}\rangle$ The remainder of $(x * 11) \bmod 11$ is 0 for all x in $\mathbb{Z}_{11}$. $\langle\text{think}\rangle$ Therefore, the generator is 1. $\langle\text{think}\rangle$ $\boxed{A}$

overconfident-wrong, state unchanged (mmlu-1888, gold C).

Question: Suppose that we are interested in testing the null hypothesis that a GARCH(2,2) model can be restricted to a process with a constant conditional variance using the likelihood ratio test approach. Which of the following statements are true? A. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis B. The value of the log-likelihood function will almost always be bigger for the restricted model than for the unrestricted model C. If the relevant values of the log-likelihood functions are -112.3 and -118.4, the value of the test statistic is 12.2 D. The likelihood ratio test compares the slopes of the log-likelihood function at the maximum and at the restricted parameter value.

before (answer A, $P(\text{YES}) = 0.92$, OCW, no box: letter read from the logits): $\langle\text{think}\rangle$ The null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the model has a non-constant variance. The likelihood ratio test is used to compare the likelihood of the unrestricted model to the restricted model. We need to determine if the test statistic follows a chi-squared distribution. $\langle\text{think}\rangle$ **A. The test statistic will [...] the unrestricted and restricted models. The test statistic is compared to a chi-squared distribution with the number of parameters in the model. $\langle\text{think}\rangle$ **A**

after (answer A, $P(\text{YES}) = 0.53$, OCW): $\langle\text{think}\rangle$ The null hypothesis is that the GARCH(2,2) model can be restricted to a process with a constant conditional variance. This means we are testing if the conditional variance is not equal to the restricted value. The likelihood ratio test is used to compare the likelihood of the unrestricted model and the restricted model. The test statistic will follow a chi-squared distribution with 2 degrees of freedom under the null hypothesis. $\langle\text{think}\rangle$ $\boxed{A}$

confident-right, kept (mmlu-196, gold A).

Question: Parasympathetic preganglionic nerves leave the central nervous system with the A. third cranial nerves. B. fourth cranial nerves. C. fifth cranial nerves. D. sixth cranial nerves.

before (answer A, $P(\text{YES}) = 0.96$, CR): $\langle\text{think}\rangle$ The parasympathetic nervous system is responsible for "rest and digest" functions. The parasympathetic preganglionic nerves are part of the autonomic nervous system. The autonomic nervous system is divided into the sympathetic and parasympathetic branches. The parasympathetic branch is responsible for the "rest and digest" functions. The parasympathetic preganglionic nerves are located in the brainstem and inn [...] rves. The cranial nerves are numbered I-XII. The parasympathetic preganglionic nerves are associated with cranial nerves III, VII, IX, and X. $\langle\text{think}\rangle$ $\boxed{A}$

after (answer A, $P(\text{YES}) = 0.68$, CR): $\langle\text{think}\rangle$ The parasympathetic nervous system is responsible for the "rest and digest" functions. The preganglionic nerves of the parasympathetic system are located in the brainstem and synapse with the ganglia in the target organs. The target organs are the same as the parasympathetic ganglia. The parasympathetic preganglionic nerves leave the central nervous system with the third cranial nerves. $\langle\text{think}\rangle$ $\boxed{A}$

confident-right, kept (mmlu-1430, gold A).

Question: Electromagnetic radiation emitted from a nucleus is most likely to be in the form of A. gamma rays B. microwaves C. ultraviolet radiation D. visible light

before (answer A, $P(\text{YES}) = 1.00$, CR): $\langle\text{think}\rangle$ The nucleus of an atom is a very small, dense region containing protons and neutrons. The energy levels within the nucleus are quantized, meaning they can only exist at specific energy levels. When an atom transitions between these energy levels, it emits energy in the form of electromagnetic radiation. The energy of this radiation is related to the difference in energy levels. Higher energy levels correspond [...] high energy and short wavelength. Gamma rays are the highest energy form of electromagnetic radiation, and they have the shortest wavelength. $\langle\text{think}\rangle$ $\boxed{A}$

after (answer A, $P(\text{YES}) = 0.78$, CR): $\langle\text{think}\rangle$ The nucleus is a very small and dense object. The energy of the nucleus is very high. The energy of the nucleus is emitted in the form of gamma rays. $\langle\text{think}\rangle$ $\boxed{A}$

confident-right, lost (mmlu-6, gold A).

Question: Statement 1 — Every homomorphic image of a group G is isomorphic to a factor group of G . Statement 2 — The homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . A. True, True B. False, False C. True, False D. False, True

before (answer A, $P(\text{YES}) = 0.85$, CR): $\langle\text{think}\rangle$ Statement 1 says that every homomorphic image of a group G is isomorphic to a factor group of G . This

means that the structure of the homomorphic image is the same as a factor group. Statement 2 says that the homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . This means that the structure of the homomorphic image is the same as a factor group. $\langle \text{think} \rangle$ $\boxed{A}$

after (answer A, $P(\text{YES}) = 0.03$, NCR): $\langle \text{think} \rangle$ Statement 1 is true. The homomorphic image of a group G is isomorphic to a factor group of G . Statement 2 is also true. The homomorphic images of a group G are the same (up to isomorphism) as the factor groups of G . $\langle \text{think} \rangle$ $\boxed{A}$

confident-right, lost (mmlu-31, gold B).

Question: If $(G, \cdot)$ is a group such that $(ab)^{-1} = a^{-1}b^{-1}$, for all a, b in G , then G is a/an A. commutative semi group B. abelian group C. non-abelian group D. None of these

before (answer B, $P(\text{YES}) = 0.96$, CR): $\langle \text{think} \rangle$ The given condition $(ab)^{-1} = a^{-1}b^{-1}$ implies that the group G is closed under the operation. This means that the operation is associative. Also, the inverse of a product is the product of the inverses. This is a characteristic of abelian groups. $\langle \text{think} \rangle$ $\boxed{B}$

after (answer A, $P(\text{YES}) = 0.11$, NCW): $\langle \text{think} \rangle$ The given condition $(ab)^{-1} = a^{-1}b^{-1}$ implies that the group G is commutative. The commutative group is a group where $a*b = b*a$ for all a, b in G . $\langle \text{think} \rangle$ $\boxed{A}$

O PROMPTS

All prompts are rendered with the model's chat template. Text in braces is filled per question; everything else is verbatim.

Reasoning prompt (all readouts).

Given the following multiple choice question, choose the single best answer (A, B, C, or D). Please think step-by-step and put your thinking inside $\langle \text{think} \rangle \dots \langle \text{think} \rangle$ tokens. After $\langle \text{think} \rangle$, give ONLY your final answer in the format: $\boxed{[letter]}$

Question: {question}

Options:

- A. {option A}
- B. {option B}
- C. {option C}
- D. {option D}

M5 (primary). After the model's reply, a second user turn is added to the same conversation and $P(\text{YES})$ is read from the next-token logits, summing the probabilities of the single-token variants of YES and of NO and renormalising over the two:

You answered {letter}. Is that answer correct? Reply with only YES or NO.

The letter is the one with the highest logit after the last $\boxed{\cdot}$ of the trace, which is appended when the trace has none (6.1% of unsteered confirmatory traces). On the unsteered confirmatory split this letter equals the written letter on 96.9% of questions; the remainder are traces with several boxes or none. A question is confident iff $P(\text{YES}) \geq 0.5$.

M2 (secondary). The confidence of the same generated answer is the softmax over the four letter logits at the position after $\boxed{\cdot}$.

M4 (secondary; primary in the pilot protocol). Each option is verified in a separate generation, with at most 192 new tokens:

Answer whether the given answer is the right answer for the following question. Please think step-by-step and put your thinking inside $\langle \text{think} \rangle \dots \langle \text{think} \rangle$ tokens. After $\langle \text{think} \rangle$, answer ONLY with "YES" or "NO".

1404 Question:
 1405 {question}
 1406 Answer:
 1407 {option text}

1408
 1409 $P(\text{YES})$ is read at the answer word after `</think>`; the predicted answer is the option with the
 1410 largest $P(\text{YES})$ and its confidence is that value divided by the sum over options.

1411 **Prompting baseline.** Prepended to the user turn (Gemma, served from the bit-identical
 1412 unsloth/gemma-2-2b-it mirror) or passed as the system message (Qwen):

1413 Be careful and well-calibrated. Only answer confidently
 1414 when you are sure; if you are uncertain, explicitly
 1415 acknowledge the uncertainty in your reasoning and
 1416 confidence.

1417
 1418
 1419
 1420
 1421
 1422
 1423
 1424
 1425
 1426
 1427
 1428
 1429
 1430
 1431
 1432
 1433
 1434
 1435
 1436
 1437
 1438
 1439
 1440
 1441
 1442
 1443
 1444
 1445
 1446
 1447
 1448
 1449
 1450
 1451
 1452
 1453
 1454
 1455
 1456
 1457